

Introduction to Profile Areas in Bioinformatics

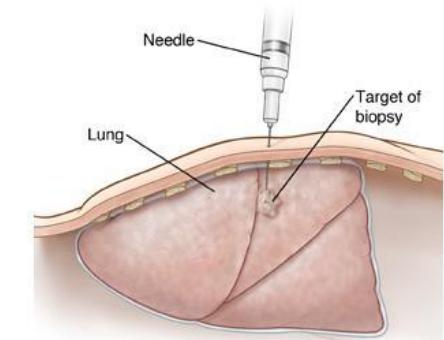
Lecturer: Kathleen Gallo

Session 3

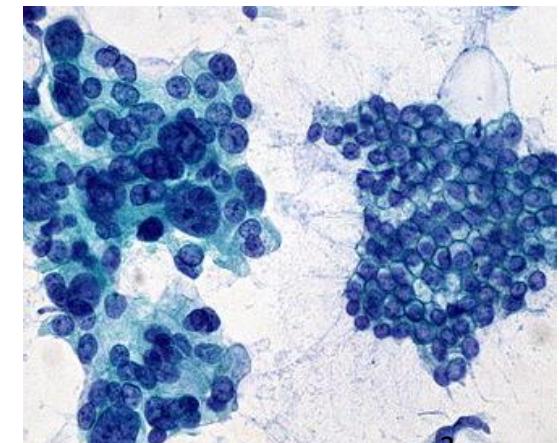
- kathleen.gallo@online.de
- Source code and report via mail until 13.11., 18h
- No peer review
- Tasks will be evaluated by me
- Short feedback at the beginning of the lecture on 8.11.

LAST WEEK's Project Overview

- Imagine, you are given data about patients having some suspicious tissue, e.g. in lung or breast.
- Features are computed from a digitized image of a fine needle aspirate (FNA) tissue. They describe characteristics of the cell nuclei present in the image.
- You want to build a tool that can help medical doctors to decide whether a sample is benign or malign.
- How would you do that?



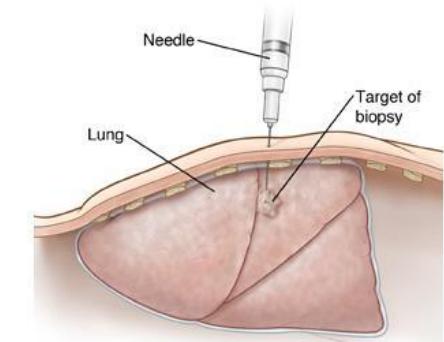
For the biopsy, a thin needle is used to remove samples of tissue from an abnormal area in the lung.



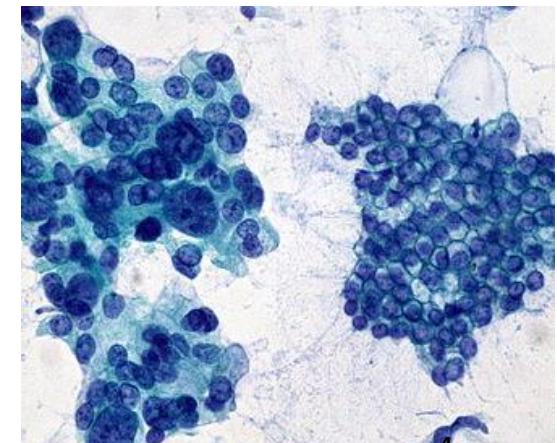


THIS WEEK's Project Overview

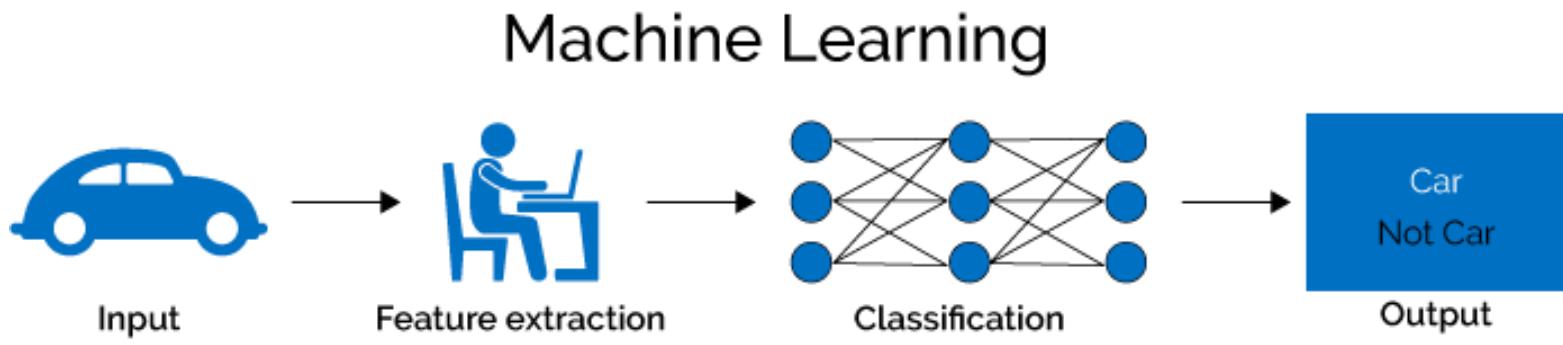
- Imagine, you are given data about patients having some suspicious tissue, e.g. in lung or breast.
- ~~Features are computed from a digitized image of a fine needle aspirate (FNA) tissue. They describe characteristics of the cell nuclei present in the image.~~
- You want to build a tool that can help medical doctors to decide whether a sample is benign or malign.
- How would you do that?



For the biopsy, a thin needle is used to remove samples of tissue from an abnormal area in the lung.

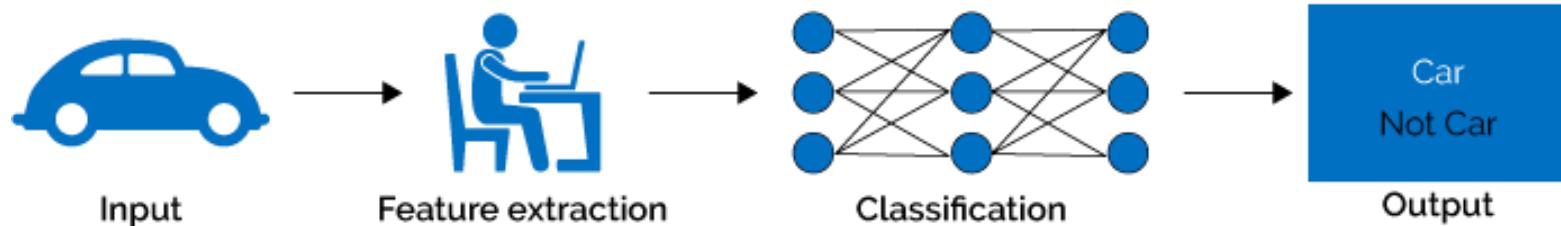


Two main approaches

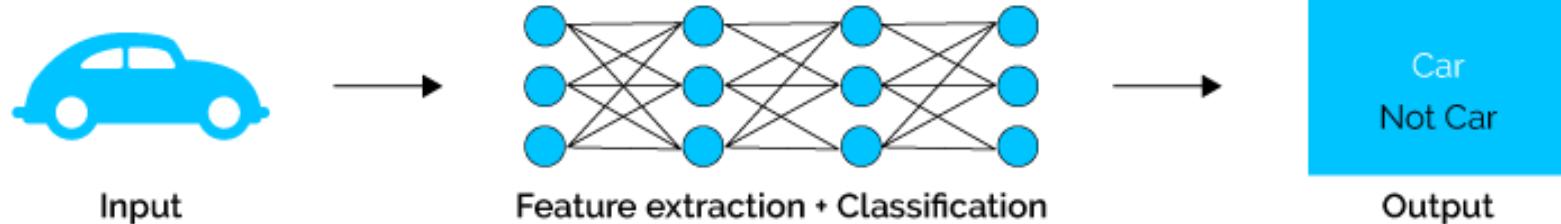


Two main approaches

Machine Learning



Deep Learning



THIS WEEK's Project Overview

- The Data:
 - LAST WEEK: Breast Cancer Wisconsin (Diagnostic) Data Set
 - THIS WEEK: A data-set containing microscopy images
- The task: develop a classifier to diagnose the state of a sample.
 - LAST WEEK: Using extracted features
 - THIS WEEK: Using the images directly
- What to deliver: The source code and a report.

Deep learning

From Wikipedia, the free encyclopedia

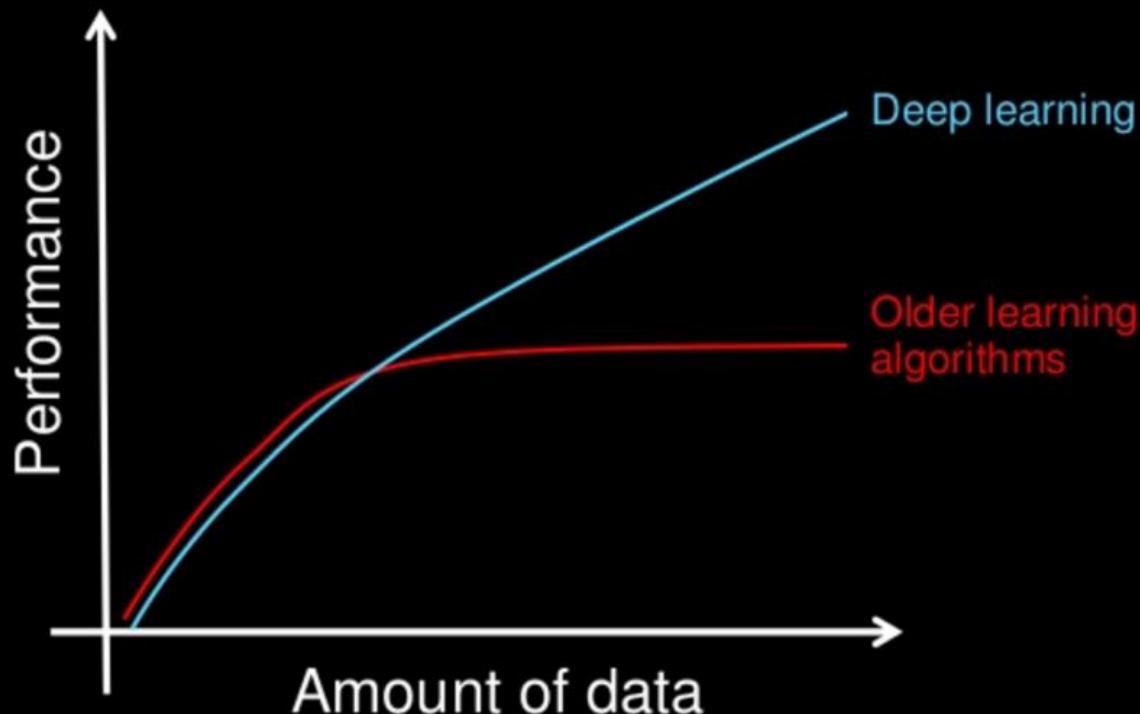
„.... is a branch of machine learning [...]“

“... has been characterized as a buzzword, or a rebranding of neural networks.”

“... is often presented as a step towards realizing strong [generalized] AI [...]”

Deep Learning

Why deep learning



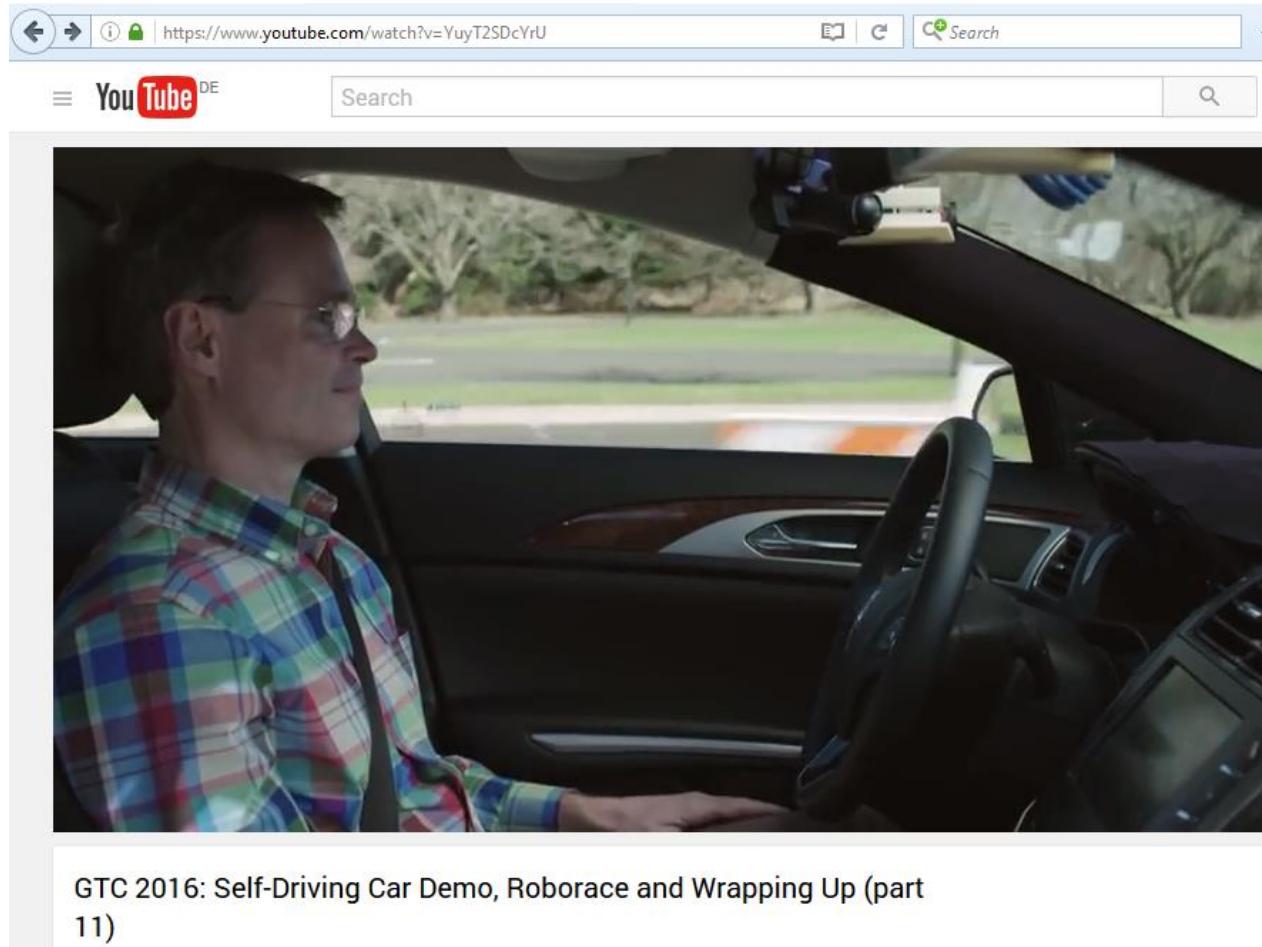
How do data science techniques scale with amount of data?

- “There’s a distinction [...] between **generalized AI** and **specialized AI**.
- In science fiction, what you hear about is **generalized AI**, right? Computers start getting smarter than we are [...]
- My impression, based on talking to my top science advisers, is that we’re still a reasonably long way away from that.”



<https://www.wired.com/2016/10/president-obama-mit-joi-ito-interview/>

Driving a Car – nothing but: $y=f(x)$?



GTC 2016: Self-Driving Car Demo, Roborace and Wrapping Up (part
11)

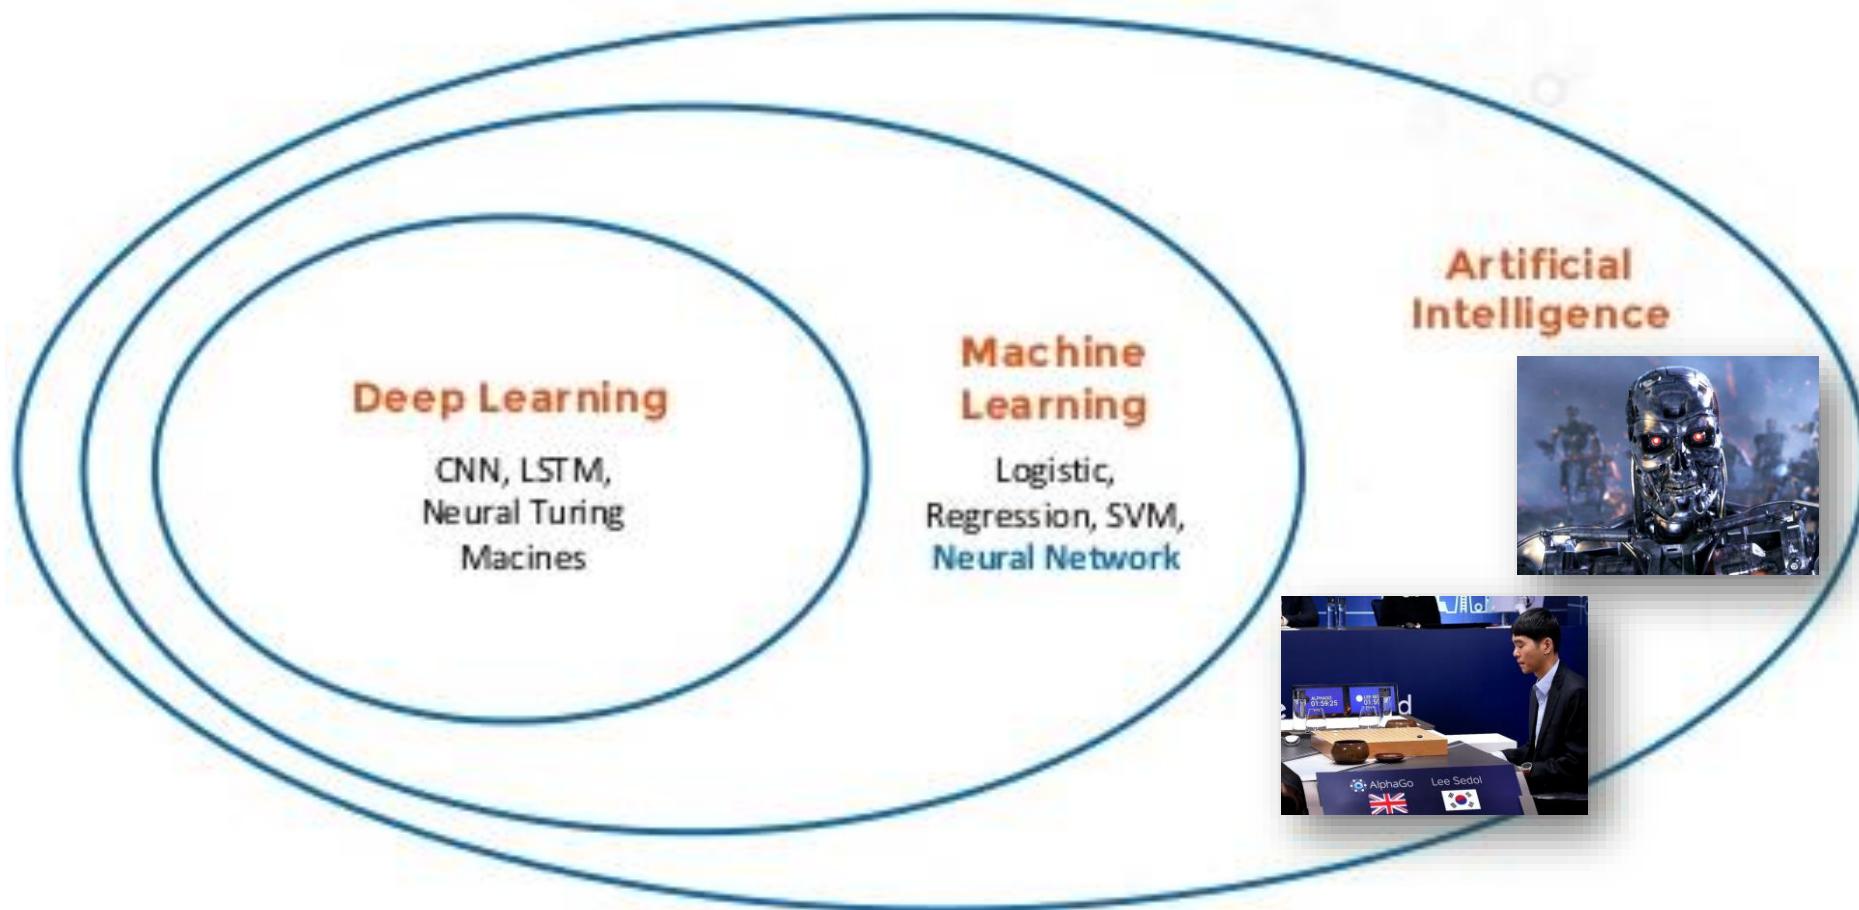
<https://www.youtube.com/watch?v=YuyT2SDcYrU>

- “[...] **specialized AI** [...] is about using algorithms and computers to figure out increasingly complex tasks.
- We've been seeing specialized AI in every aspect of our lives, from medicine and transportation to how electricity is distributed, [...]”



<https://www.wired.com/2016/10/president-obama-mit-joi-ito-interview/>

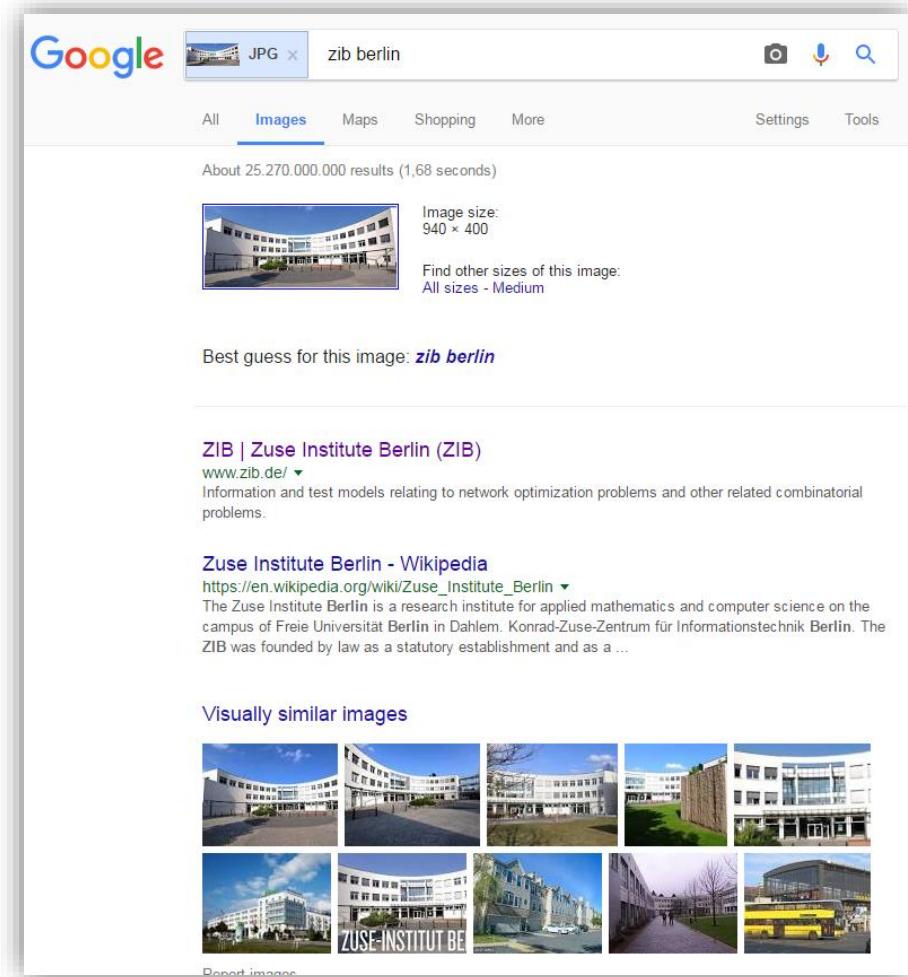
From AI to Deep Learning



Deep Learning for AI

The methods we are going to talk about today are used by several companies for a variety of applications, such as classification, retrieval, detection, etc.

It is especially well established in the area of computer vision.



Deep Learning for (specialized) AI

clarifai

Demo Solutions ▾ Pricing Developer API Resources ▾

Get your free API key Talk to us



GENERAL MODEL

PREDICTED CONCEPT	PROBABILITY
architecture	0.997
contemporary	0.987
modern	0.984
facade	0.984
house	0.979
horizontal plane	0.979
luxury	0.977
window	0.976
no person	0.970
outdoors	0.960
sky	0.957
lease	0.945
building	0.942

Deep Learning for (specialized) AI

clarifai

Demo Solutions ▾ Pricing Developer API Resources ▾

[Get your free API key](#) [Talk to us](#)



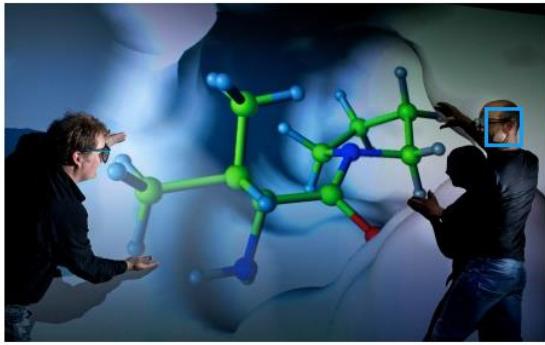
PREDICTED CONCEPT	PROBABILITY
adult	0.990
healthcare	0.986
hospital	0.982
man	0.980
woman	0.978
people	0.973
surgeon	0.973
nurse	0.962
medicine	0.959
portrait	0.953
medical practitioner	0.943
education	0.942
teamwork	0.941

Deep Learning for (specialized) AI

clarifai

Demo Solutions ▾ Pricing Developer API Resources ▾

[Get your free API key](#) [Talk to us](#)



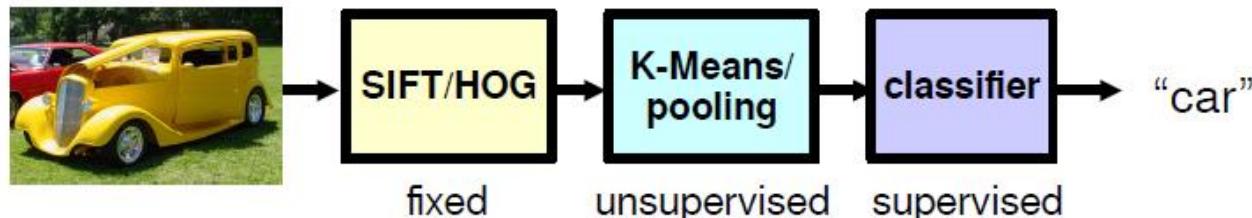
GENERAL MODEL

PREDICTED CONCEPT	PROBABILITY
people	0.986
music	0.963
business	0.959
research	0.958
man	0.955
medicine	0.934
adult	0.933
science	0.931
interaction	0.925

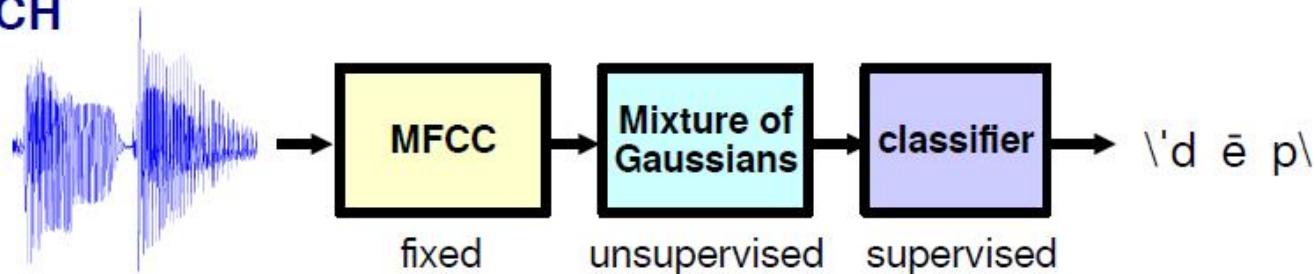
TRADITIONAL VS HIERARCHICAL LEARNING

Traditional Pattern Recognition

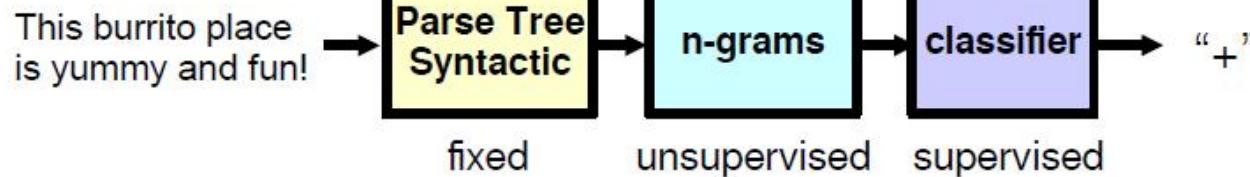
VISION



SPEECH



NLP



Hierarchical Compositionality (DEEP)

VISION

$$f(x) \approx g_1(g_2(\dots g_n(x)\dots))$$

pixels → edge → texton → motif → part → object

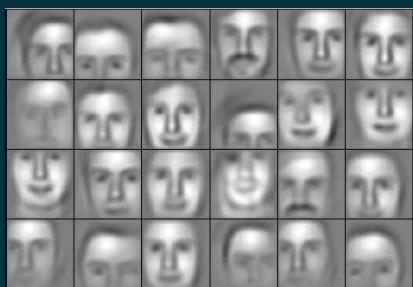
SPEECH

sample → spectral band → formant → motif → phone → word

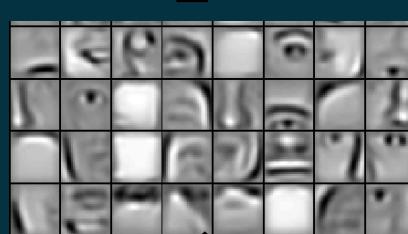
NLP

character → word → NP/VP/.. → clause → sentence → story

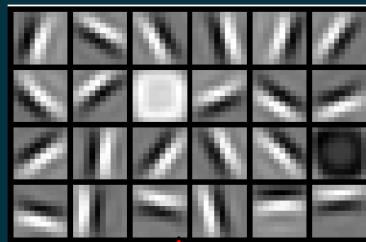
Deep Learning: learn hierarchical representations



object models



object parts
(combination
of edges)



edges

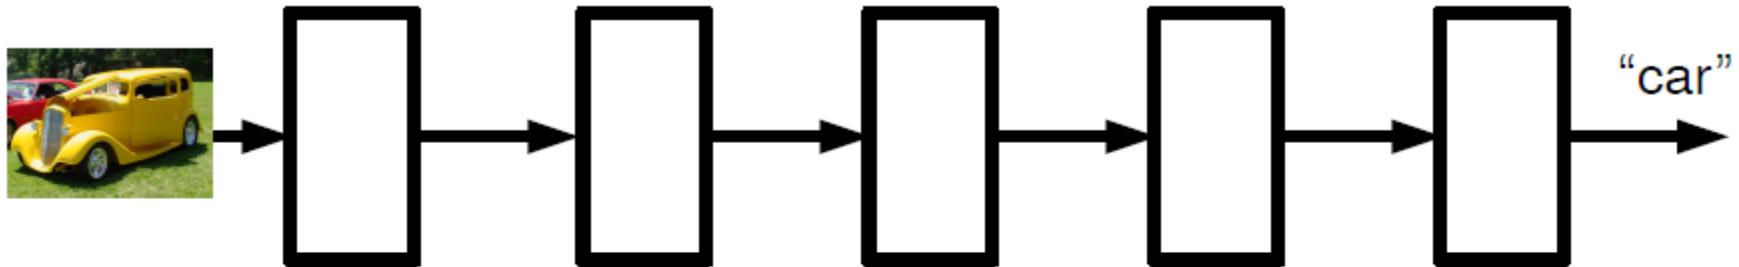


pixels



$$f(x) \approx g_1(g_2(\dots g_n(x)\dots))$$

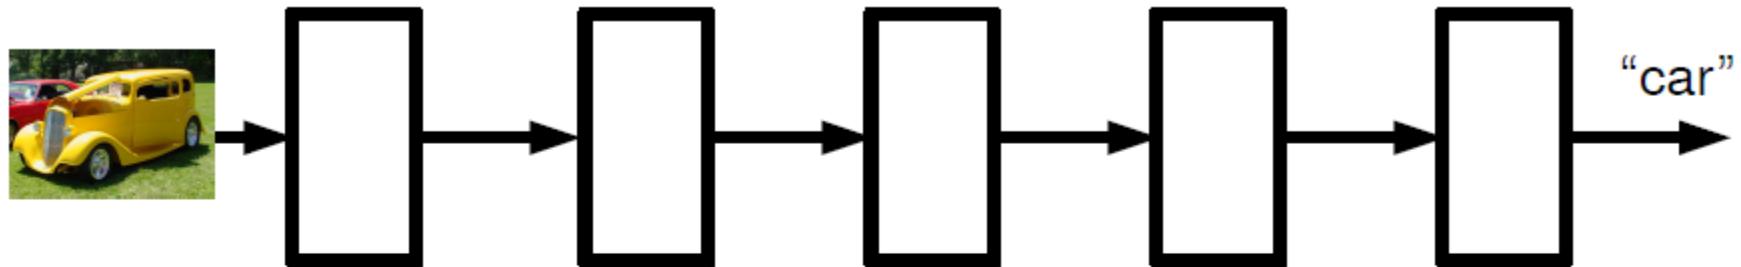
Deep Learning



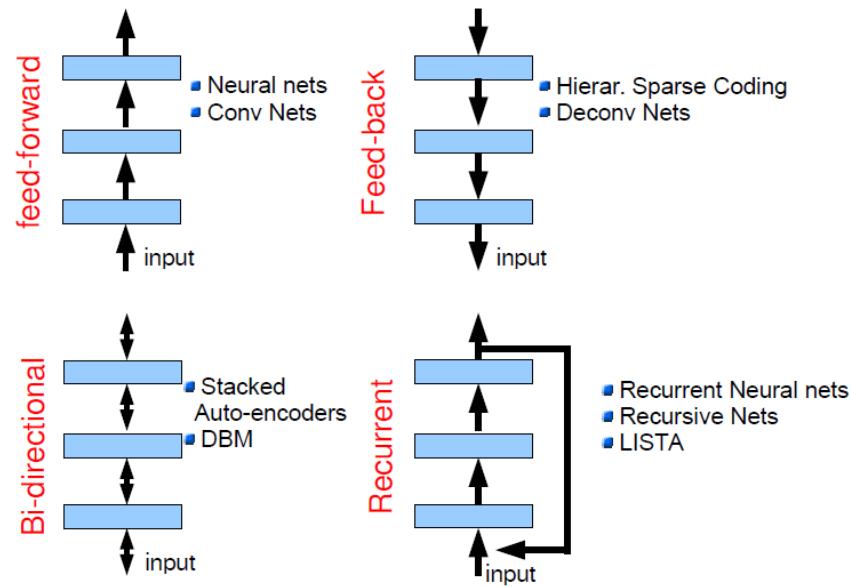
What is Deep Learning

- Cascade of non-linear transformations
- End to end learning
- General framework (any hierarchical model is deep)

Deep Learning



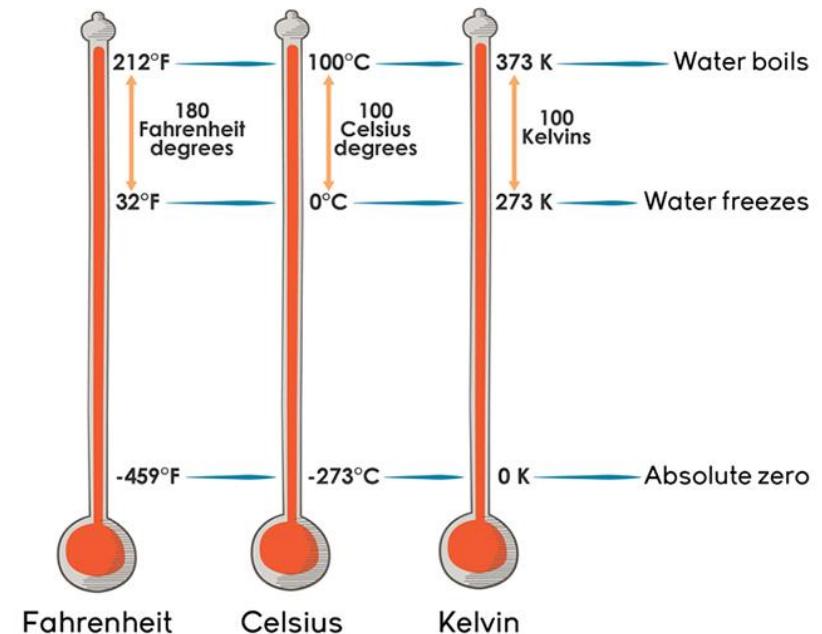
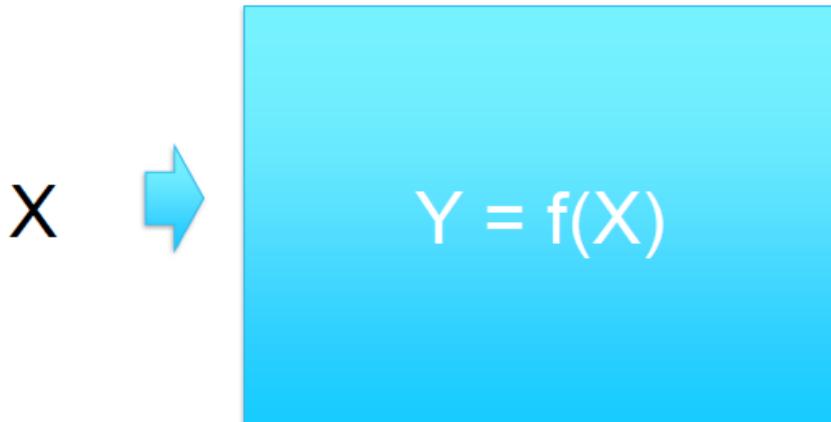
There are many architectures and flavors.



MACHINE LEARNING IS A
METHOD OF TEACHING COMPUTERS
TO MAKE PREDICTIONS BASED
ON SOME **DATA**.



Machine Learning: Find the Function



HERE'S HOW:

**Model-based vs. Data-driven
(or „deductive“ vs „inductive“)**

Classical (Deductive) Approach

Example: Converting Celsius to Fahrenheit

Requirements

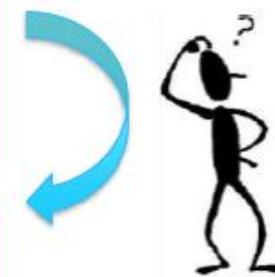
Input: C
Output: F
Where F is Fahrenheit equivalent of C in Celsius

Model

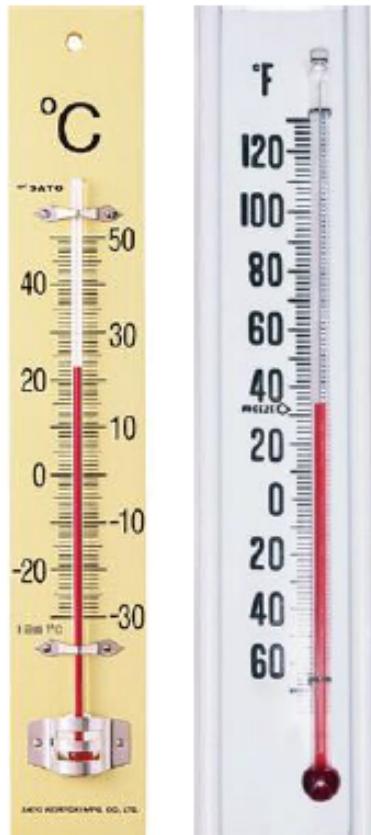
$$F = 1.8 * C + 32$$

Implementation

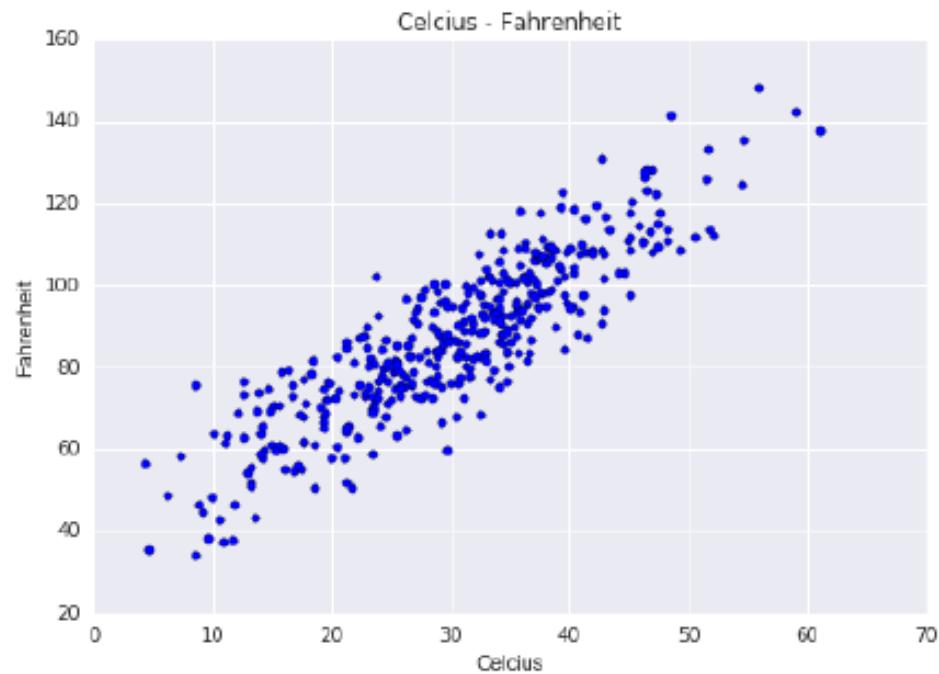
```
public class C2F {  
    public static void main(String[] args) {  
        java.util.Scanner sc = new java.util.Scanner(System.in);  
        double c = sc.nextDouble();  
        double f = 1.8 * c + 32.0;  
        System.out.println("F = " + f);  
        sc.close();  
    }  
}
```



Alternative: Data-driven, Inductive Approach

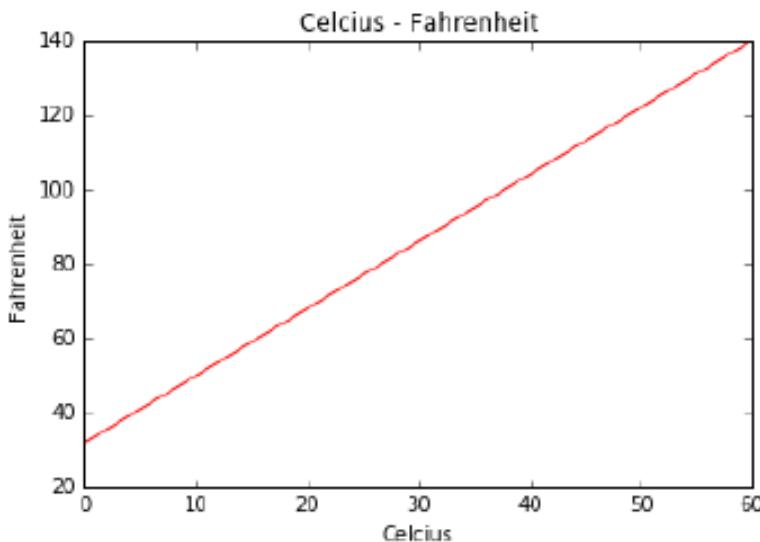


“Find function that represents this data”



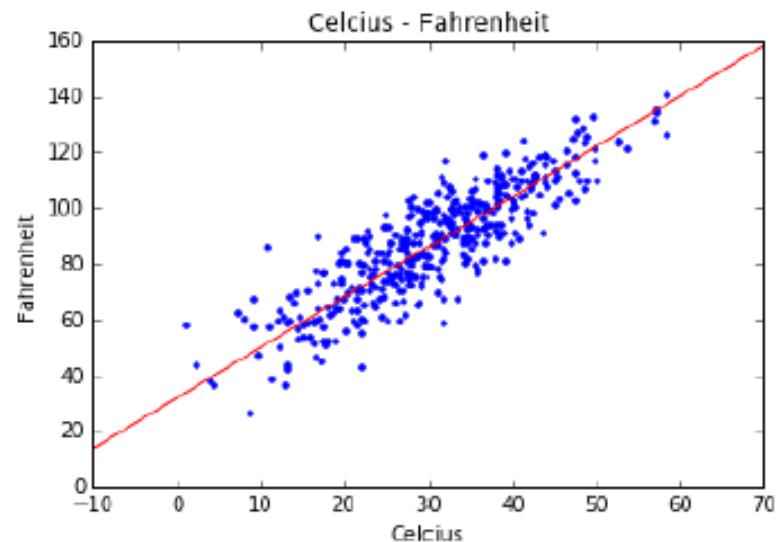
Machine Learning does exactly this

Definitive Function



$$F = 1.8 * C + 32$$

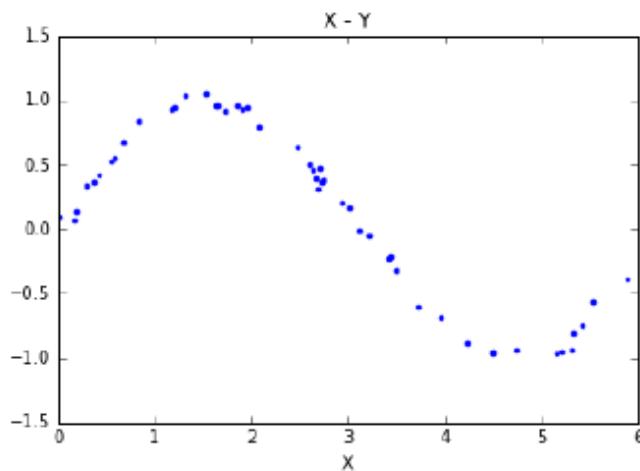
Stochastic Function



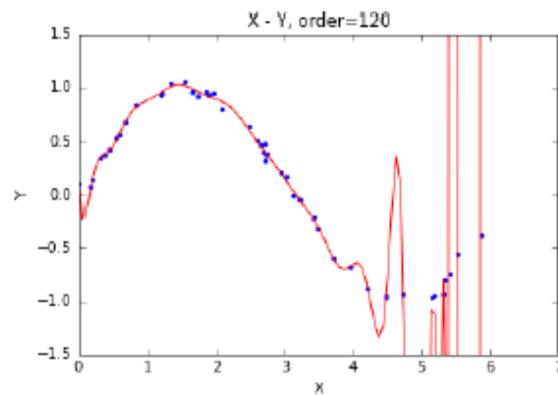
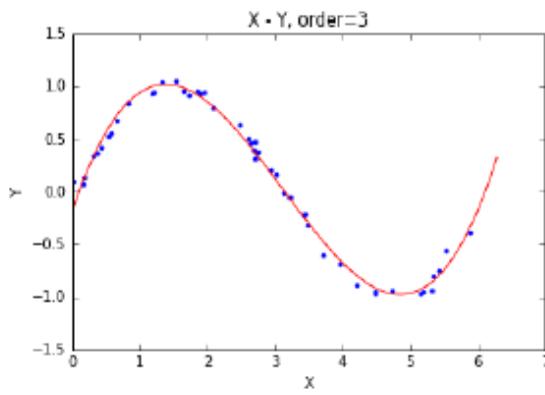
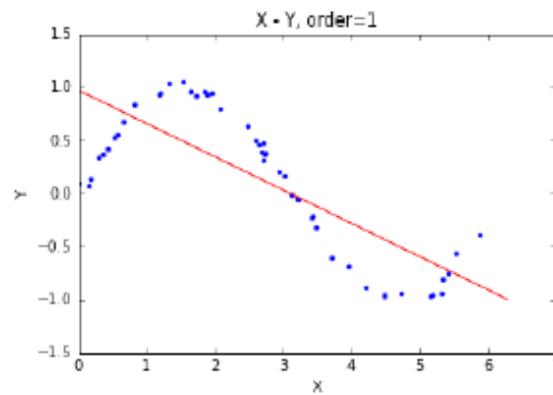
$$F = 1.8 * C + 32 + e$$
$$e \sim N(0, 10)$$

In (classical) Machine Learning, Model must be chosen a-priori

What is the function that represents this data?

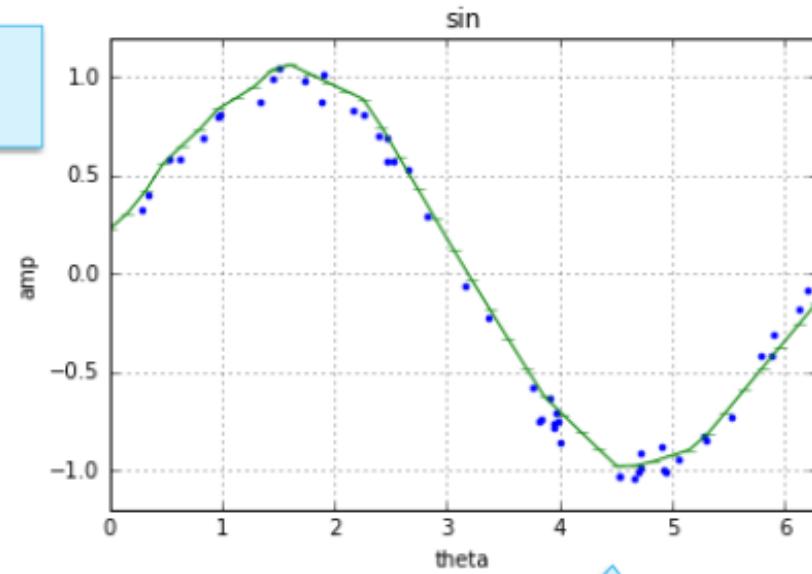
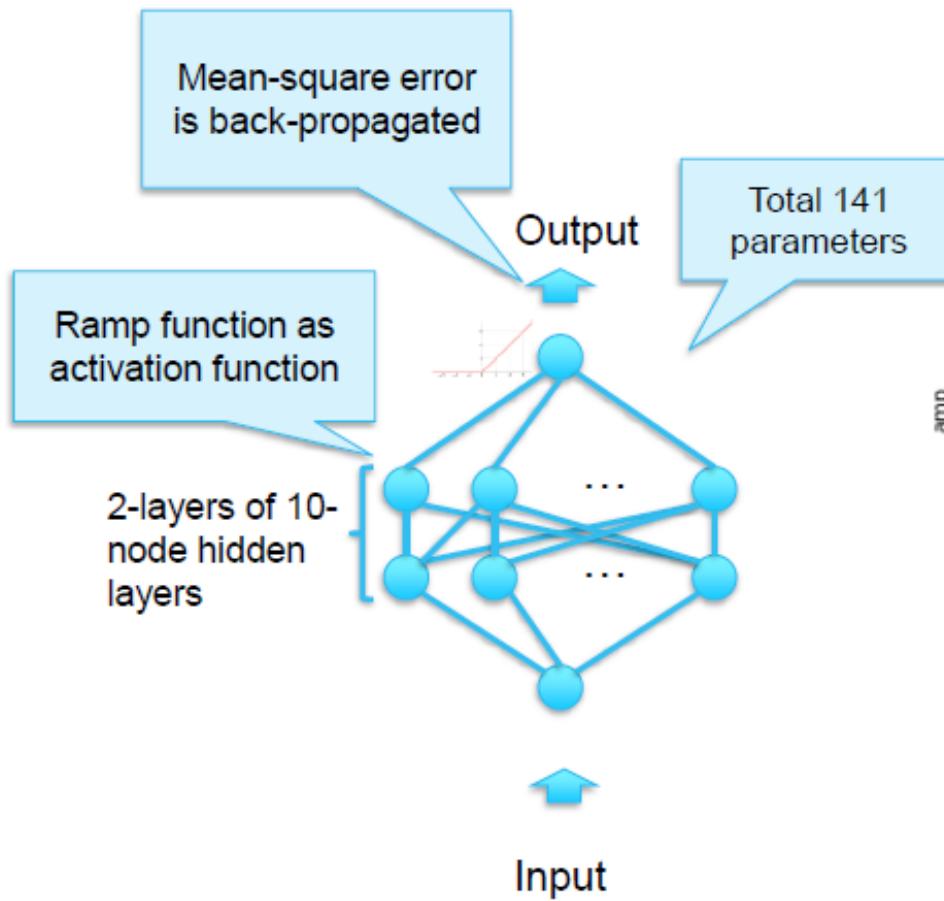


Too many parameters result in overfitting!



Choosing the right model is the key

Deep Learning can approximate any function without too much overfitting (mostly)



Good approximation without too much overfitting

So, 1. **what exactly is deep learning ?**

And, 2. **why is it generally better** than other methods on image, speech and certain other types of data?

The short answers

So, 1. **what exactly is deep learning ?**

'Deep Learning' **means** using a **neural network** with **several layers of nodes** between input and output

And, 2. **why is it generally better** than other methods on image, speech and certain other types of data?

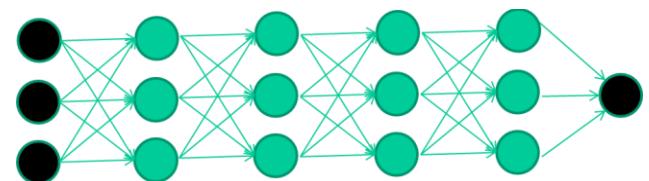
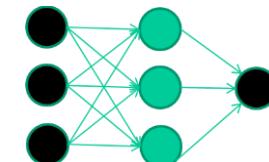
The series of layers between input & output do feature identification and processing in a series of stages, just as our brains seem to.

Multilayer neural networks have been around for 25 years.
What's actually new?

We have always had good algorithms for learning
the weights in networks with few (1) hidden layer

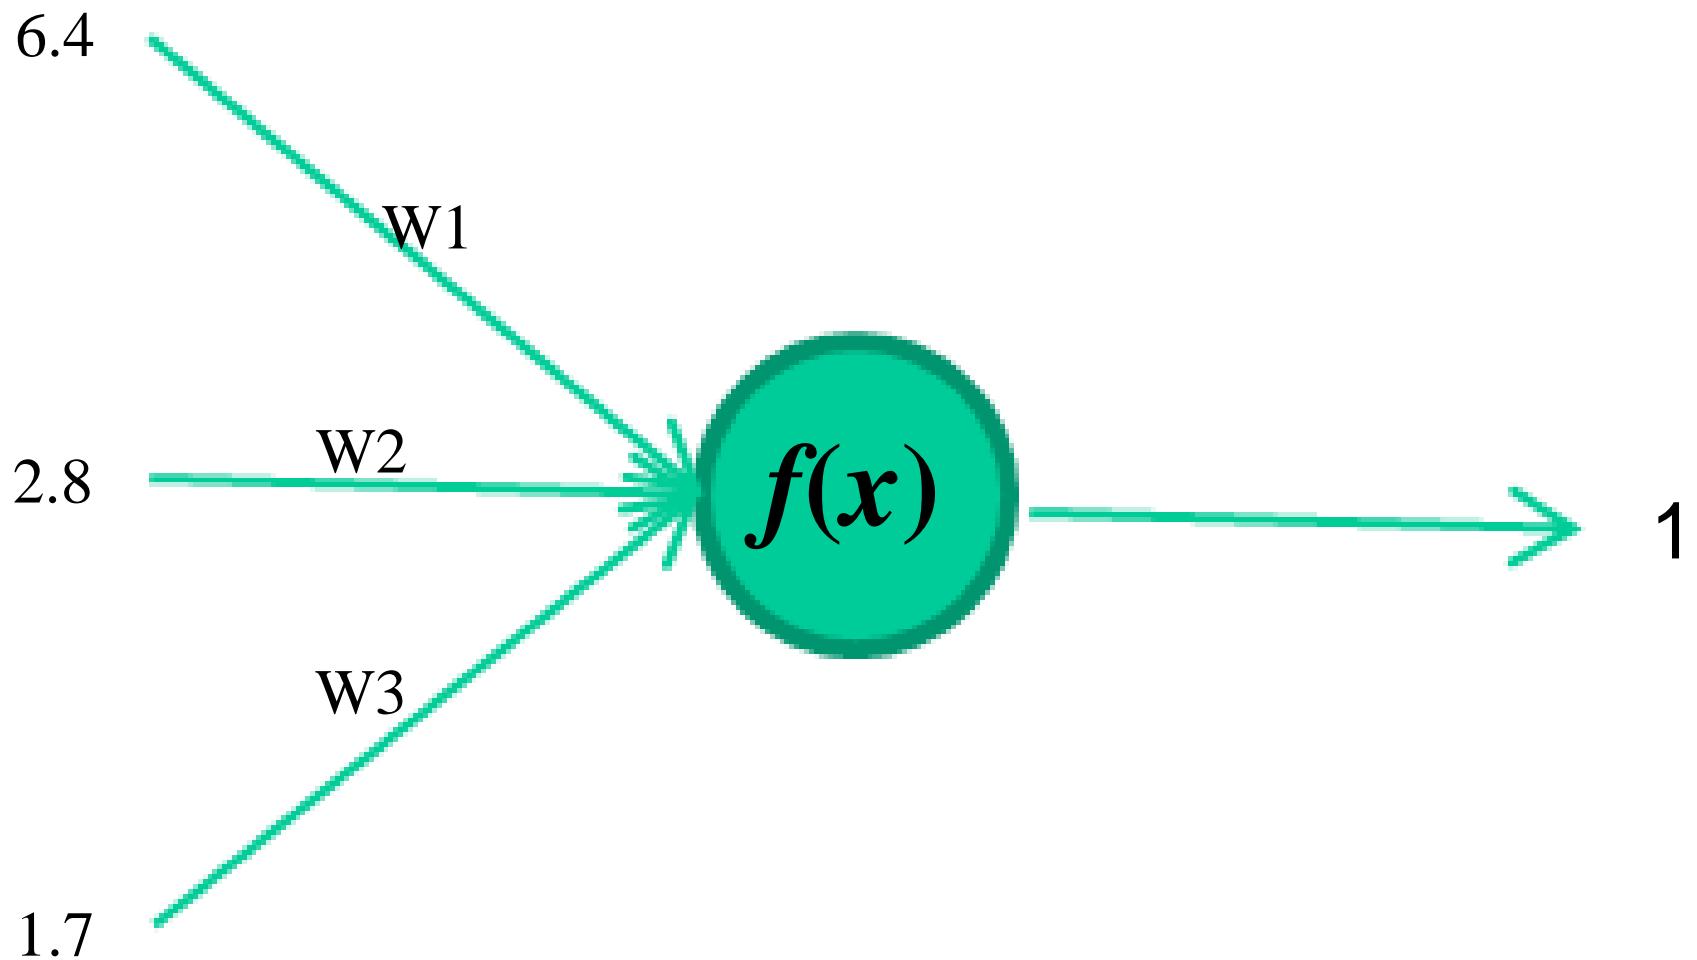
But these algorithms are not good at learning the
weights for networks with more hidden layers

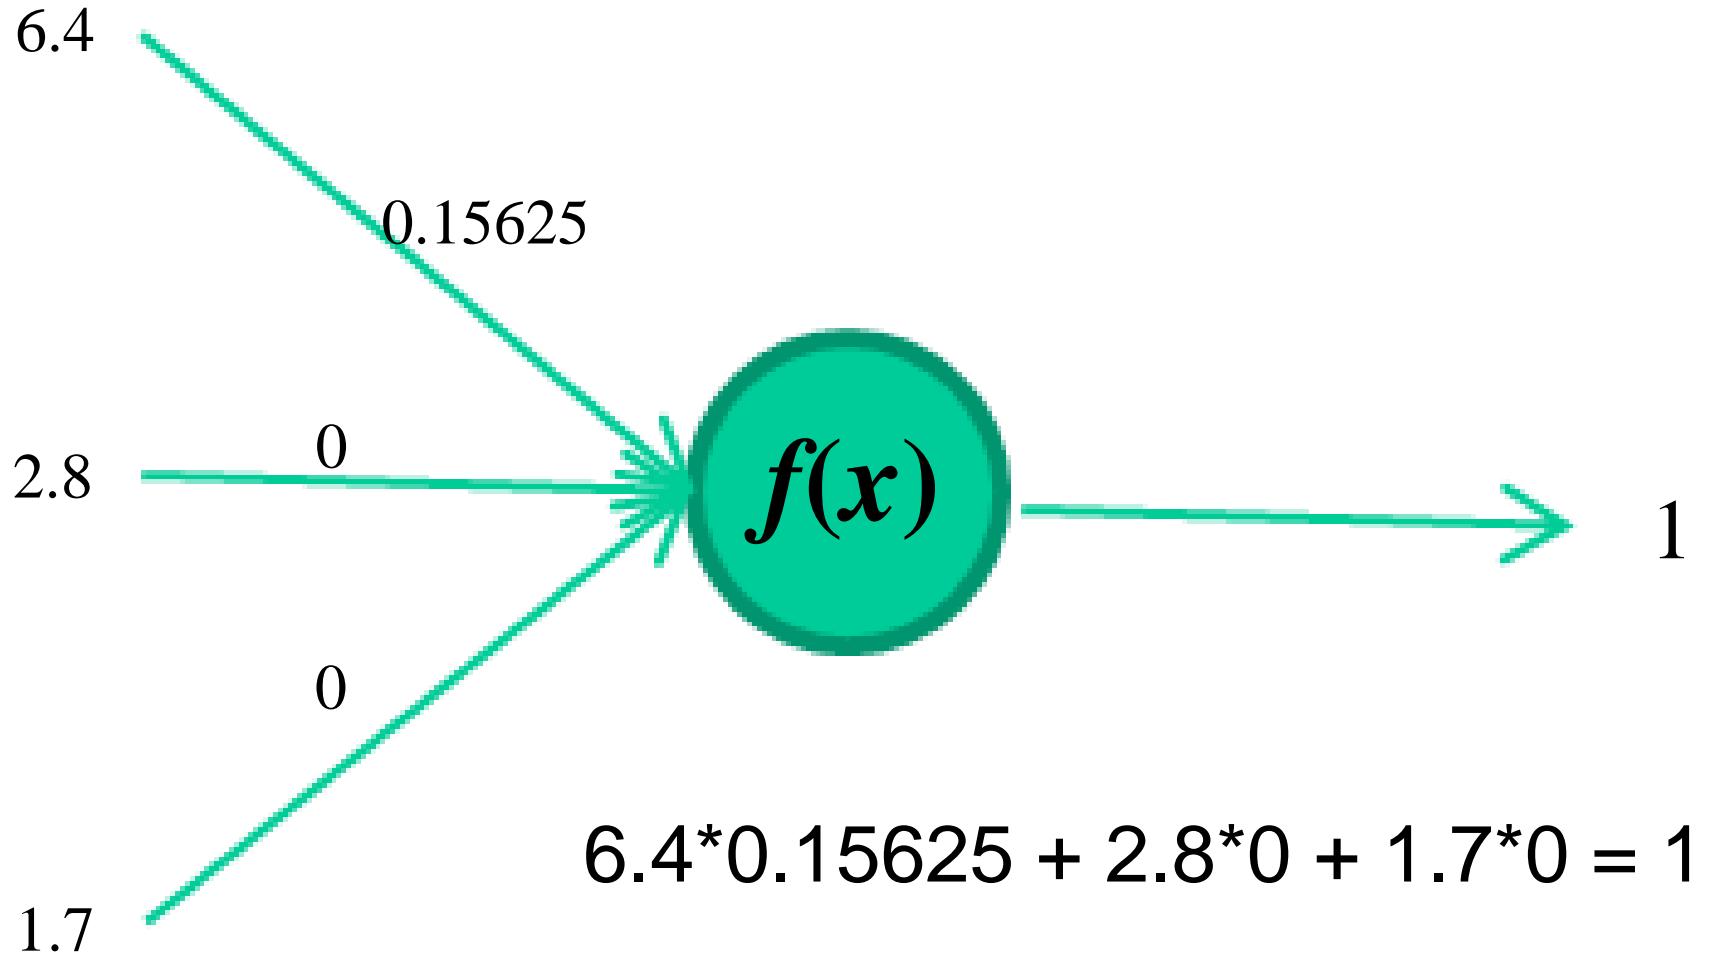
What's new is:
algorithms for training many-layer networks



Longer Answer

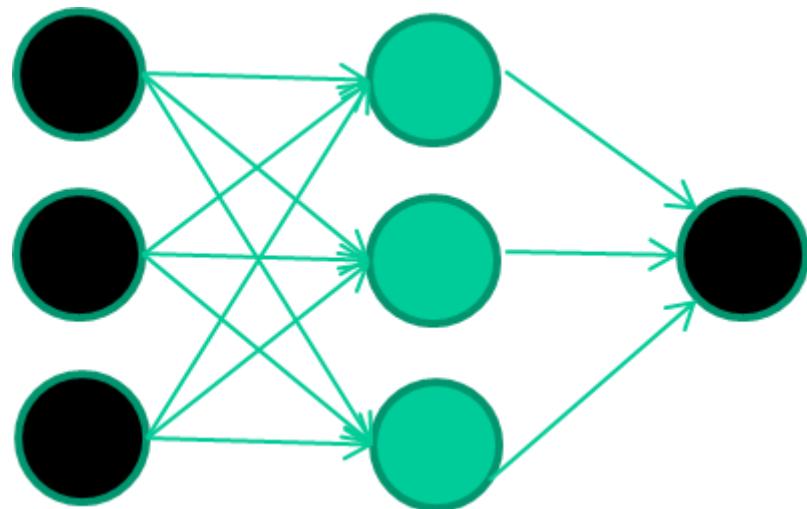
1. reminder/quick-explanation of how neural network weights are learned;
2. the idea of **unsupervised feature learning** (why ‘intermediate features’ are important for difficult classification tasks, and how NNs seem to naturally learn them)
3. The ‘breakthrough’ – the simple trick for training Deep neural networks





A dataset

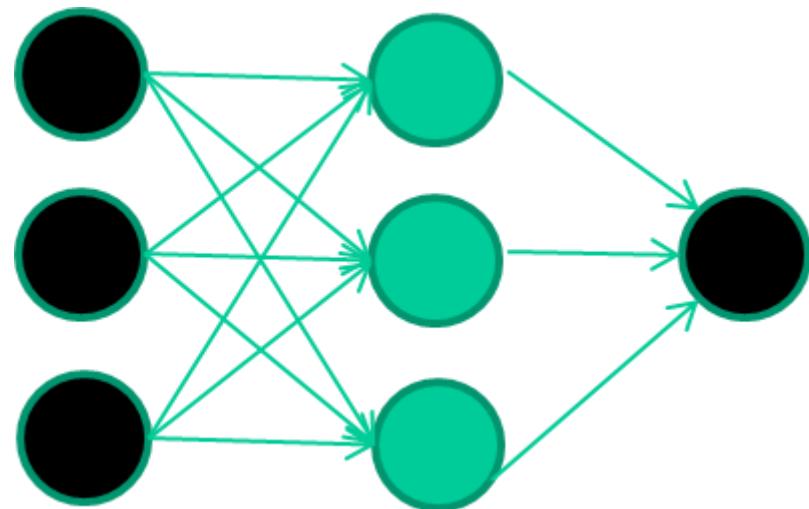
<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	



Training data

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	

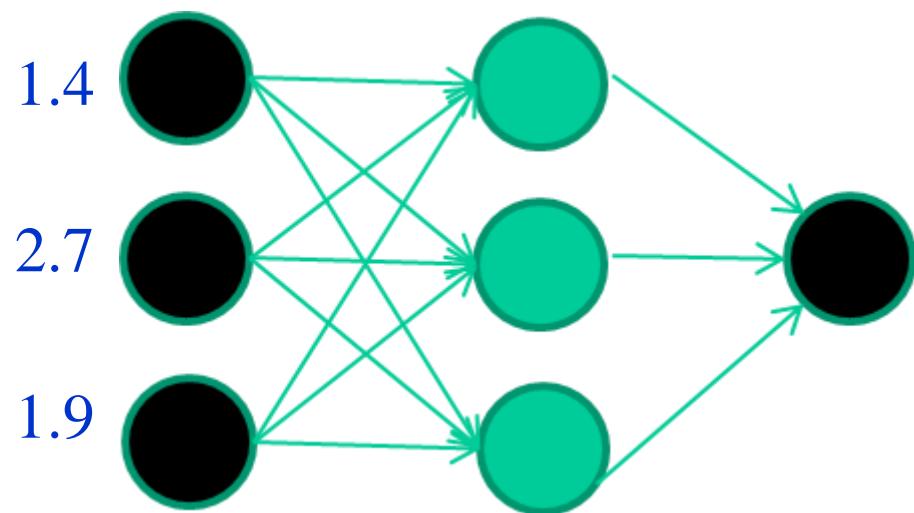
Initialise with random weights



Training data

<i>Fields</i>	<i>class</i>		
1.4	2.7	1.9	0
3.8	3.4	3.2	0
6.4	2.8	1.7	1
4.1	0.1	0.2	0
etc ...			

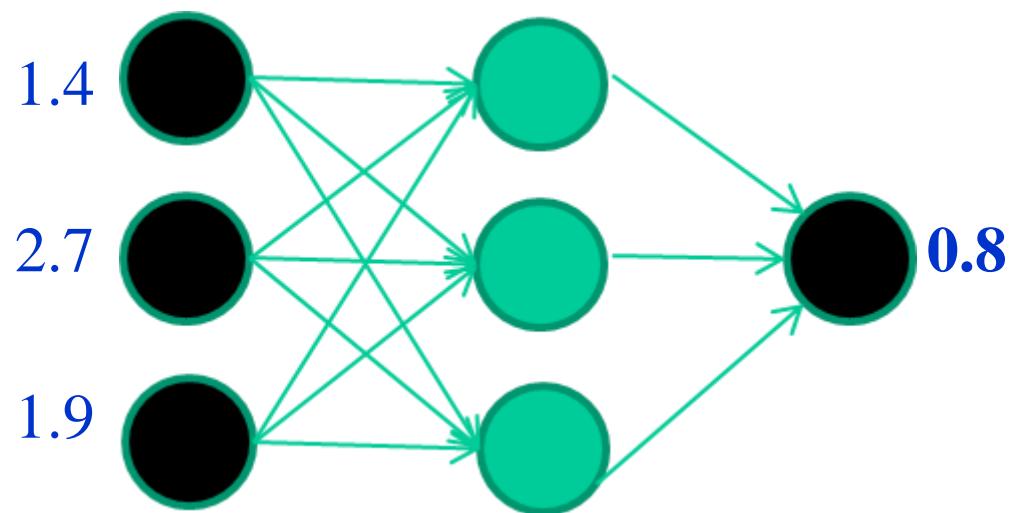
Present a training pattern



Training data

<i>Fields</i>	<i>class</i>		
1.4	2.7	1.9	0
3.8	3.4	3.2	0
6.4	2.8	1.7	1
4.1	0.1	0.2	0
etc ...			

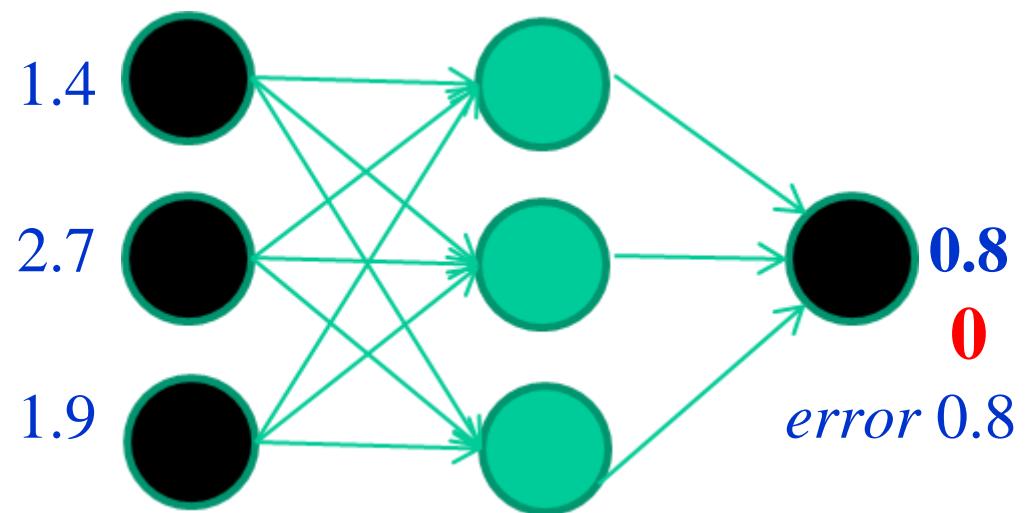
Feed it through to get output



Training data

<i>Fields</i>	<i>class</i>		
1.4	2.7	1.9	0
3.8	3.4	3.2	0
6.4	2.8	1.7	1
4.1	0.1	0.2	0
etc ...			

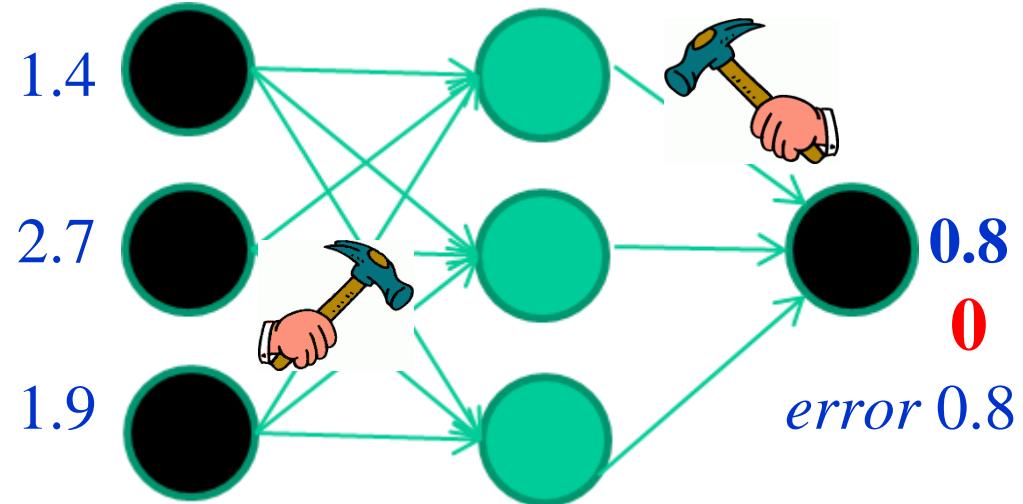
Compare with target output



Training data

<i>Fields</i>	<i>class</i>		
1.4	2.7	1.9	0
3.8	3.4	3.2	0
6.4	2.8	1.7	1
4.1	0.1	0.2	0
etc ...			

Adjust weights based on error



Training data

Fields *class*

1.4 2.7 1.9 0

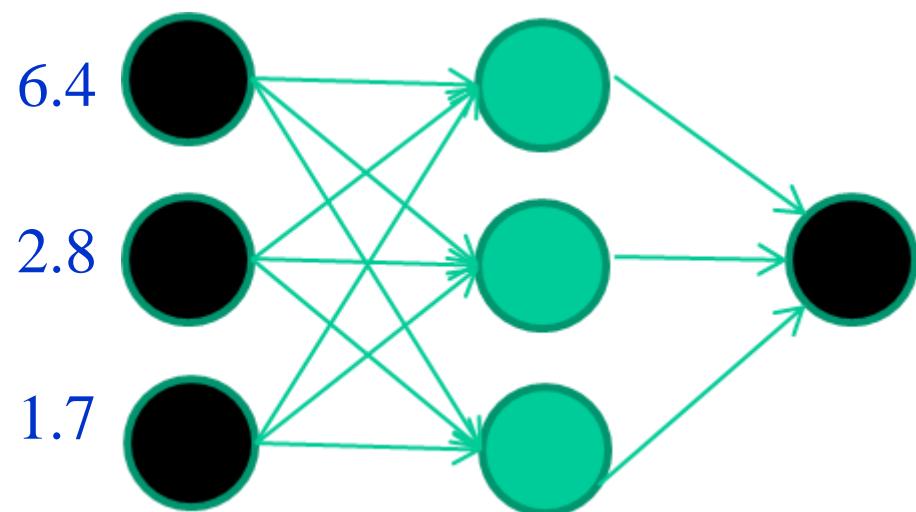
3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Present a training pattern



Training data

Fields *class*

1.4 2.7 1.9 0

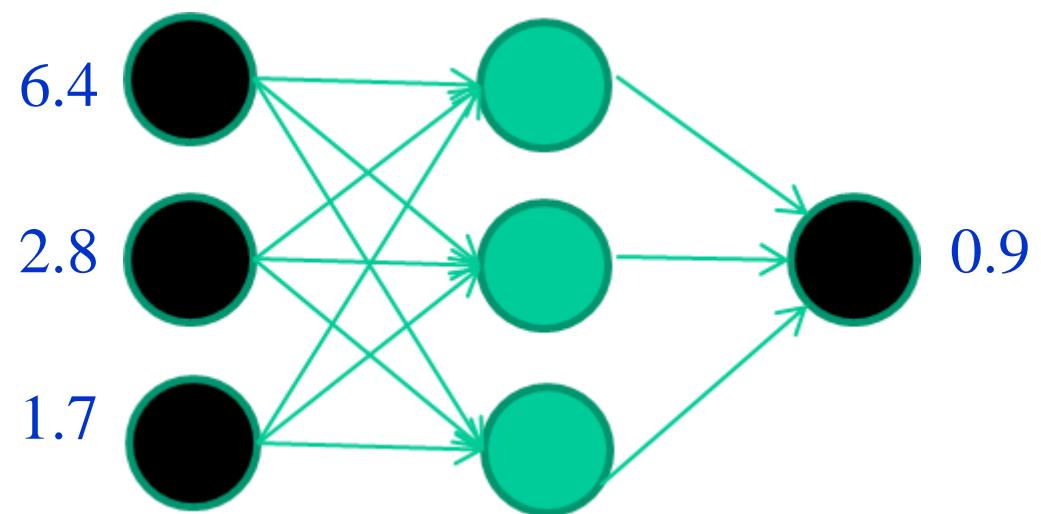
3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Feed it through to get output



Training data

Fields *class*

1.4 2.7 1.9 0

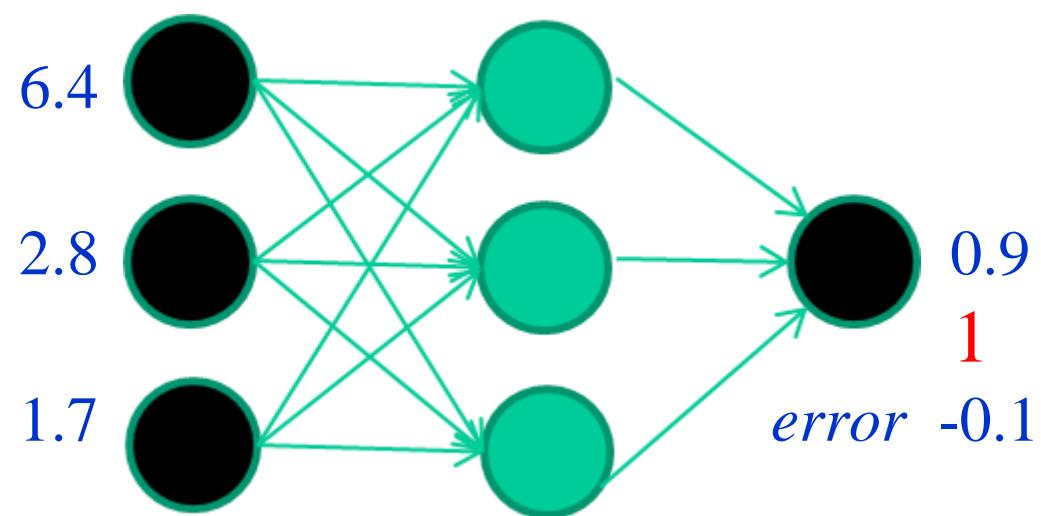
3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Compare with target output



Training data

Fields *class*

1.4 2.7 1.9 0

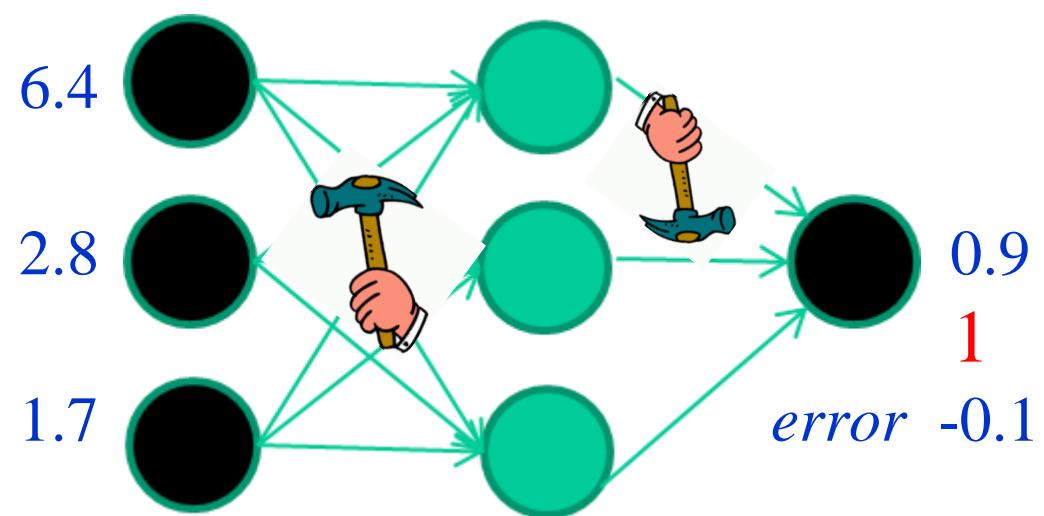
3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Adjust weights based on error



Training data

Fields *class*

1.4 2.7 1.9 0

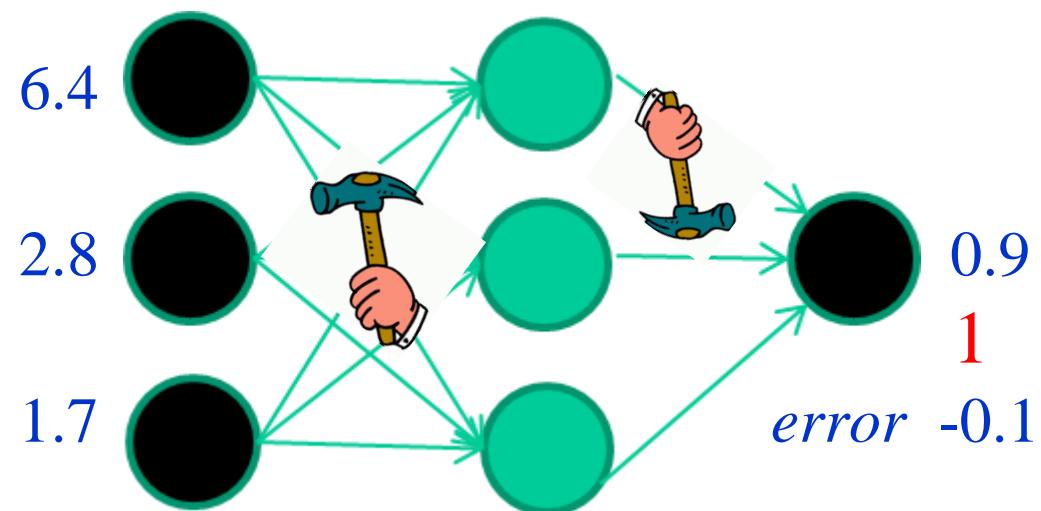
3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

And so on

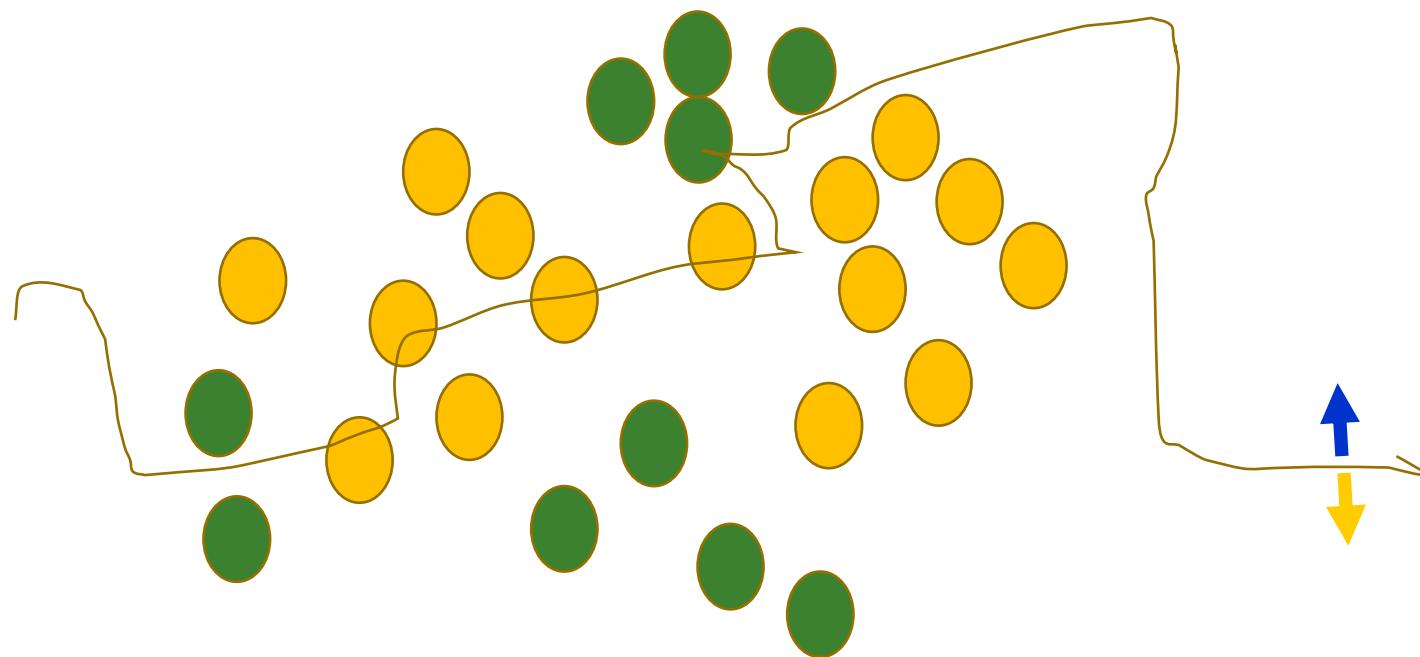


Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments.

Algorithms for weight adjustment are designed to make changes that will reduce the error.

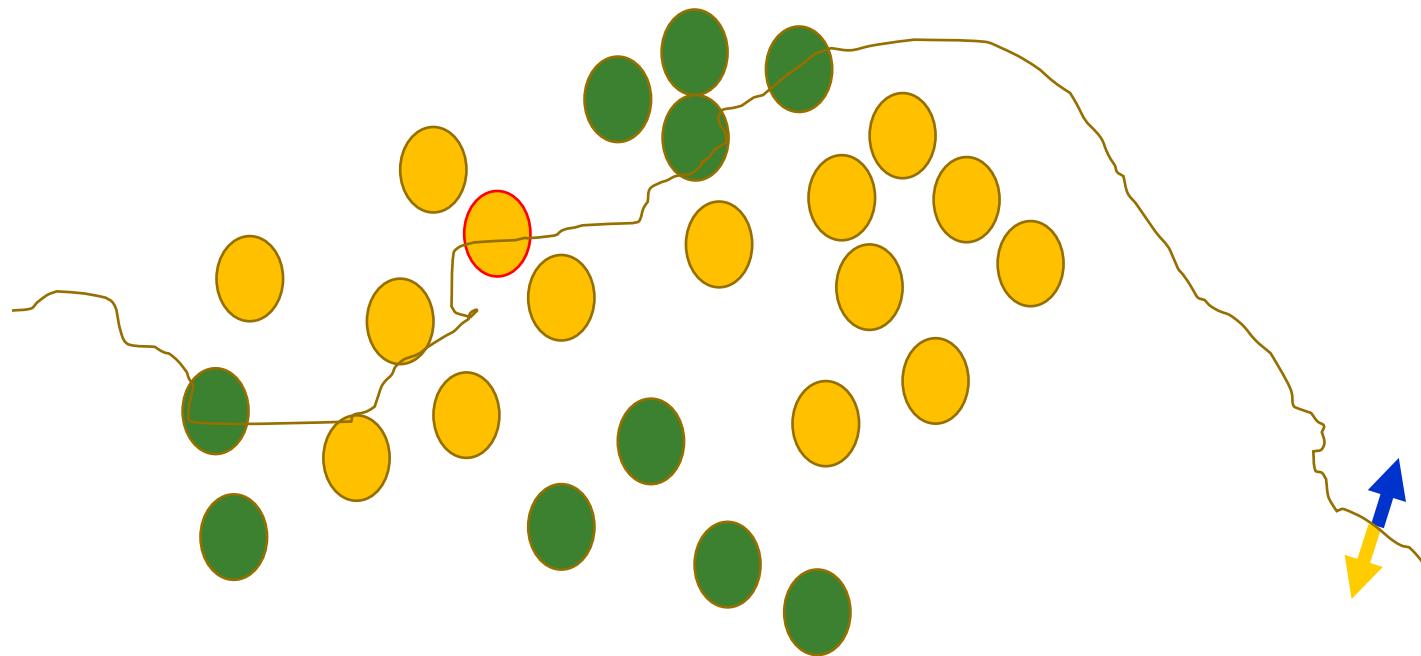
The decision boundary perspective...

Initial random weights



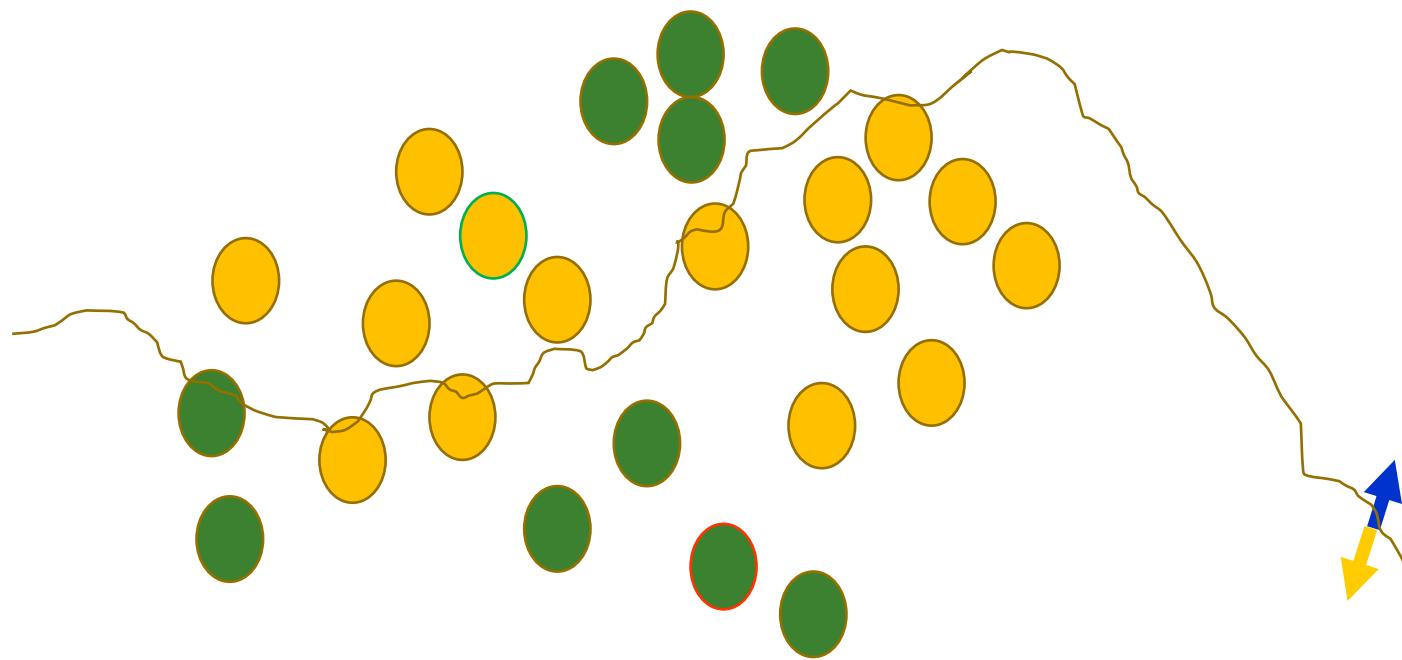
The decision boundary perspective...

Present a training instance / adjust the weights



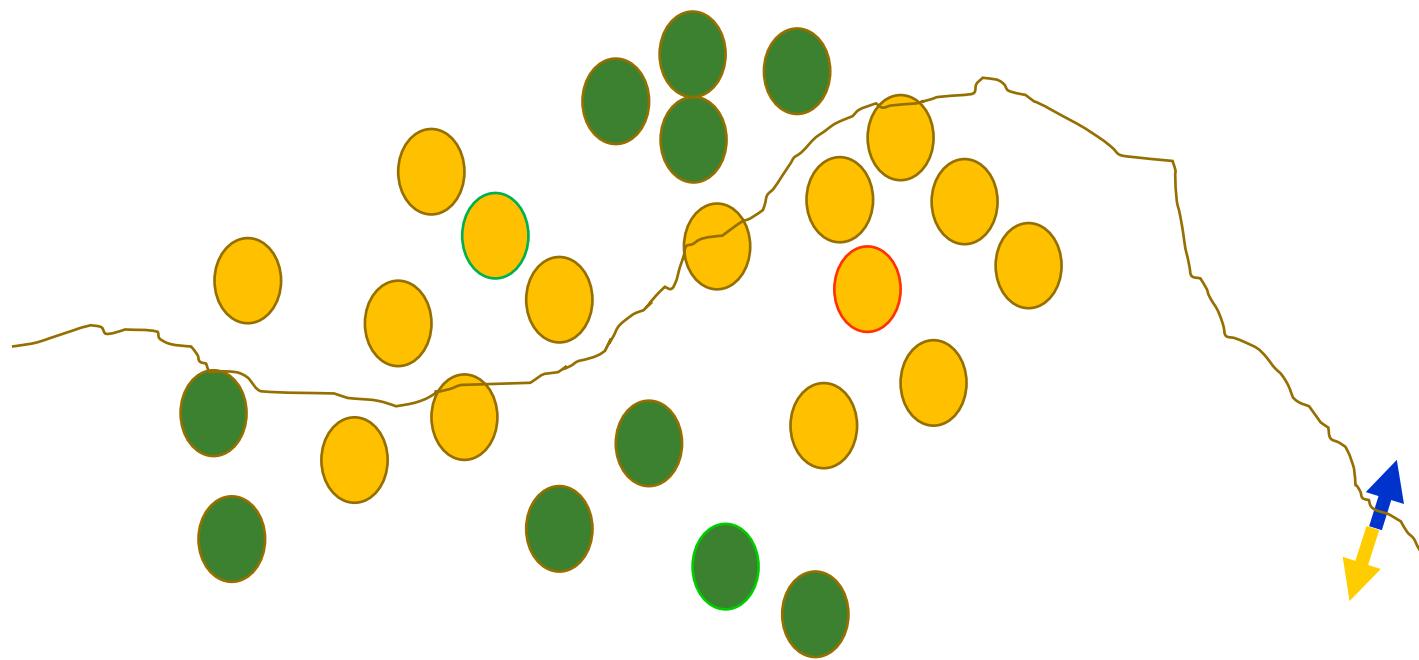
The decision boundary perspective...

Present a training instance / adjust the weights



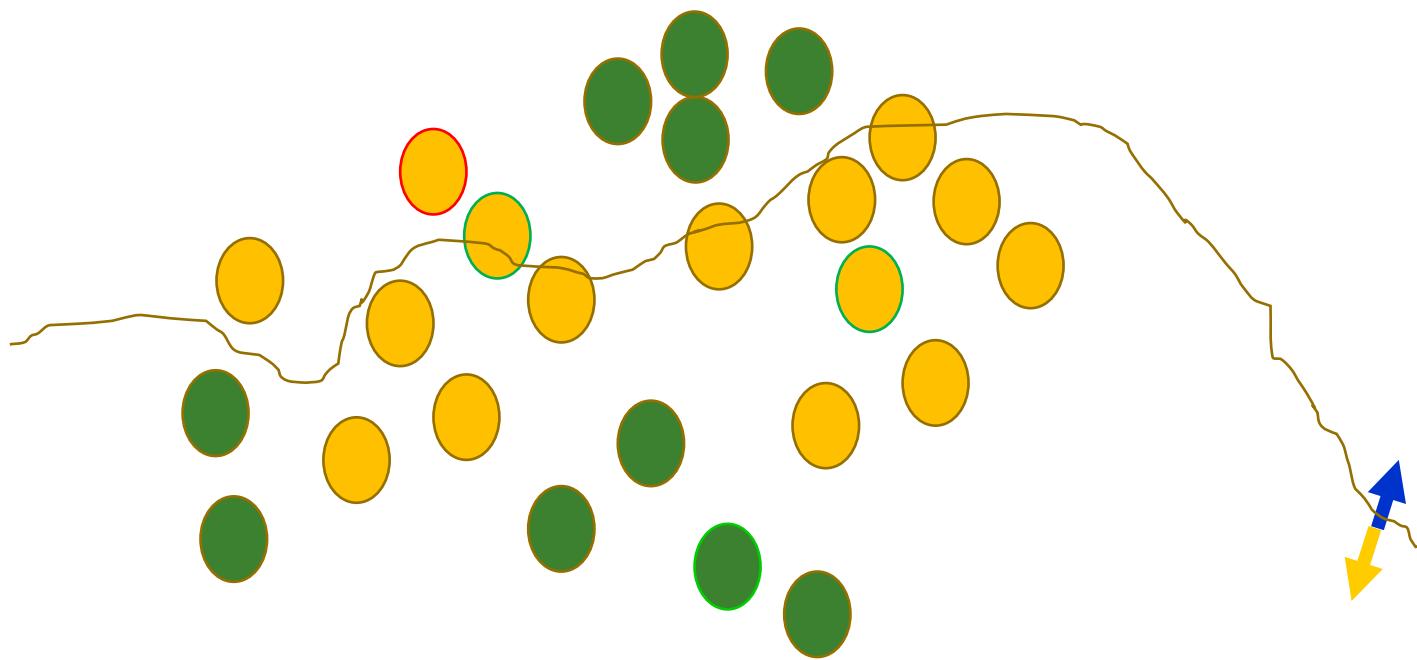
The decision boundary perspective...

Present a training instance / adjust the weights



The decision boundary perspective...

Present a training instance / adjust the weights



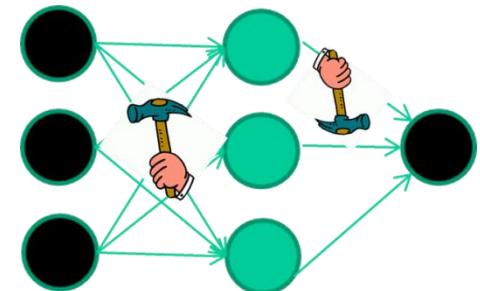
The decision boundary perspective...

Eventually



Point

- Weight-learning algorithms for NNs are dumb
- They work by making thousands and thousands of tiny adjustments, each making the network do better at the most recent pattern, but perhaps a little worse on many others
- But, by dumb luck, eventually this tends to be good enough to learn effective classifiers for many real applications



Some other points

If $f(x)$ is non-linear, a network with 1 hidden layer can, in theory, learn perfectly any classification problem. A set of weights exists that can produce the targets from the inputs.

The problem is finding them.

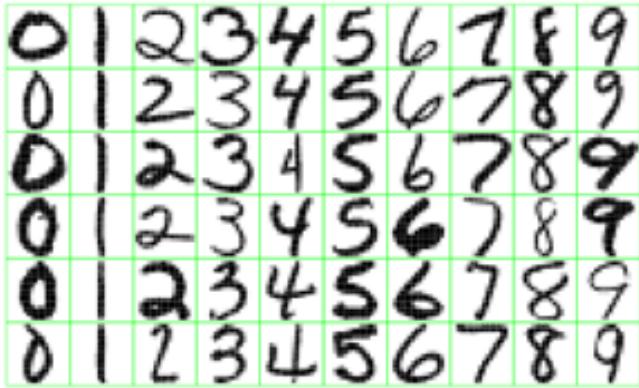
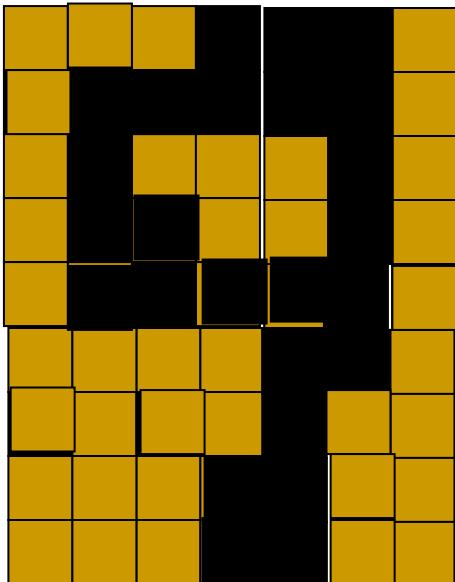
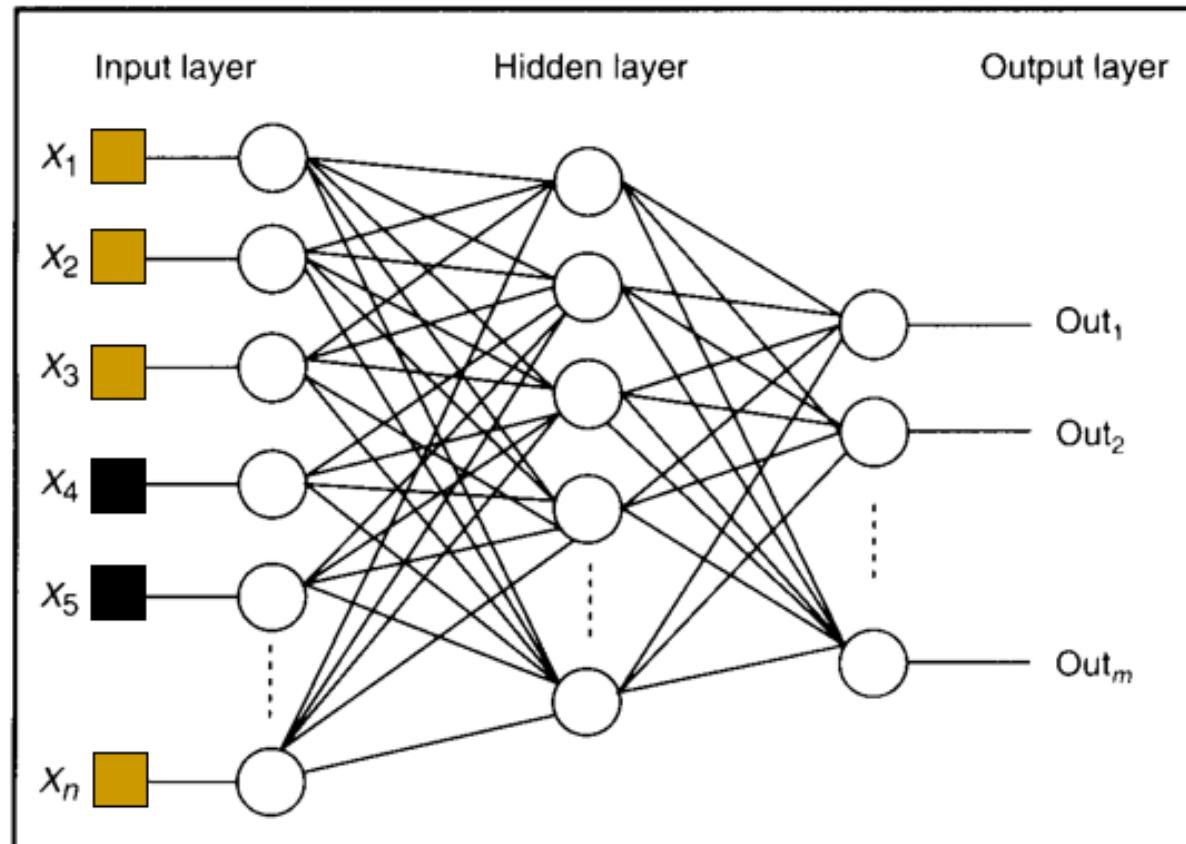


Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.



Feature detectors



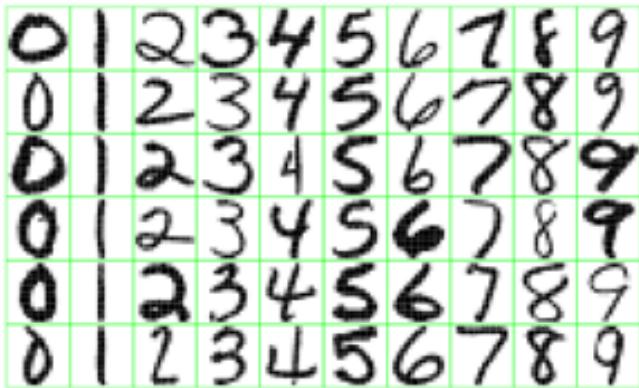
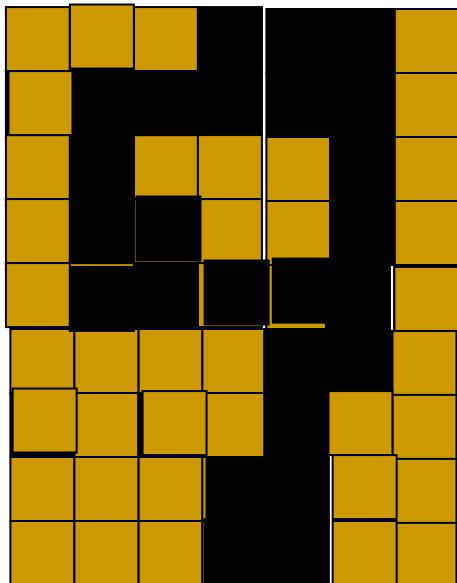
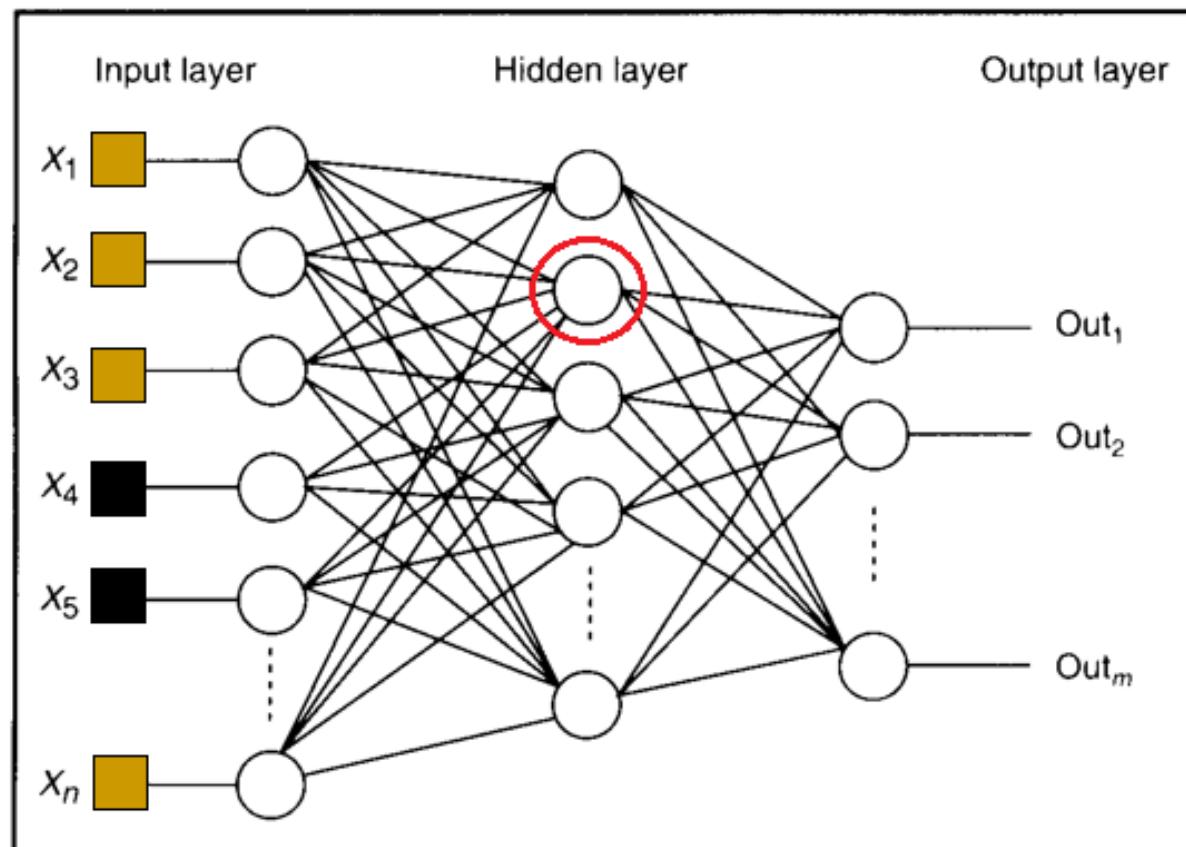


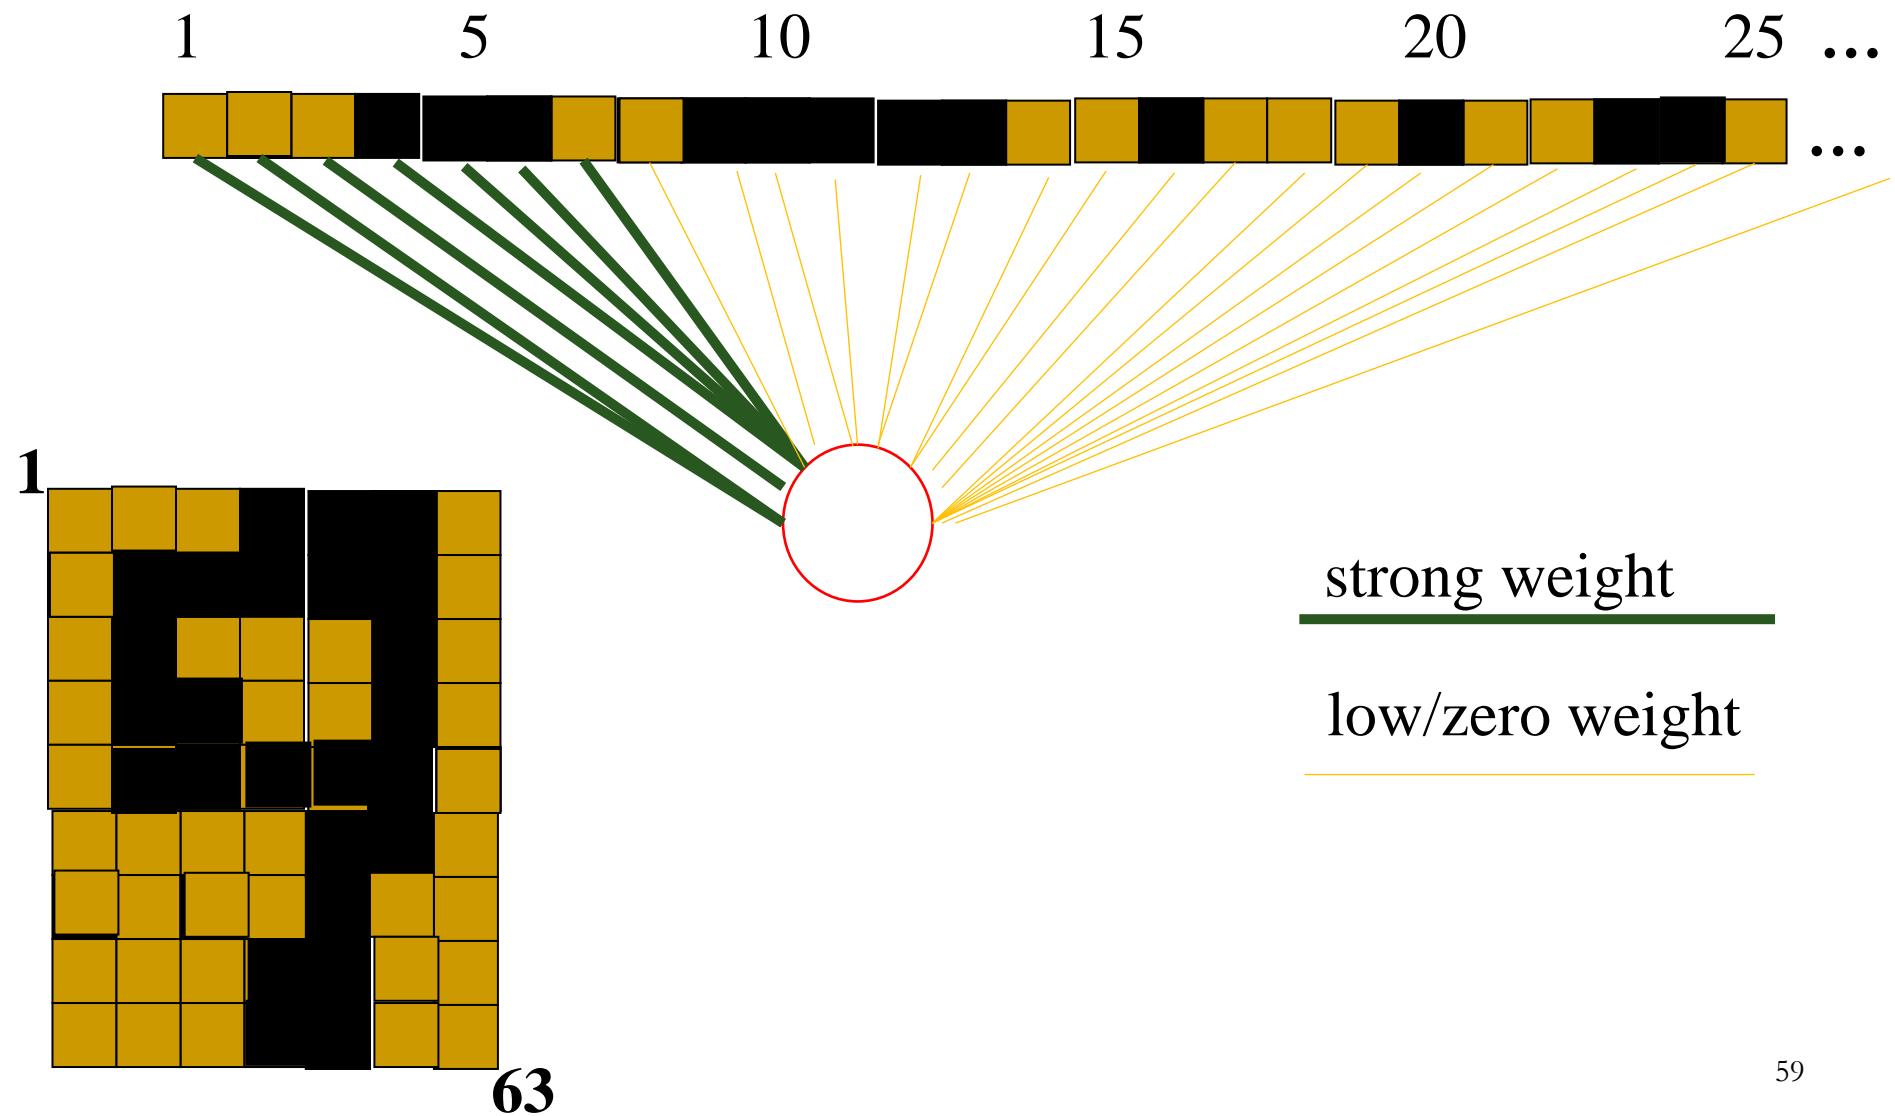
Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.



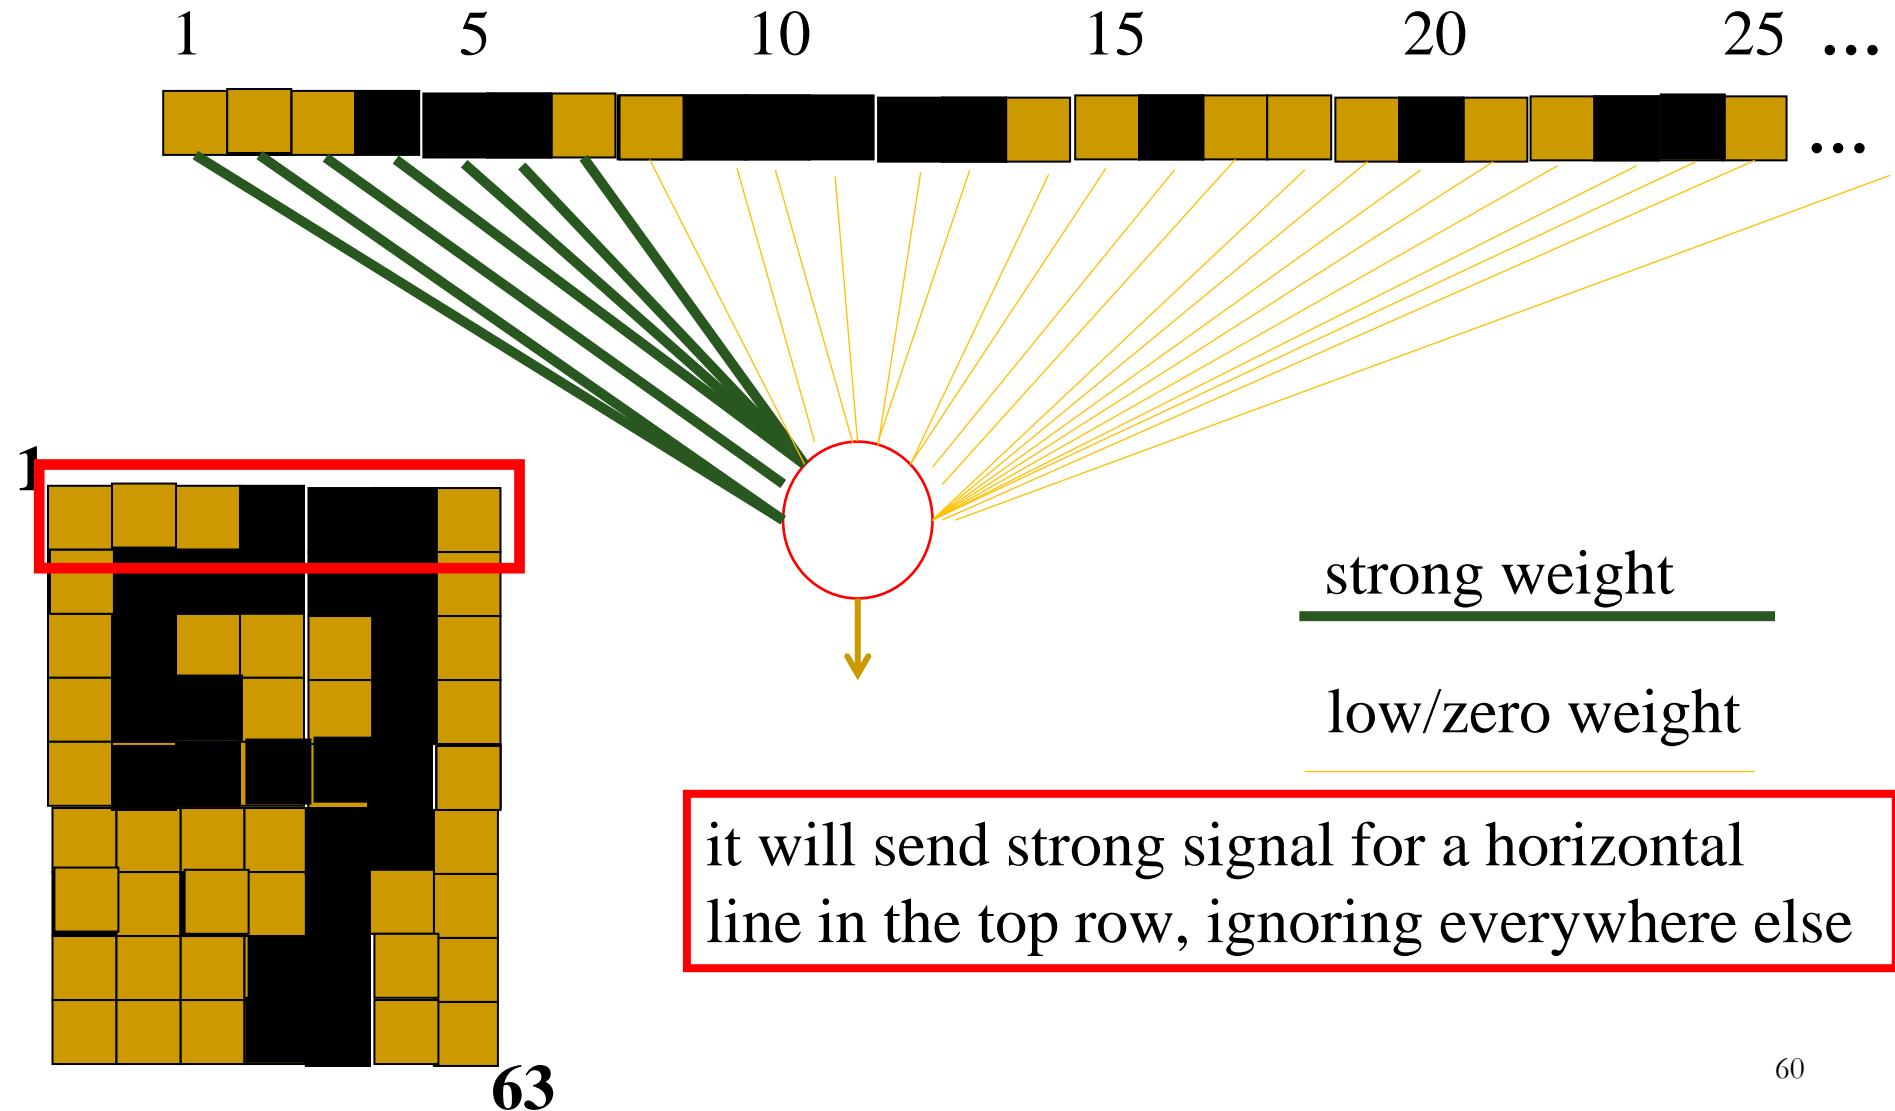
What is this unit doing?



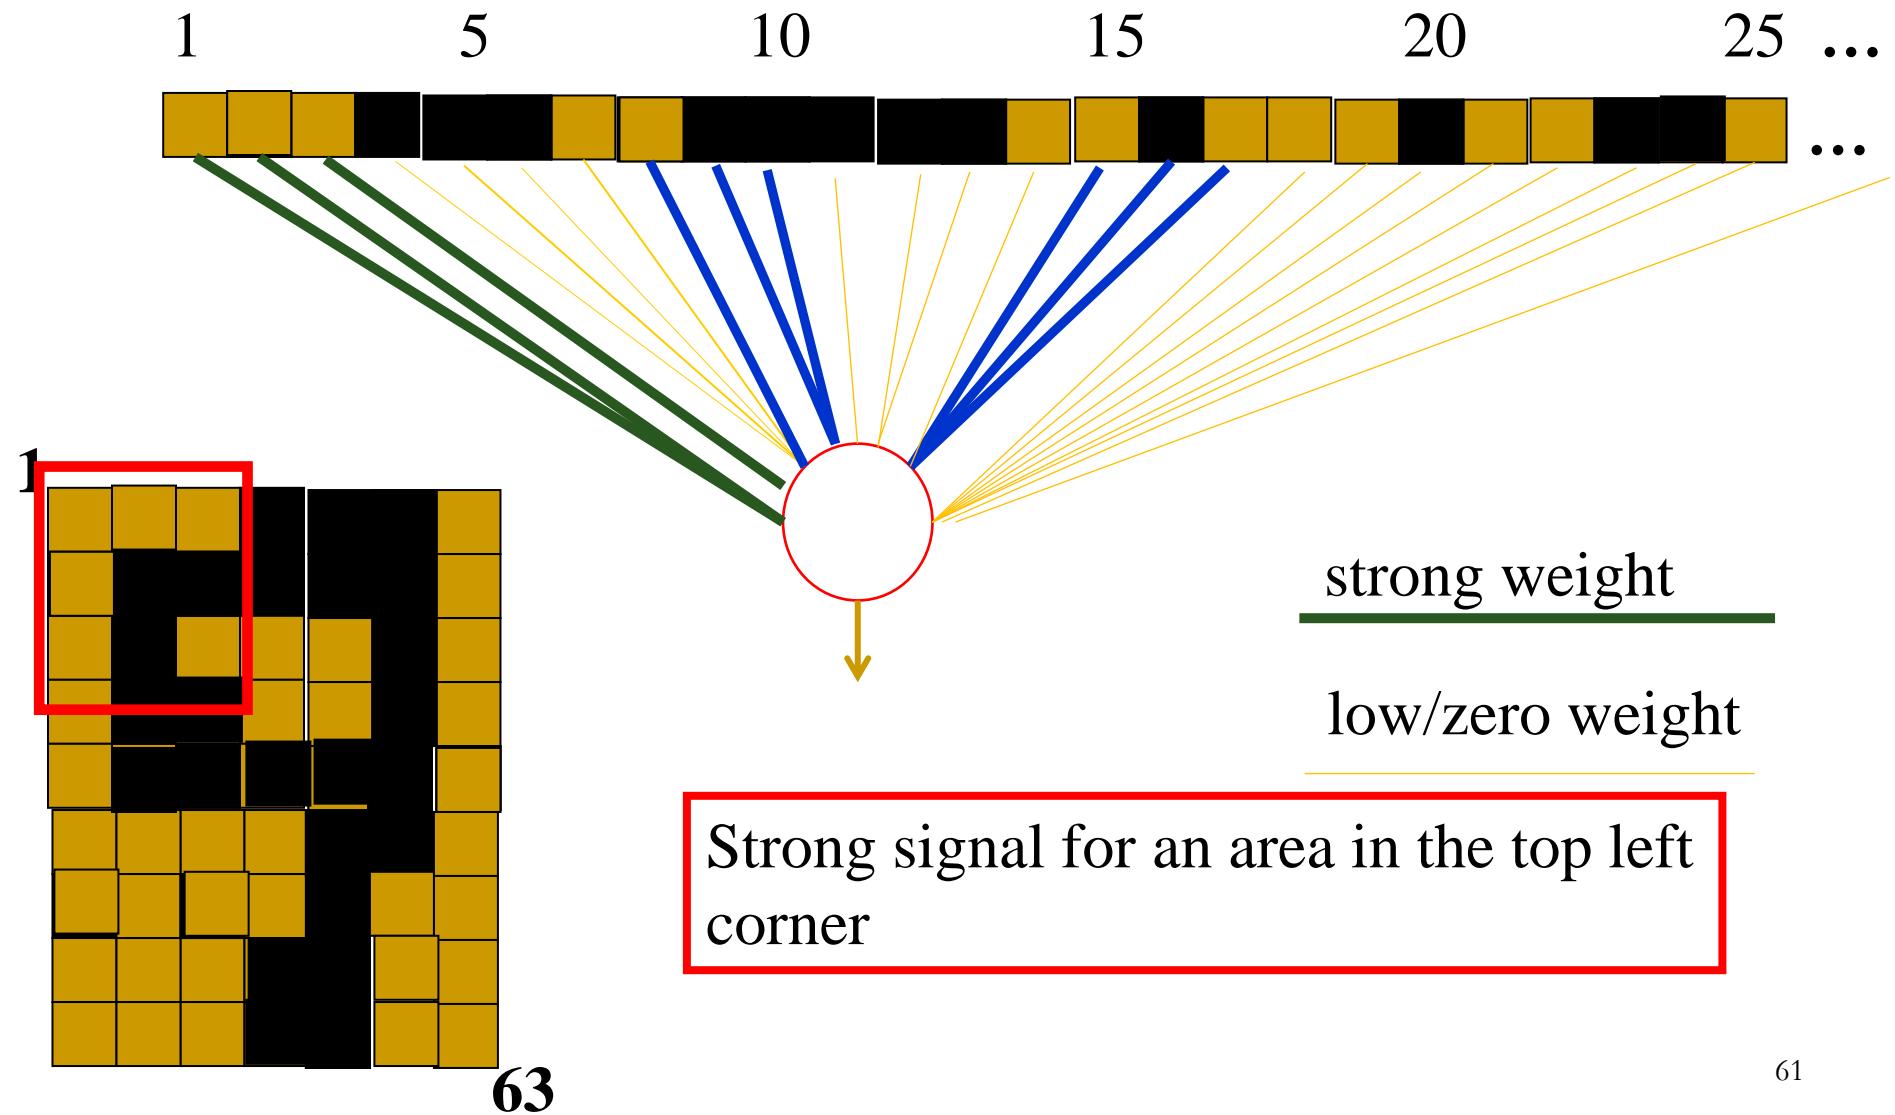
Hidden layer units become self-organised feature detectors



What does this unit detect?



What does this unit detect?



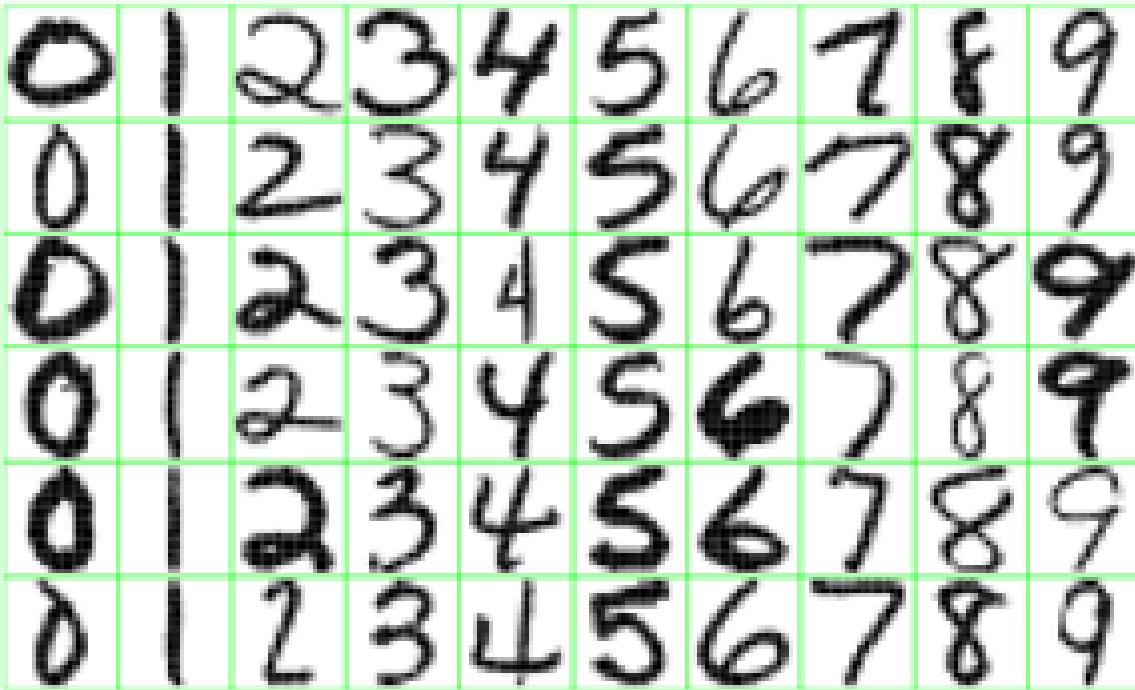


Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

What features might you expect a good NN to learn, when trained with data like this?

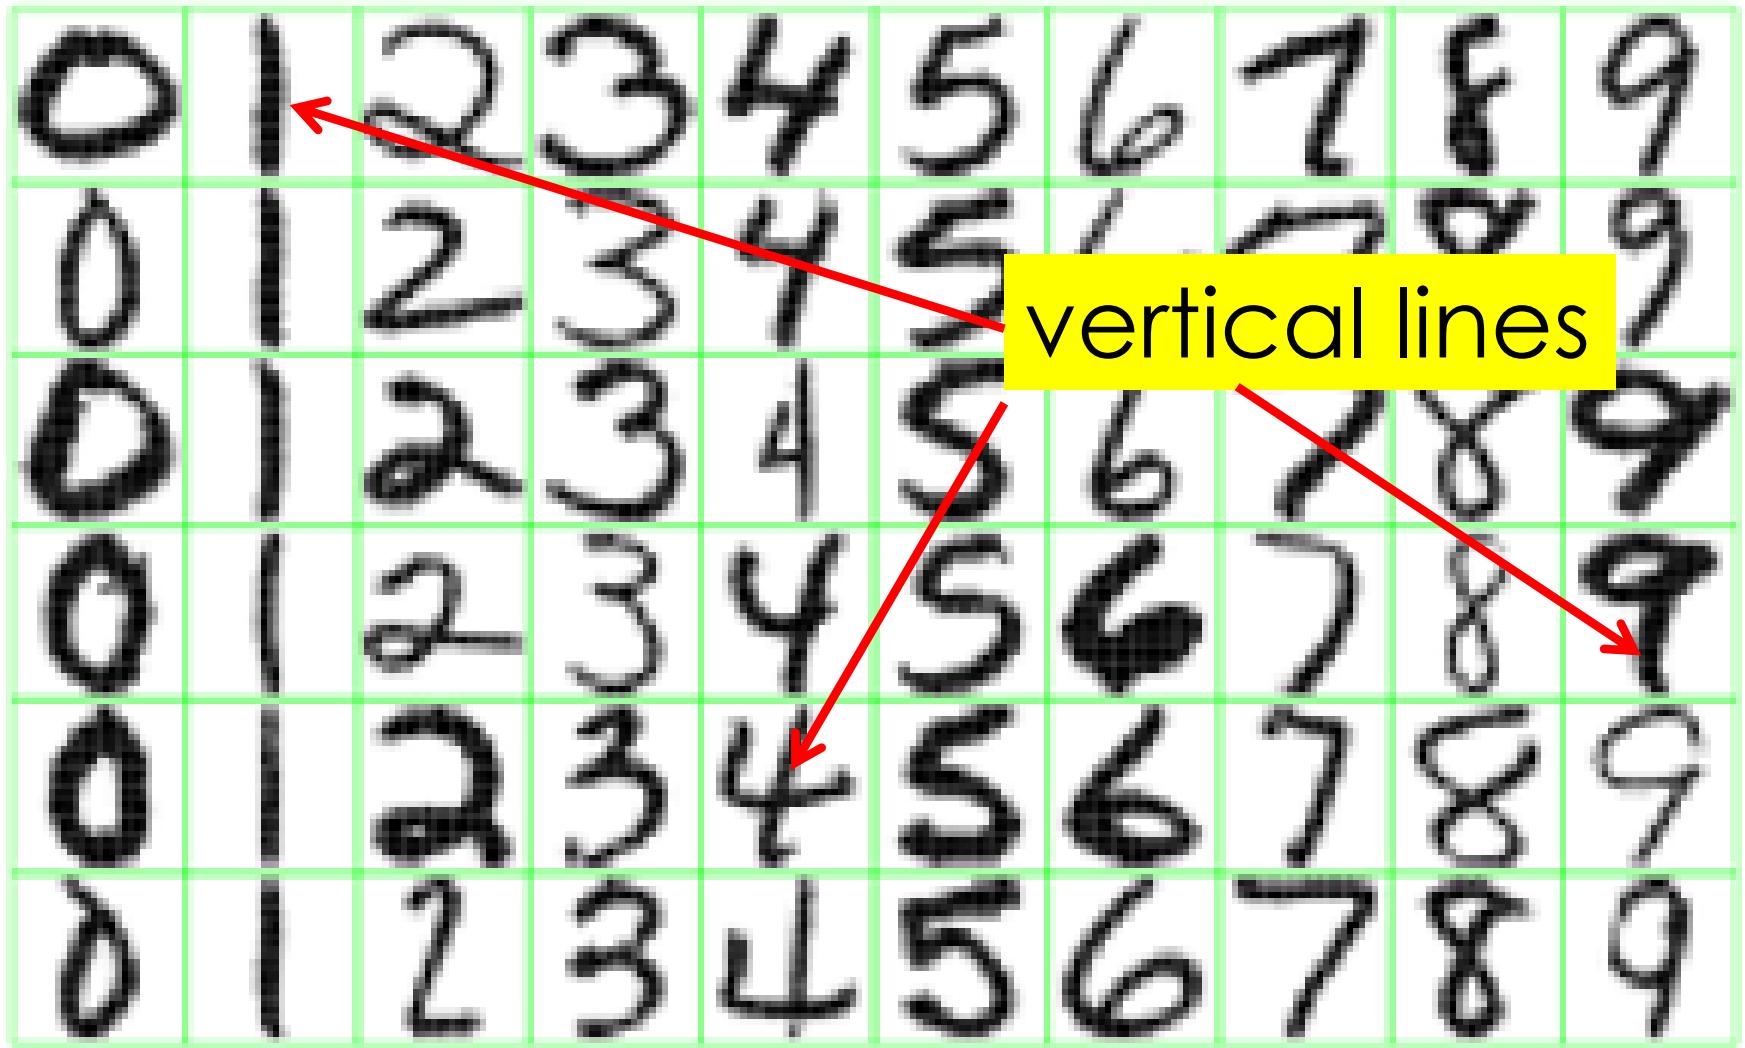


Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.

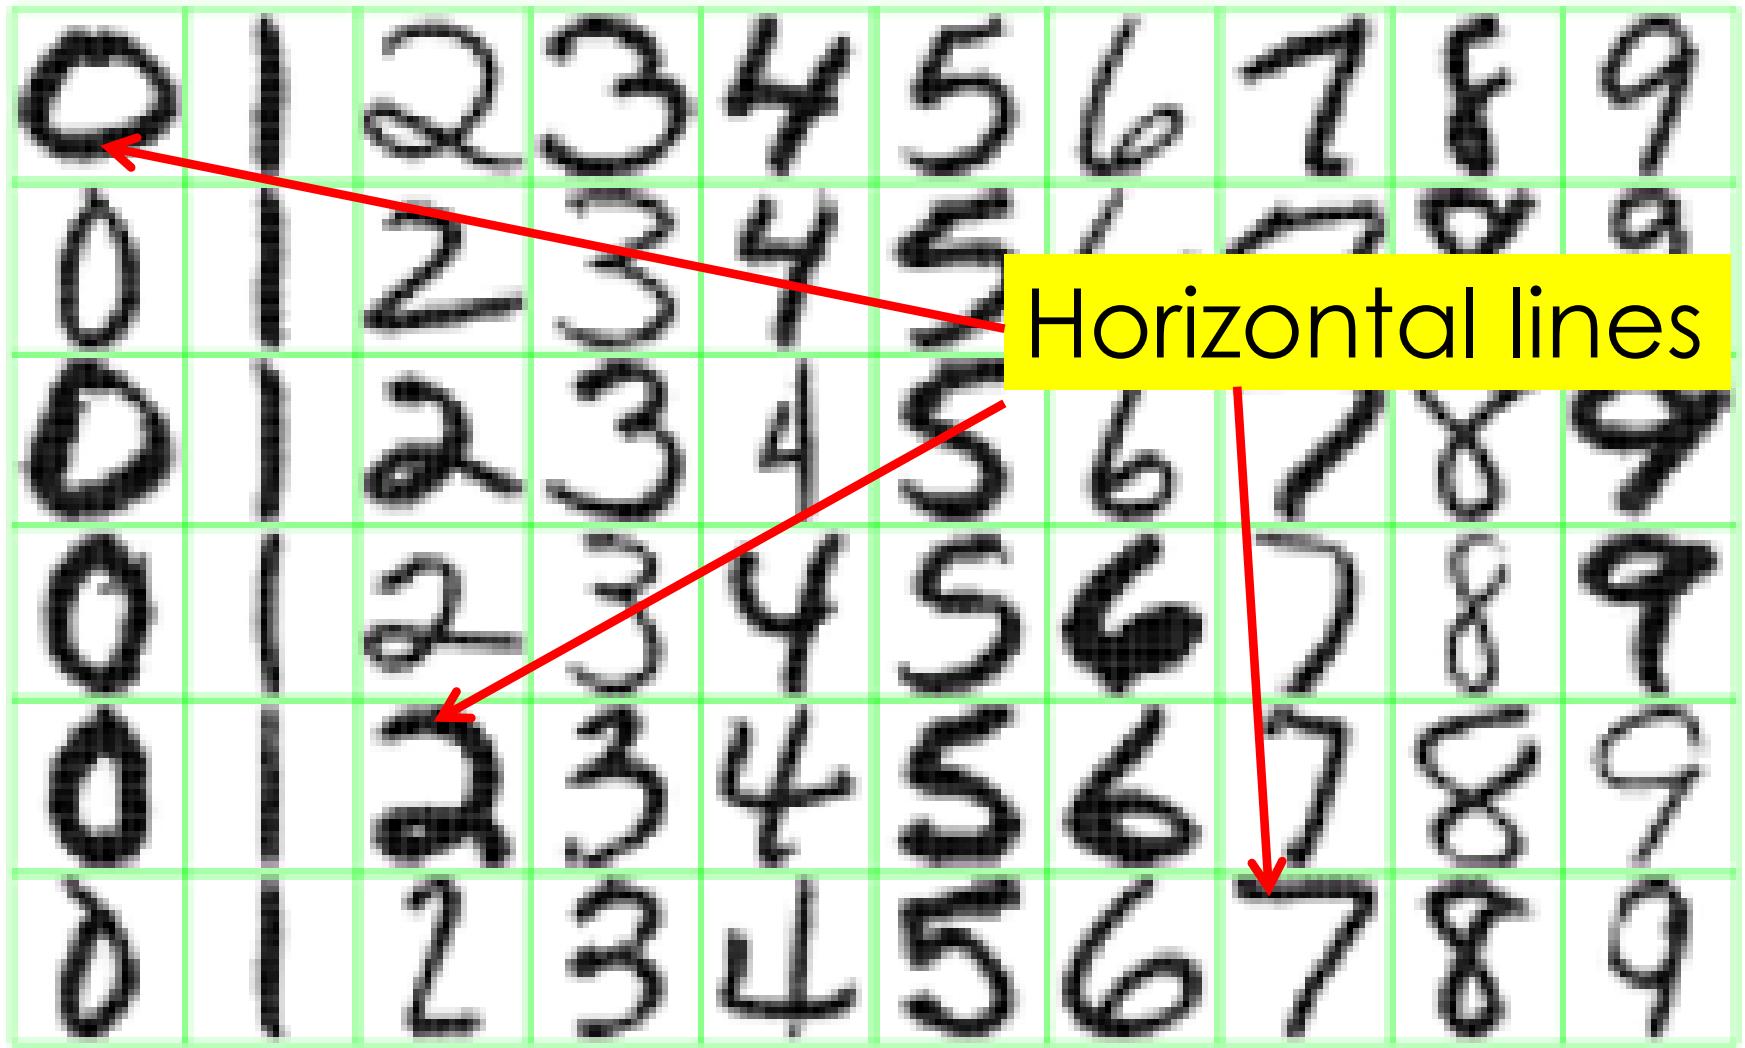


Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.

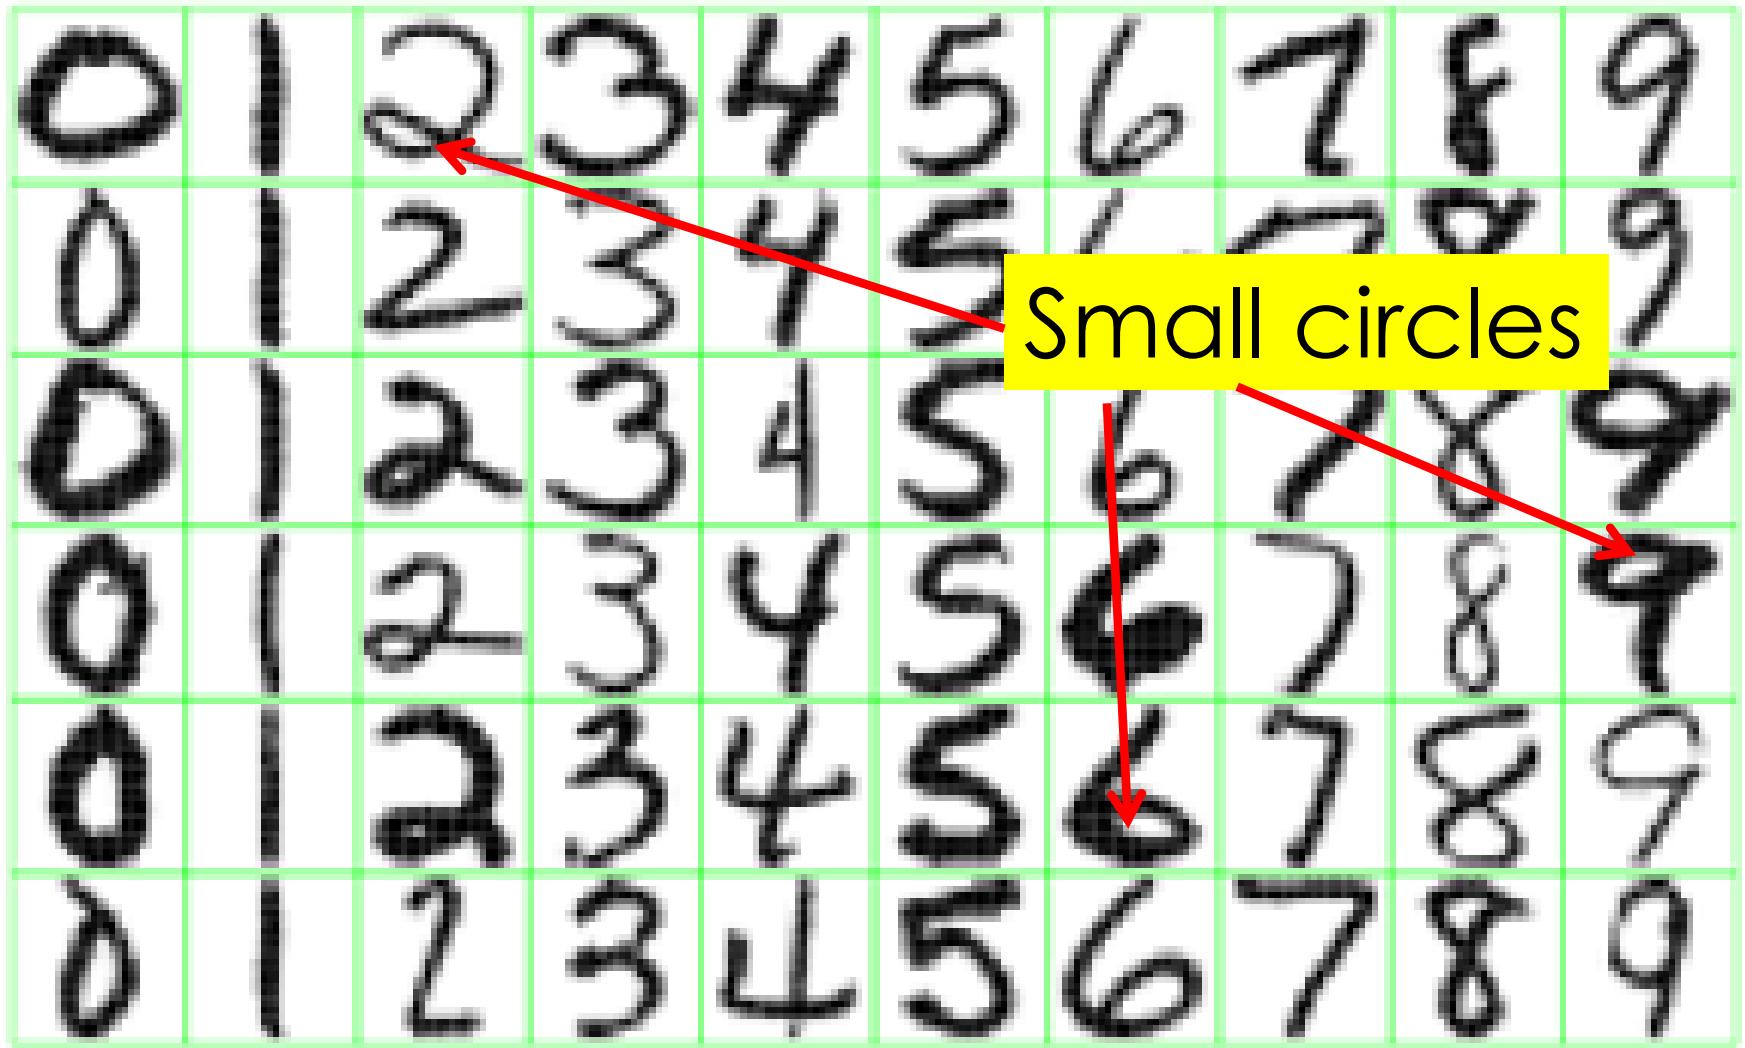
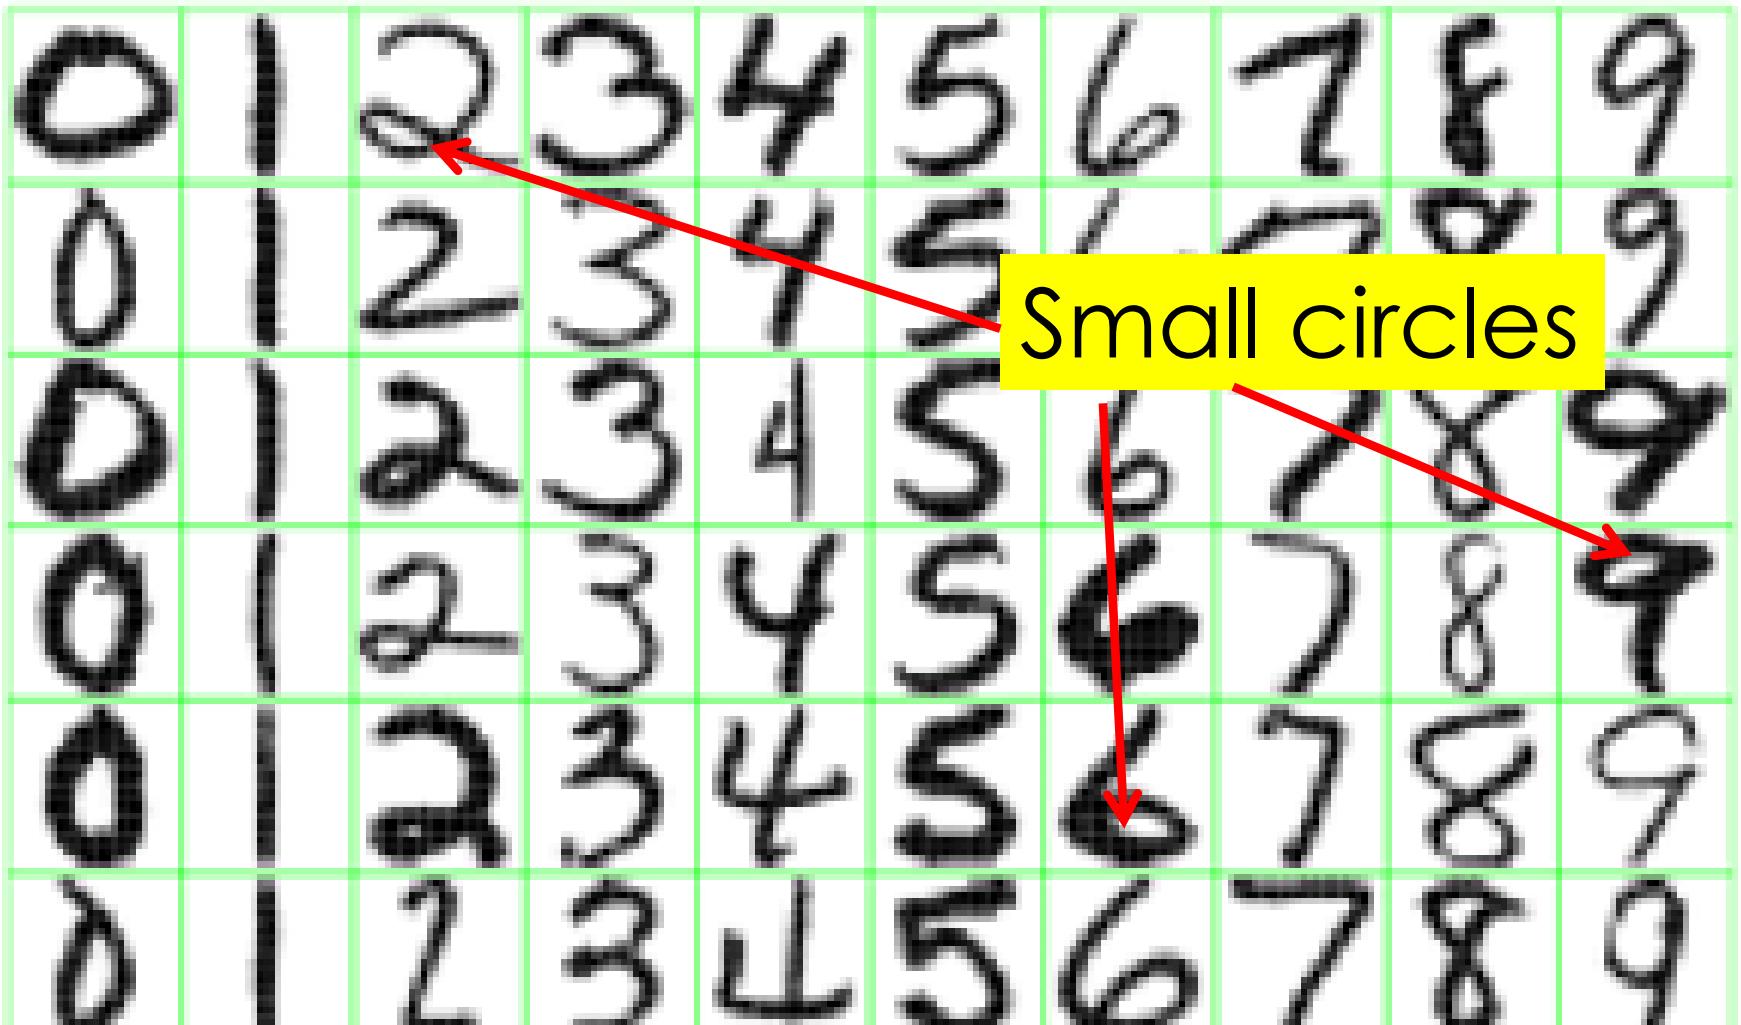


Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.

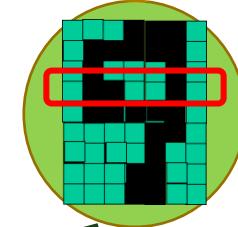
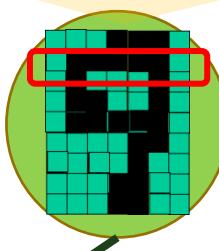
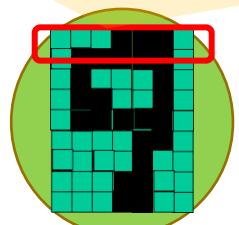


But what about position invariance ???
our example unit detectors were tied to
specific parts of the image

successive layers can learn higher-level features ...

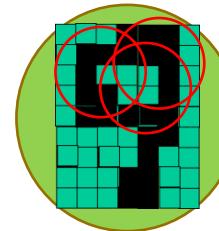
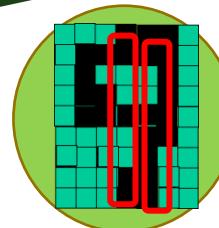
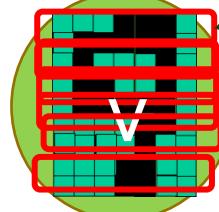


detect lines in
Specific positions



etc ...

Higher level detectors
("horizontal line",
"vertical line"
"upper loop", etc...)

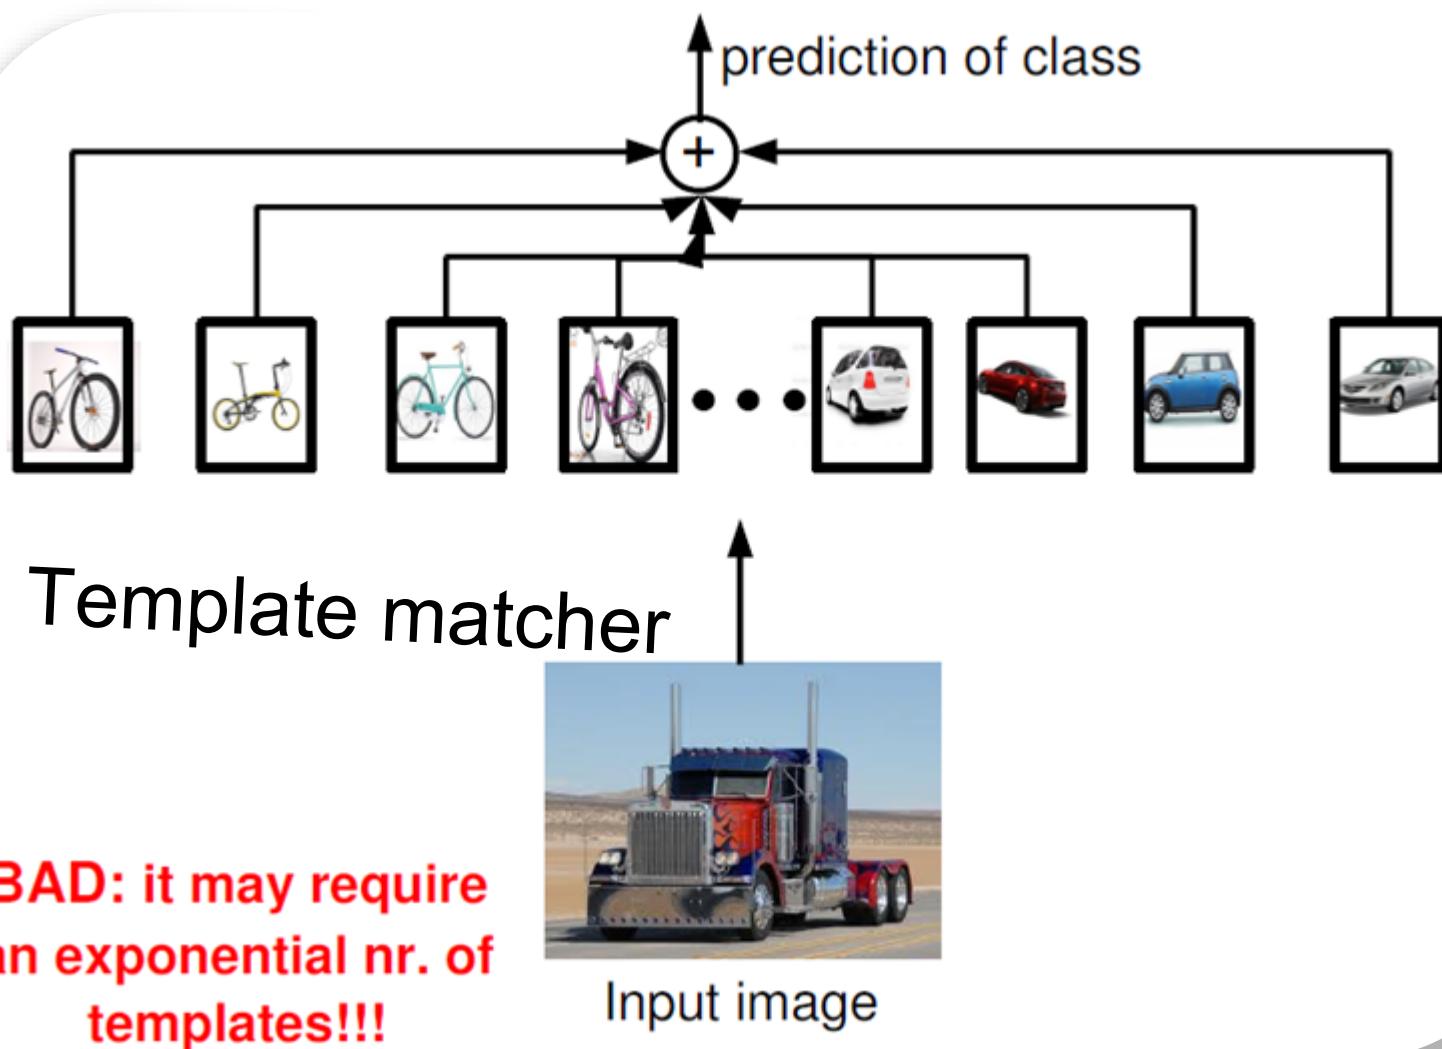


etc ...

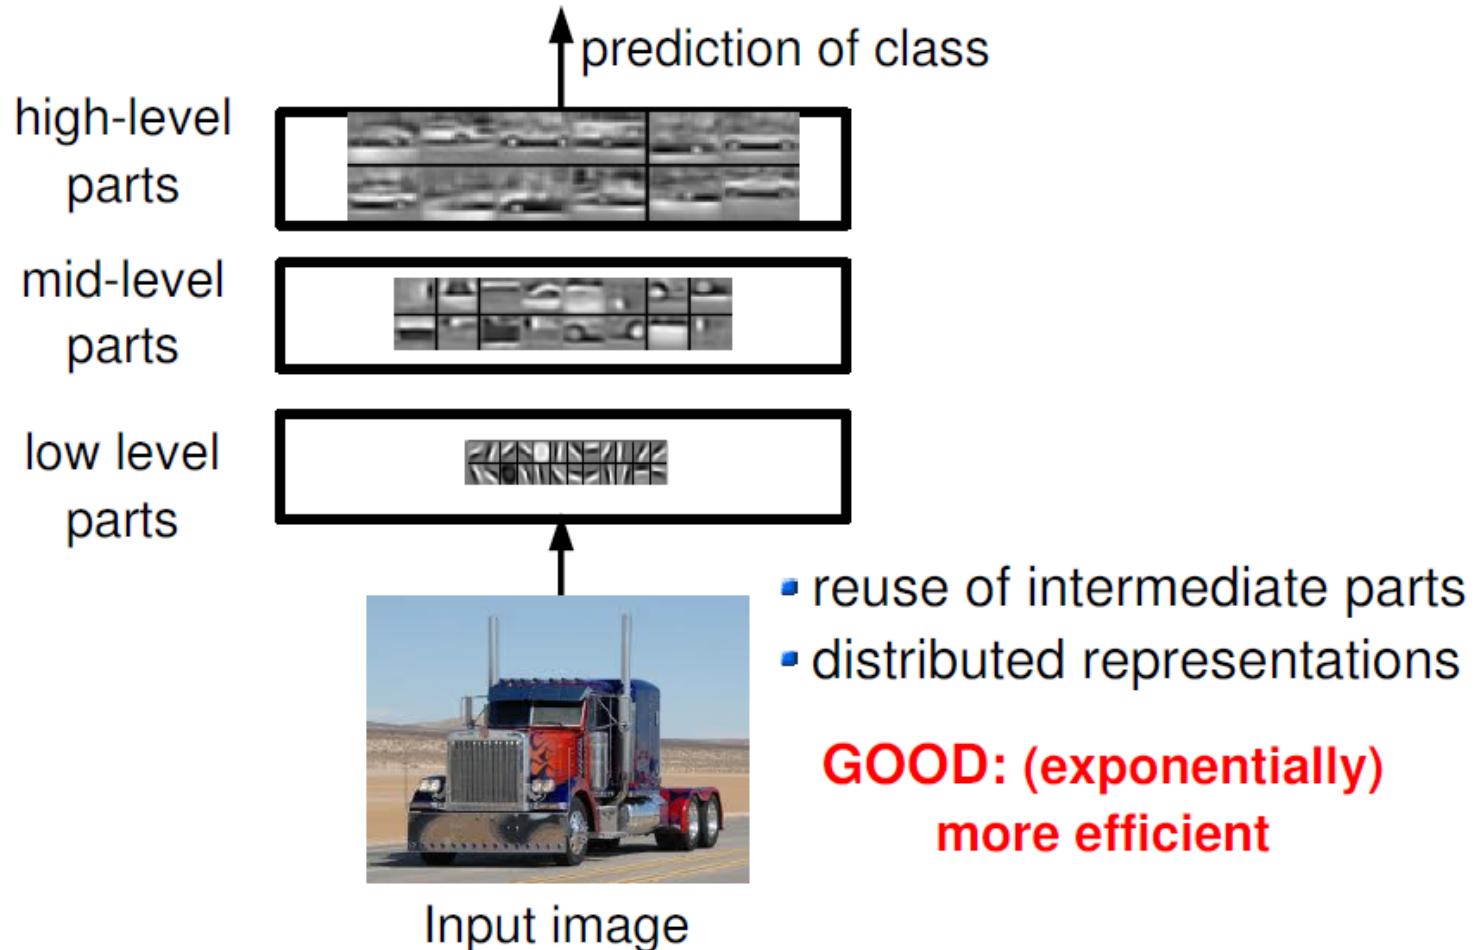


LEARNING

RECALL: the traditional approach



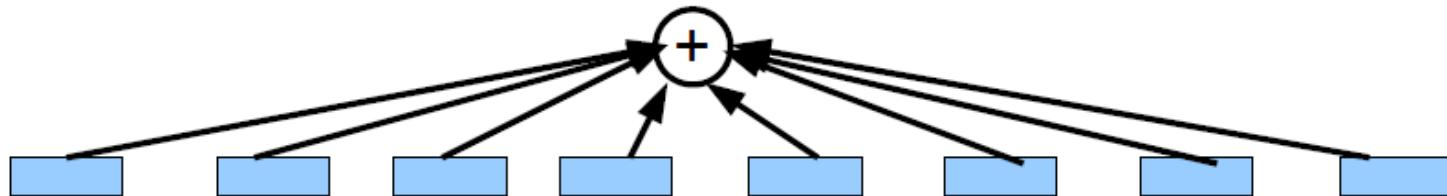
Composition



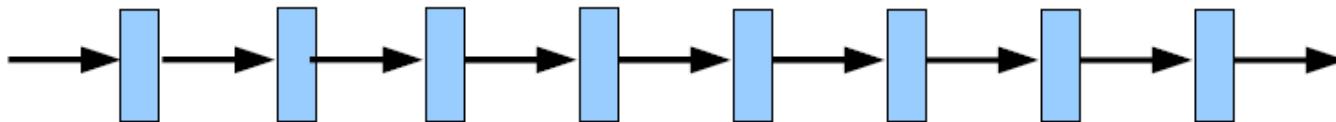
Learning

Given a dictionary of simple non-linear functions: g_1, \dots, g_n

Proposal #1: linear combination $f(x) \approx \sum_j g_j$



Proposal #2: composition $f(x) \approx g_1(g_2(\dots g_n(x)\dots))$



Learning

Given a dictionary of simple non-linear functions: g_1, \dots, g_n

Proposal #1: linear combination

$$f(x) \approx \sum_j g_j$$

- Kernel learning
- Boosting
- ...

Shallow

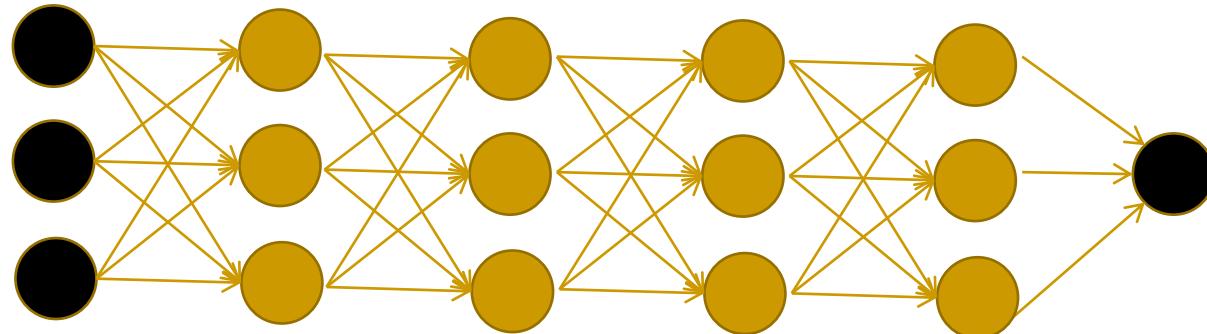
Proposal #2: composition

$$f(x) \approx g_1(g_2(\dots g_n(x)\dots))$$

- Deep learning
- Scattering networks (wavelet cascade)
- S.C. Zhou & D. Mumford “grammar”

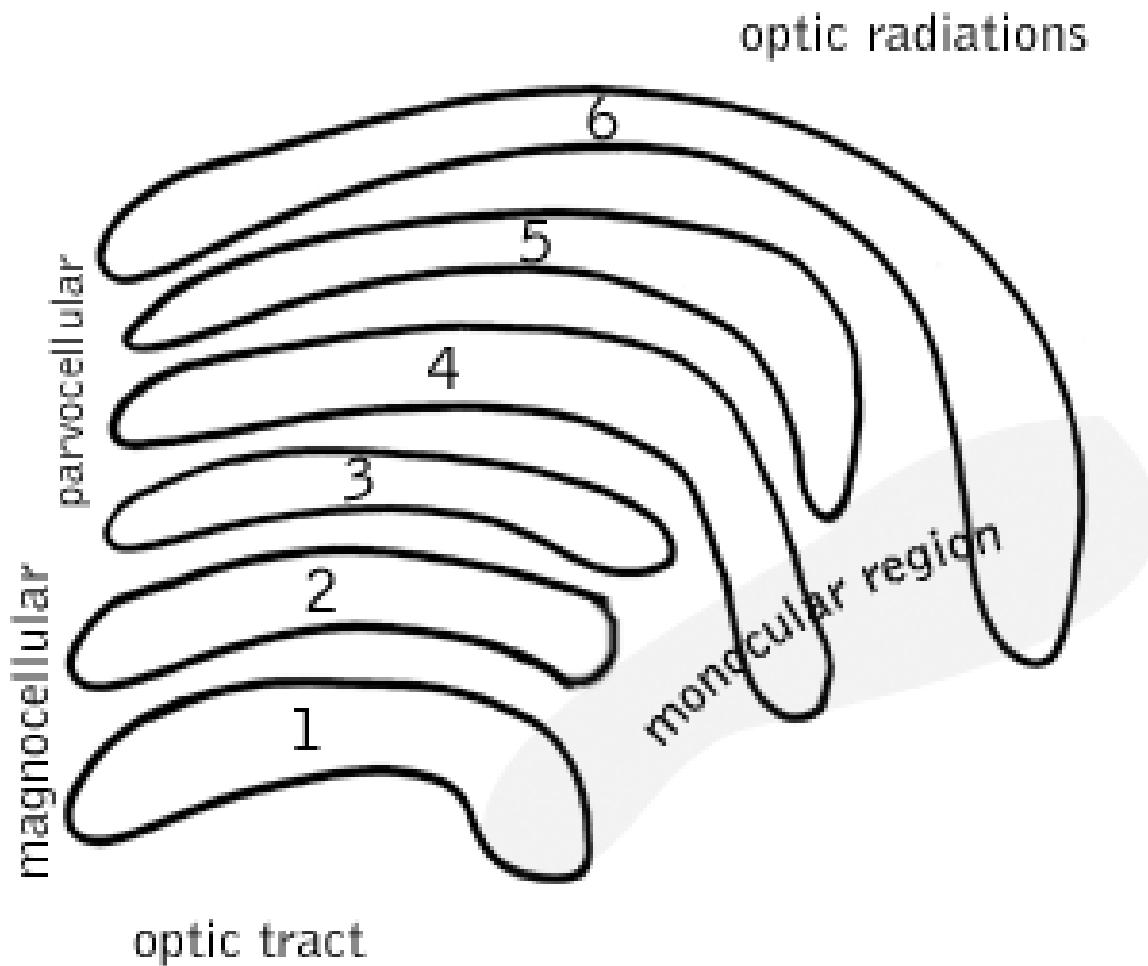
Deep

-> multiple layers



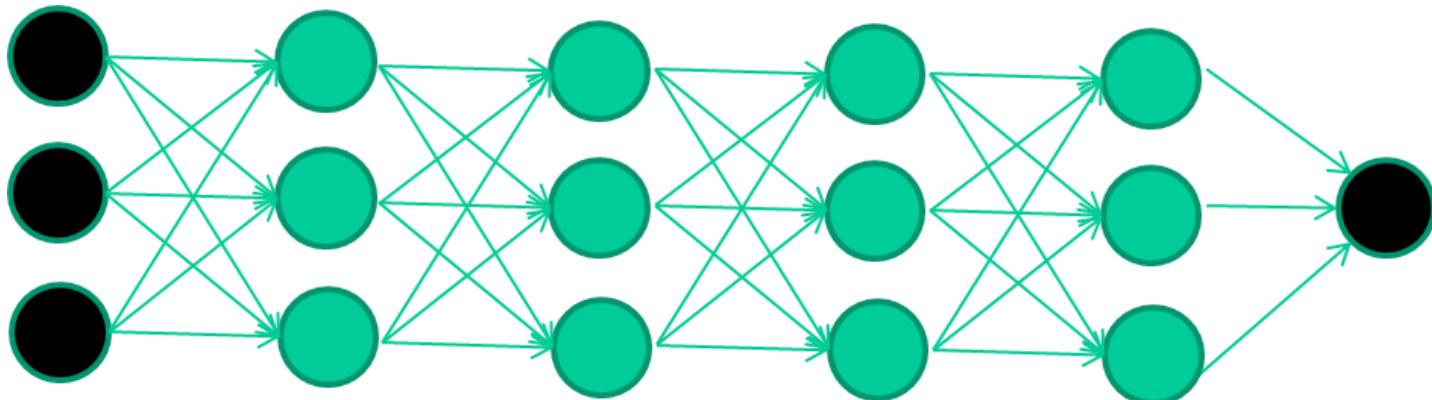
-> multiple layers

Your brain works that way

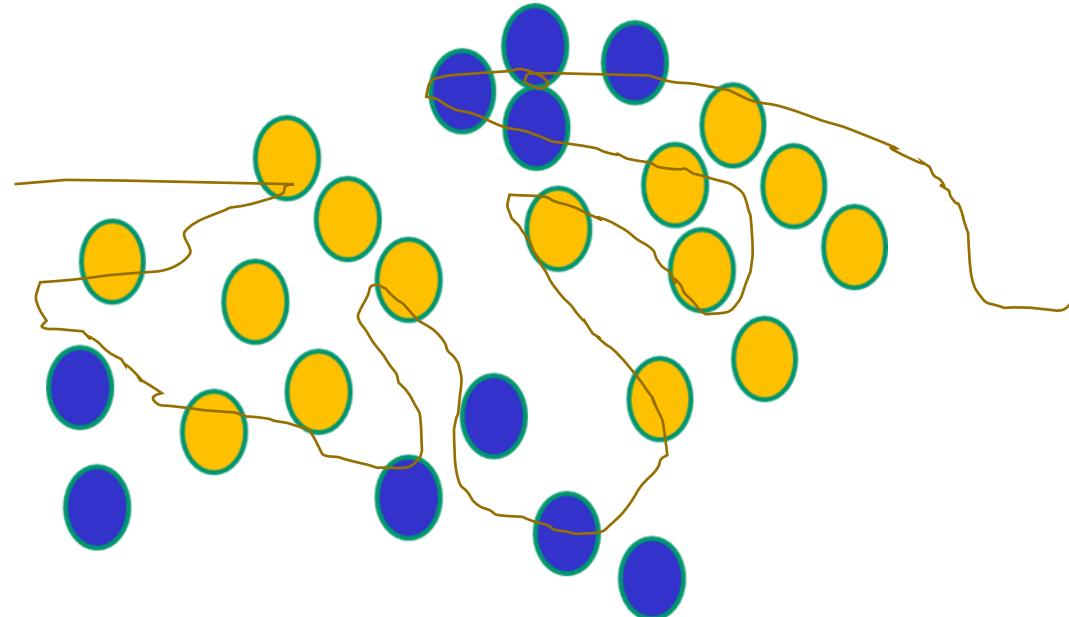
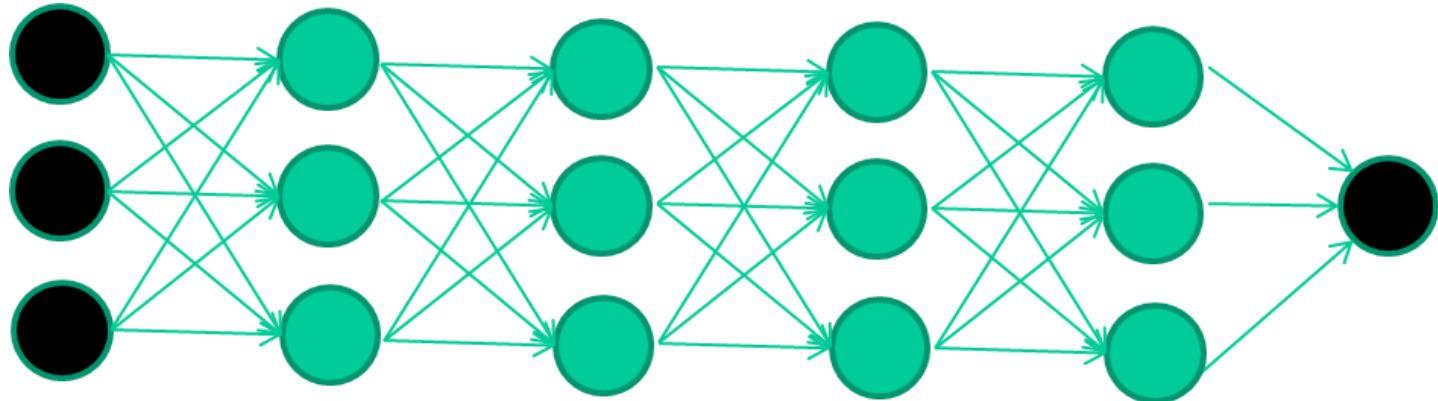


So: multiple layers make sense

Many-layer neural network architectures should be capable of learning the true underlying features and ‘feature logic’, and therefore generalise very well ...

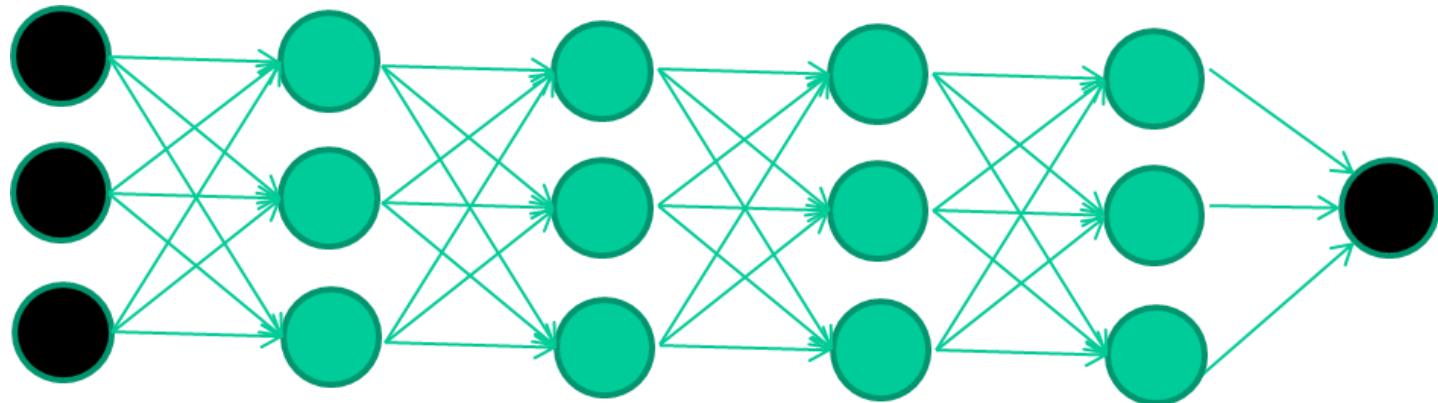


But, until very recently, our weight-learning algorithms simply did not work on multi-layer architectures

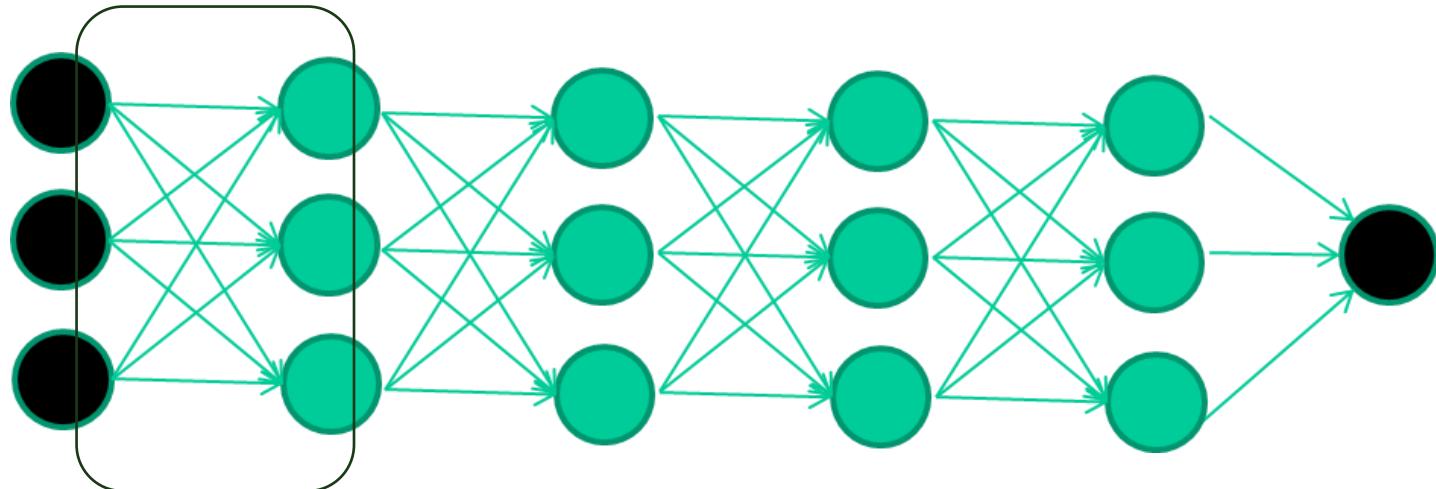




A new way to train multi-layer NNs...

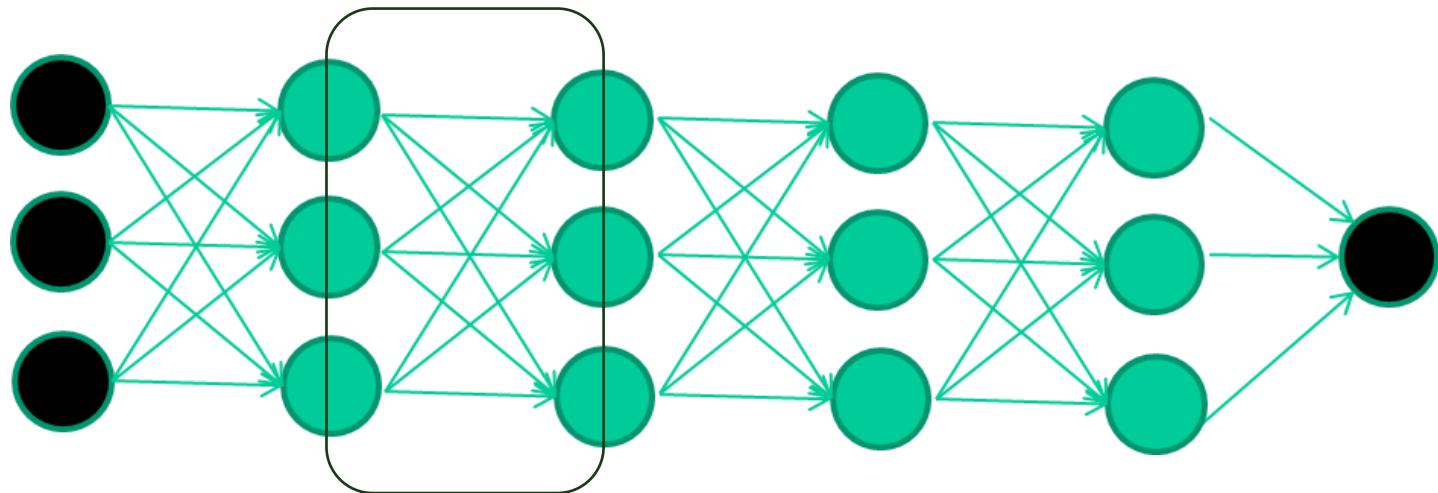


The new way to train multi-layer NNs...



Train **this** layer first

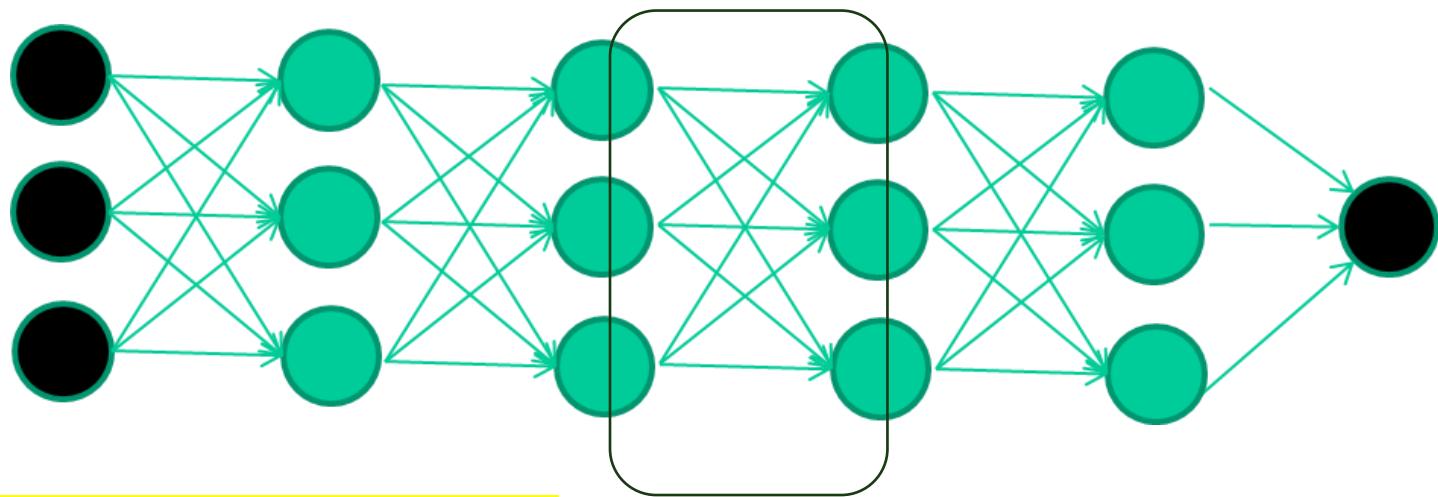
The new way to train multi-layer NNs...



Train **this** layer first

then **this** layer

The new way to train multi-layer NNs...

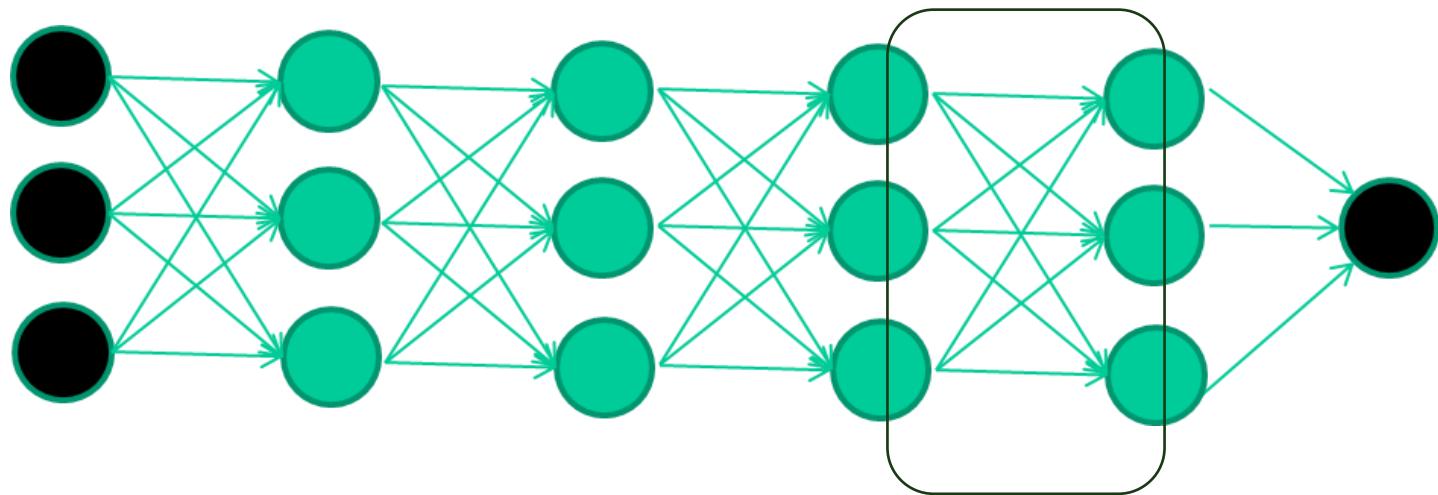


Train **this** layer first

then **this** layer

then **this** layer

The new way to train multi-layer NNs...



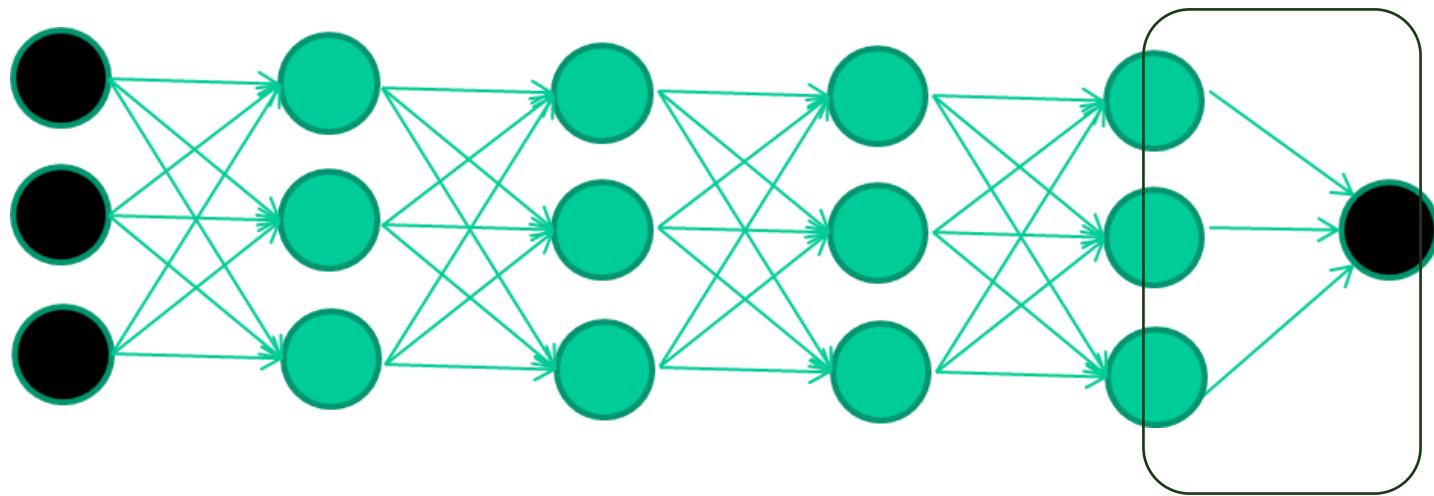
Train **this** layer first

then **this** layer

then **this** layer

then **this** layer

The new way to train multi-layer NNs...



Train **this** layer first

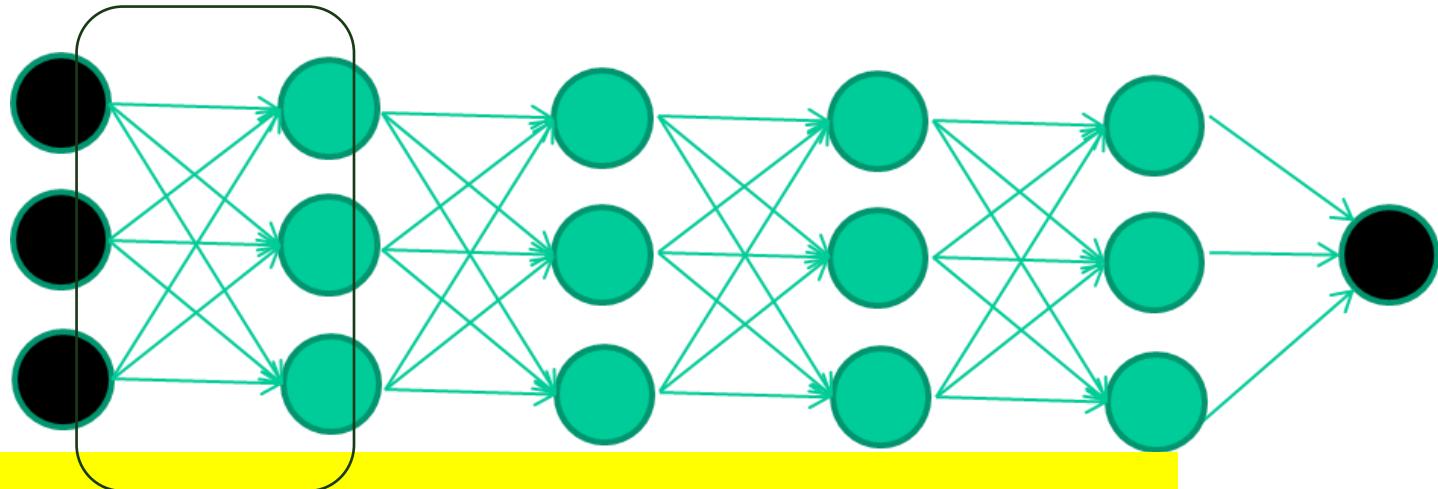
then **this** layer

then **this** layer

then **this** layer

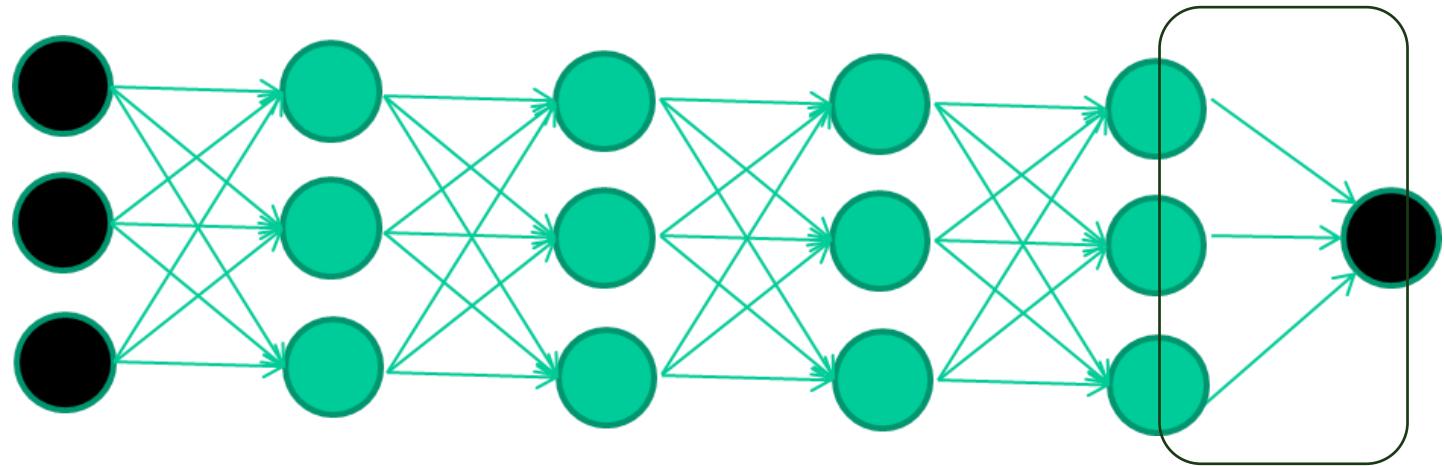
finally **this** layer

The new way to train multi-layer NNs...



*EACH of the (non-output) layers is trained
to be an **auto-encoder***

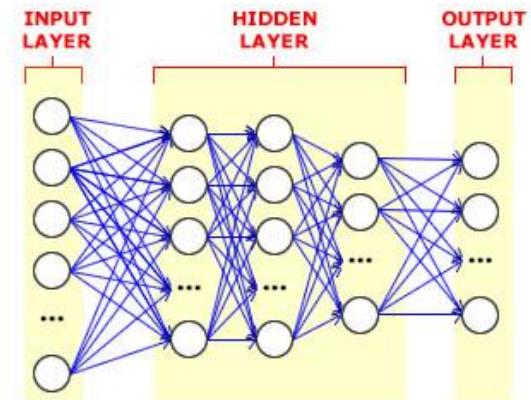
*Basically, it is forced to learn good features that
describe what comes from the previous layer*



Final layer trained to predict class based on outputs from previous layers

Basic Idea

- There are many, many types of training (learning) a deep learning network.
- Different kinds of autoencoder, variations on architectures and training algorithms, etc...
- Very fast growing area ...

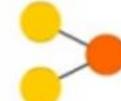


A Mostly Complete Chart of Neural Networks

(1 of 2)

- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

Perceptron (P)



Feed Forward (FF)



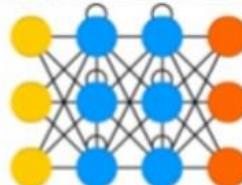
Radial Basis Network (RBF)



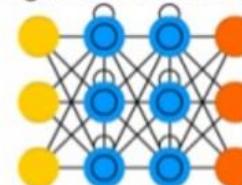
Deep Feed Forward (DFF)



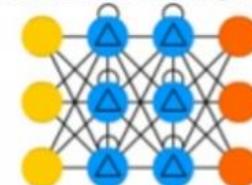
Recurrent Neural Network (RNN)



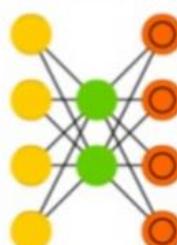
Long / Short Term Memory (LSTM)



Gated Recurrent Unit (GRU)



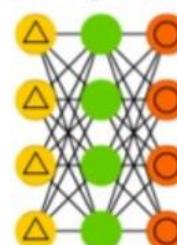
Auto Encoder (AE)



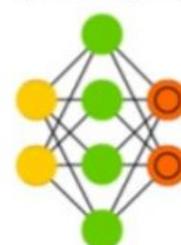
Variational AE (VAE)



Denoising AE (DAE)



Sparse AE (SAE)



Markov Chain (MC)



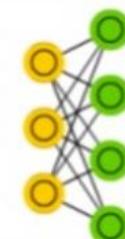
Hopfield Network (HN)



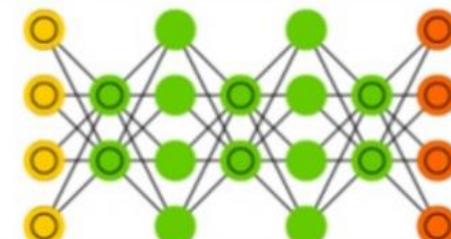
Boltzmann Machine (BM)



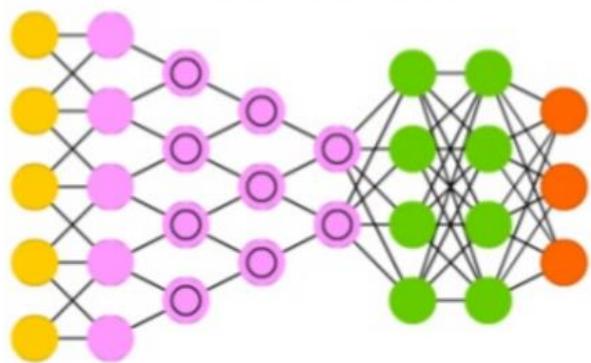
Restricted BM (RBM)



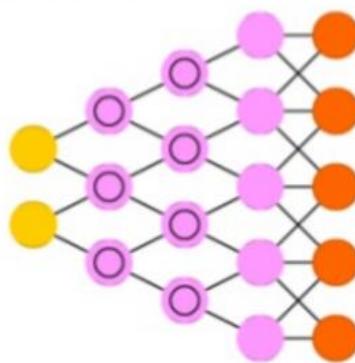
Deep Belief Network (DBN)



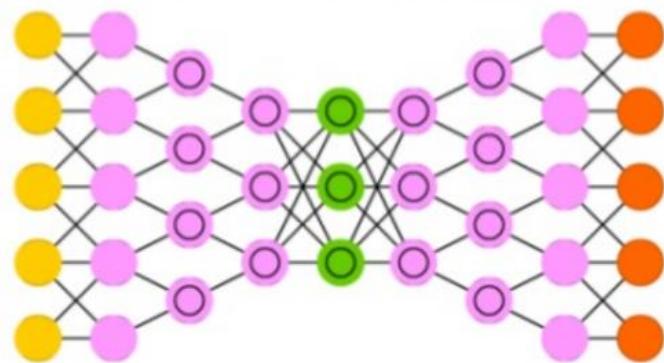
Deep Convolutional Network (DCN)



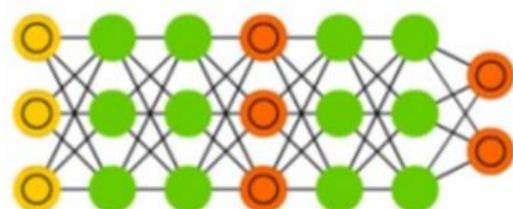
Deconvolutional Network (DN)



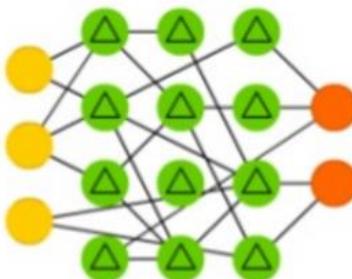
Deep Convolutional Inverse Graphics Network (DCIGN)



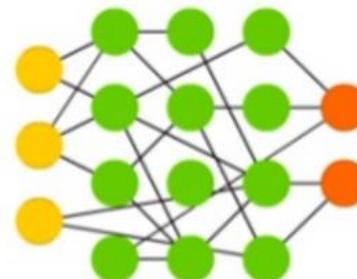
Generative Adversarial Network (GAN)



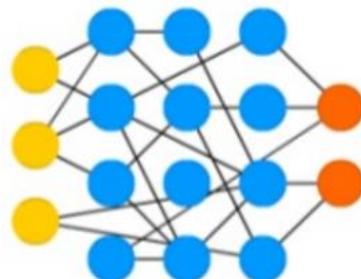
Liquid State Machine (LSM)



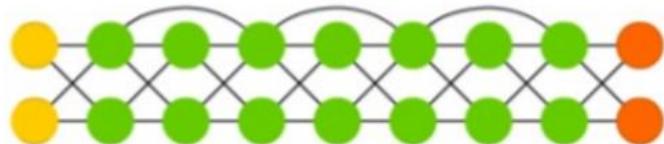
Extreme Learning Machine (ELM)



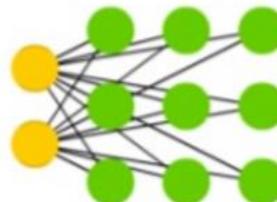
Echo State Network (ESN)



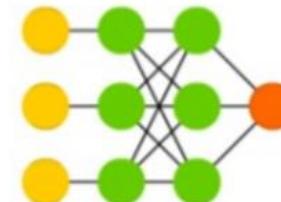
Deep Residual Network (DRN)



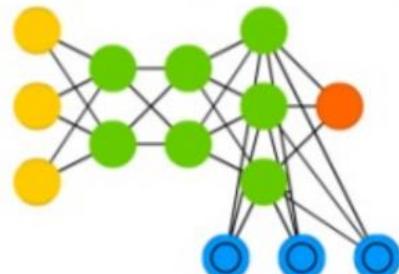
Kohonen Network (KN)



Support Vector Machine (SVM)



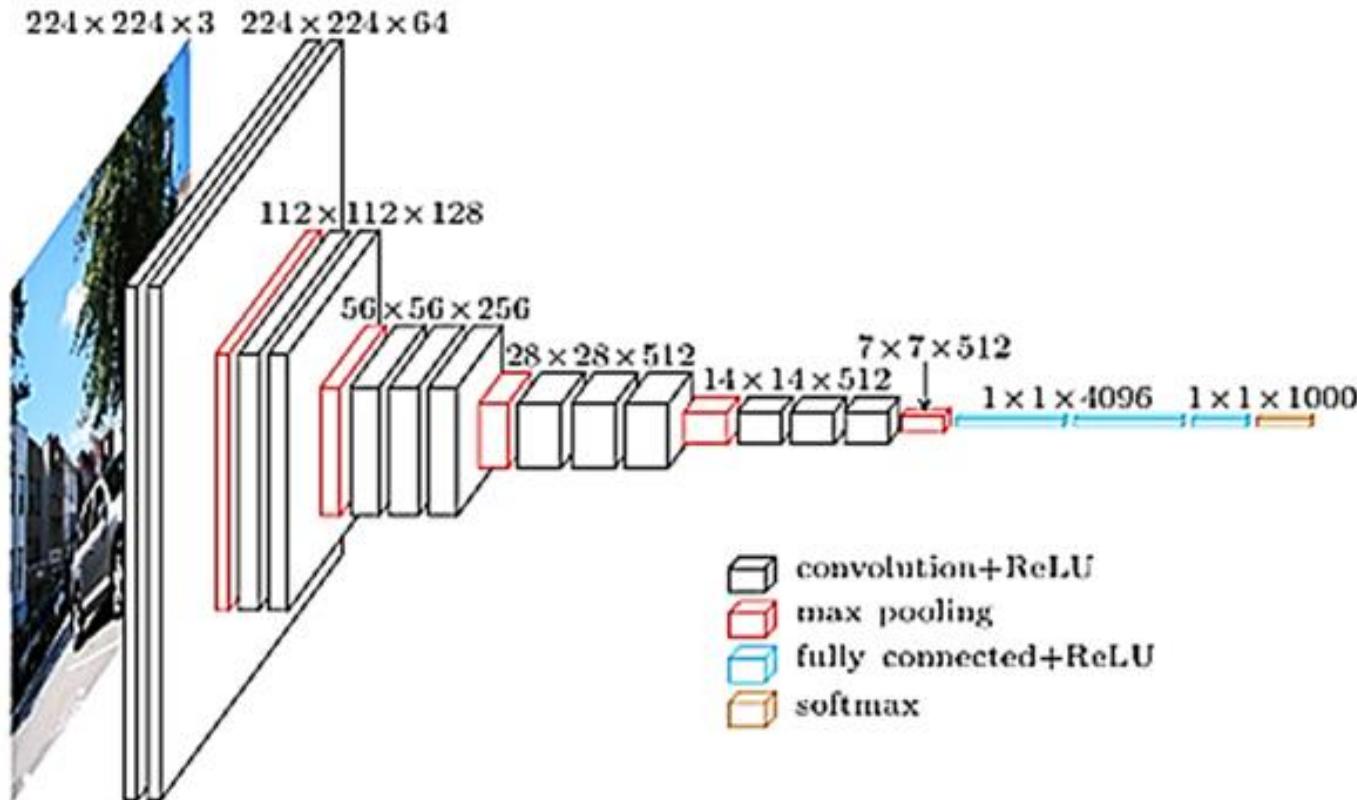
Neural Turing Machine (NTM)





CNN

CNNs



Consider learning an image:

- Some patterns are much smaller than the whole image

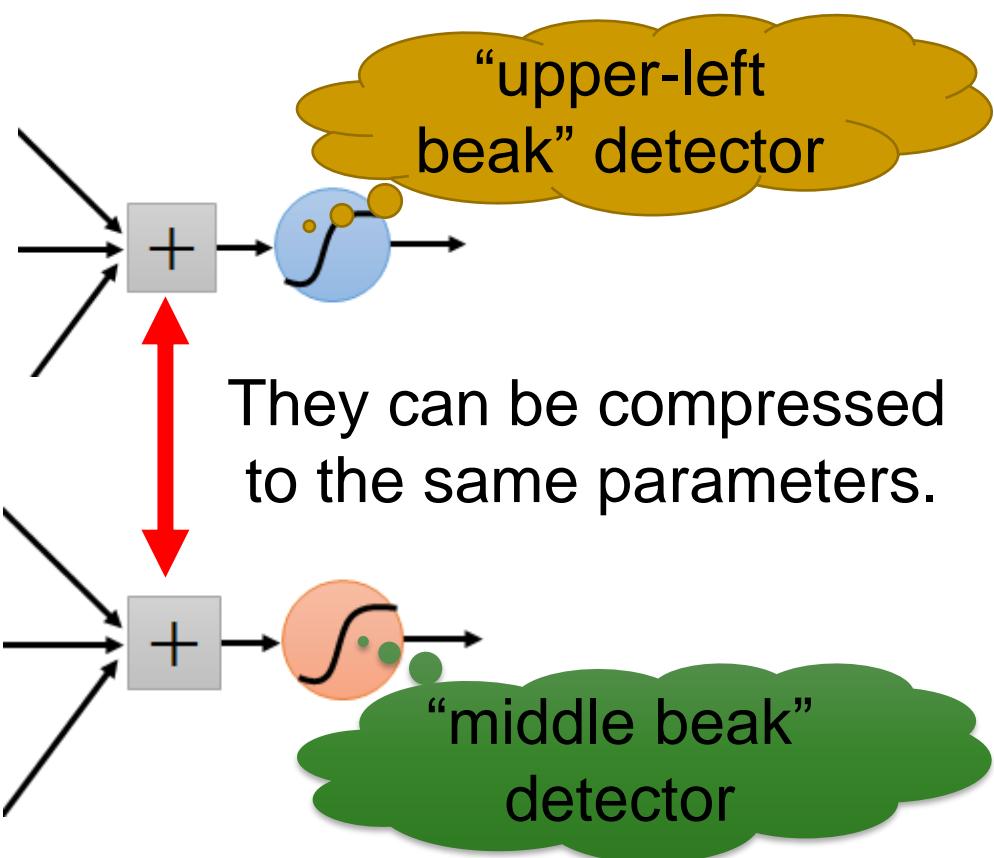
Can represent a small region with fewer parameters



Same pattern appears in different places:

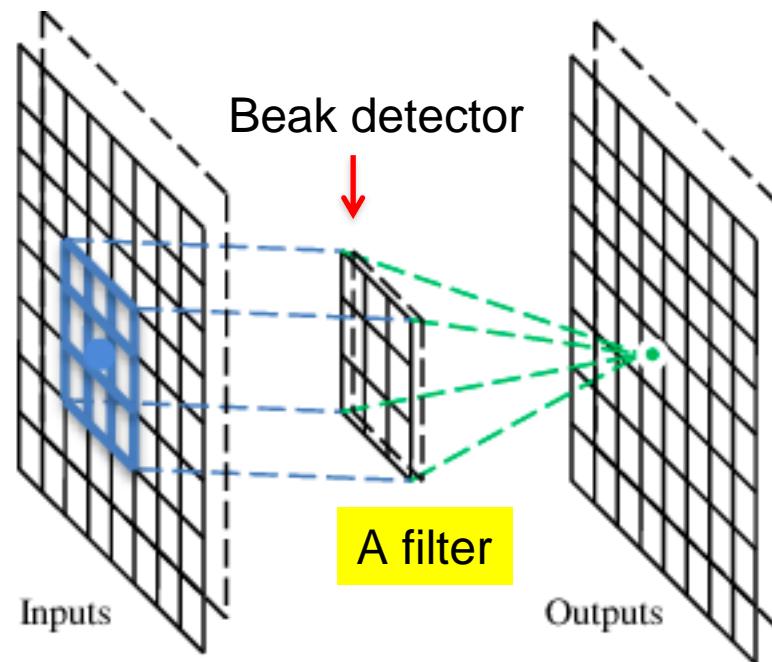
They can be compressed!

What about training a lot of such “small” detectors
and each detector must “move around”.



A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that do convolutional operation.



Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product



6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

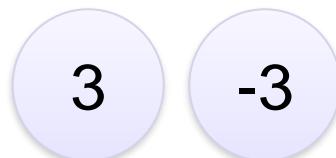
If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

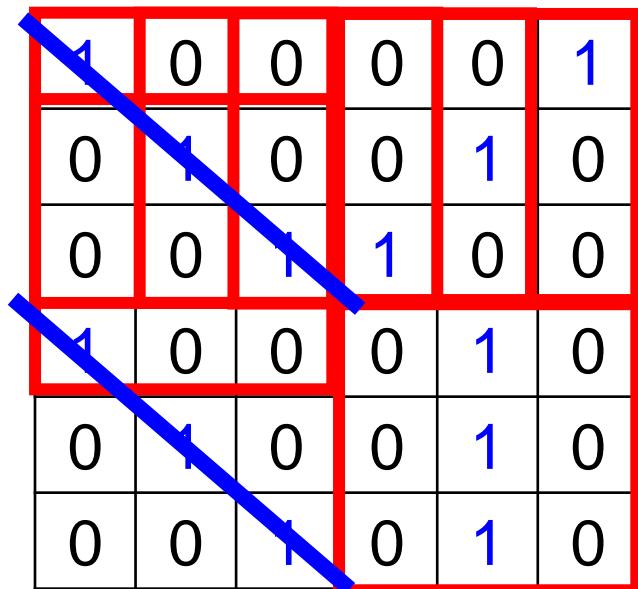
1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

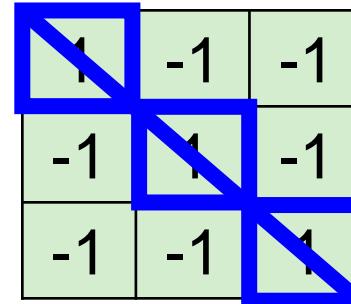


Convolution

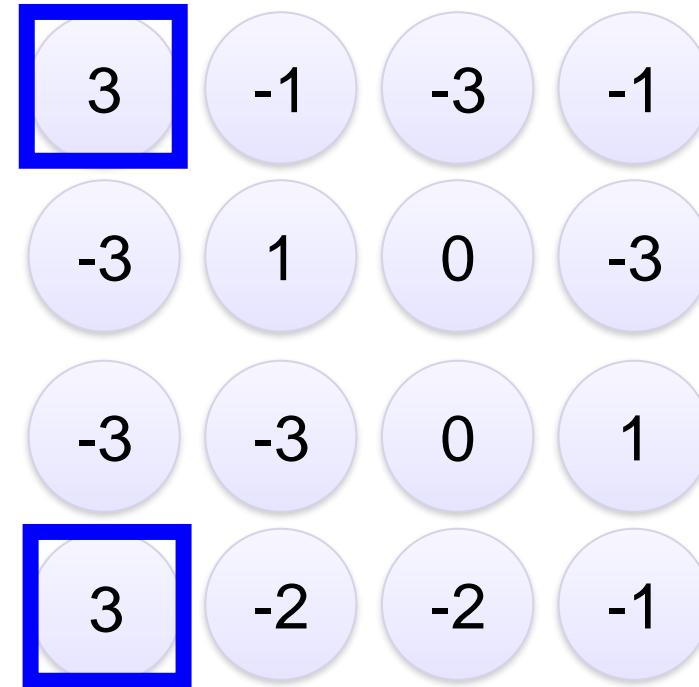
stride=1



6 x 6 image



Filter 1



Convolution

stride=1

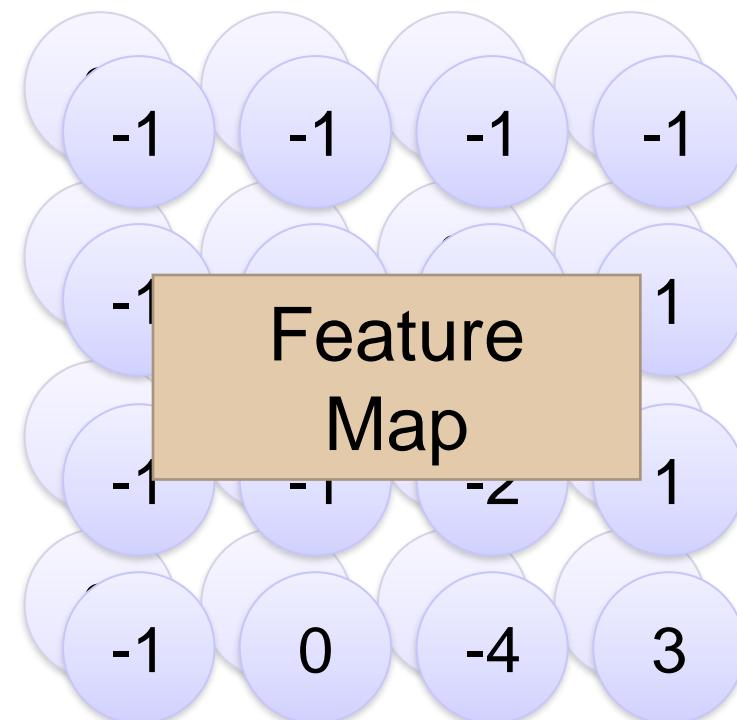
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

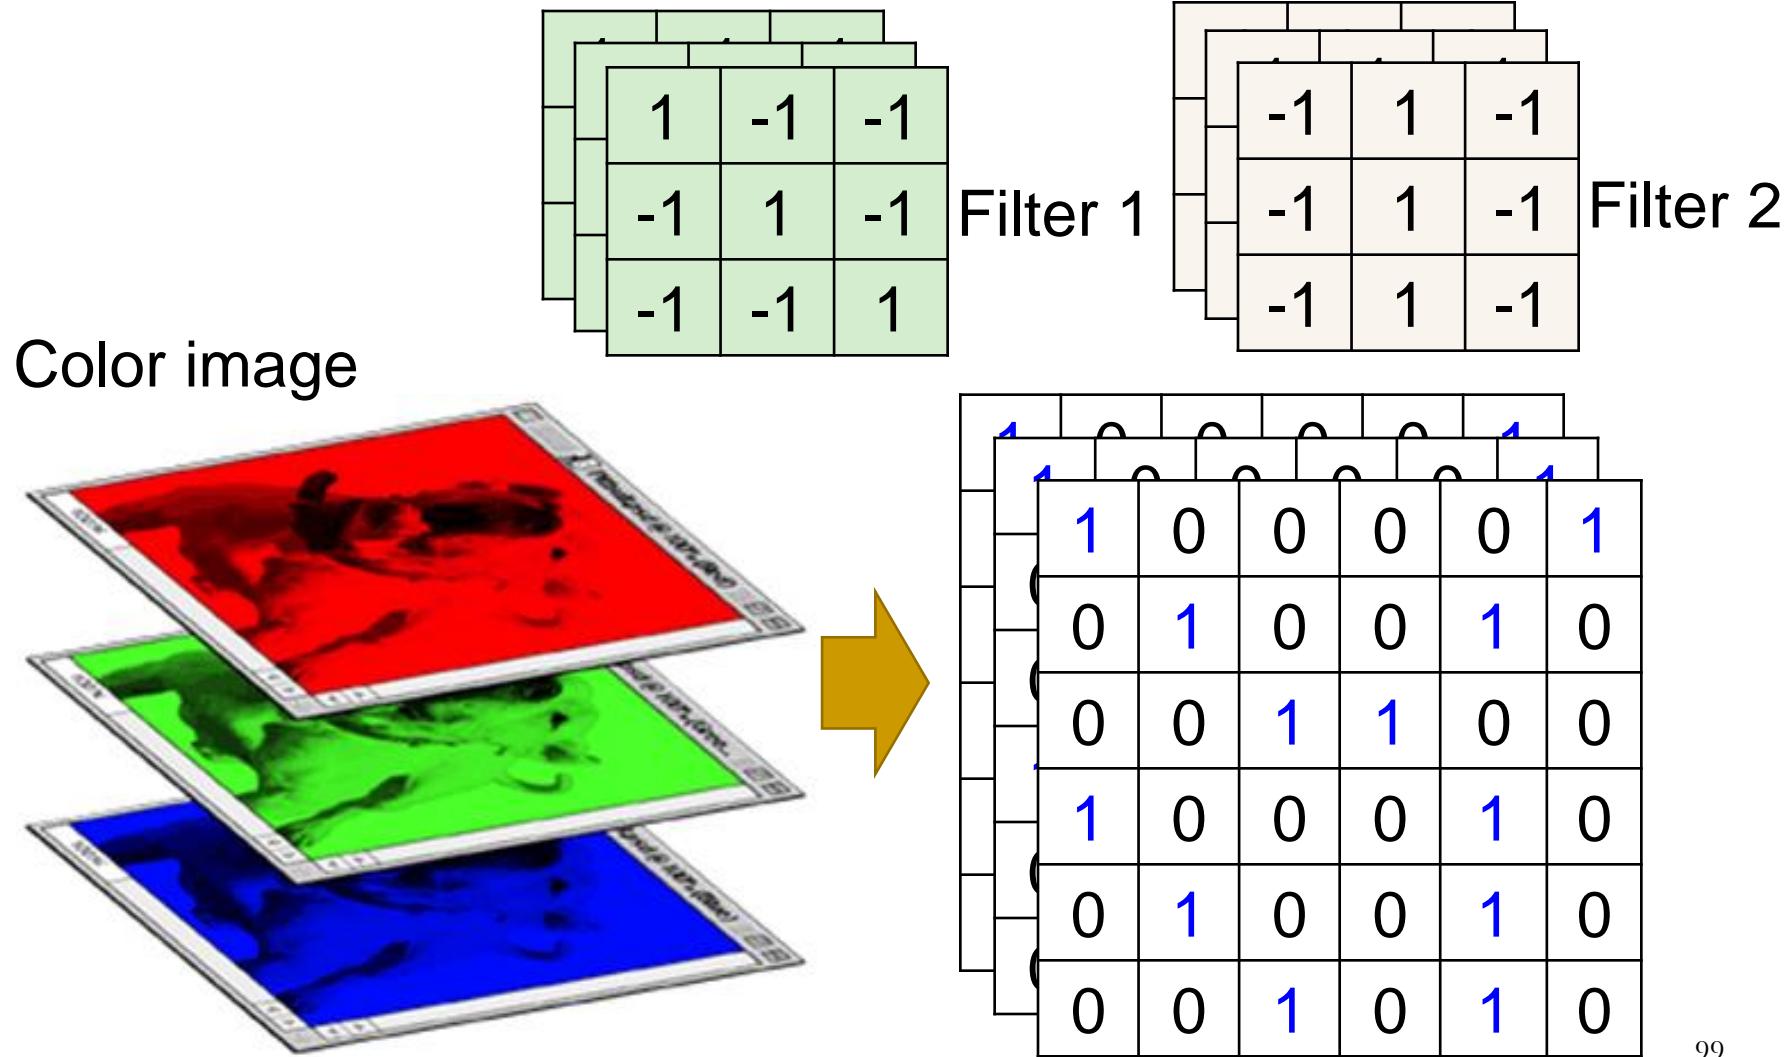
Filter 2

Repeat this for each filter

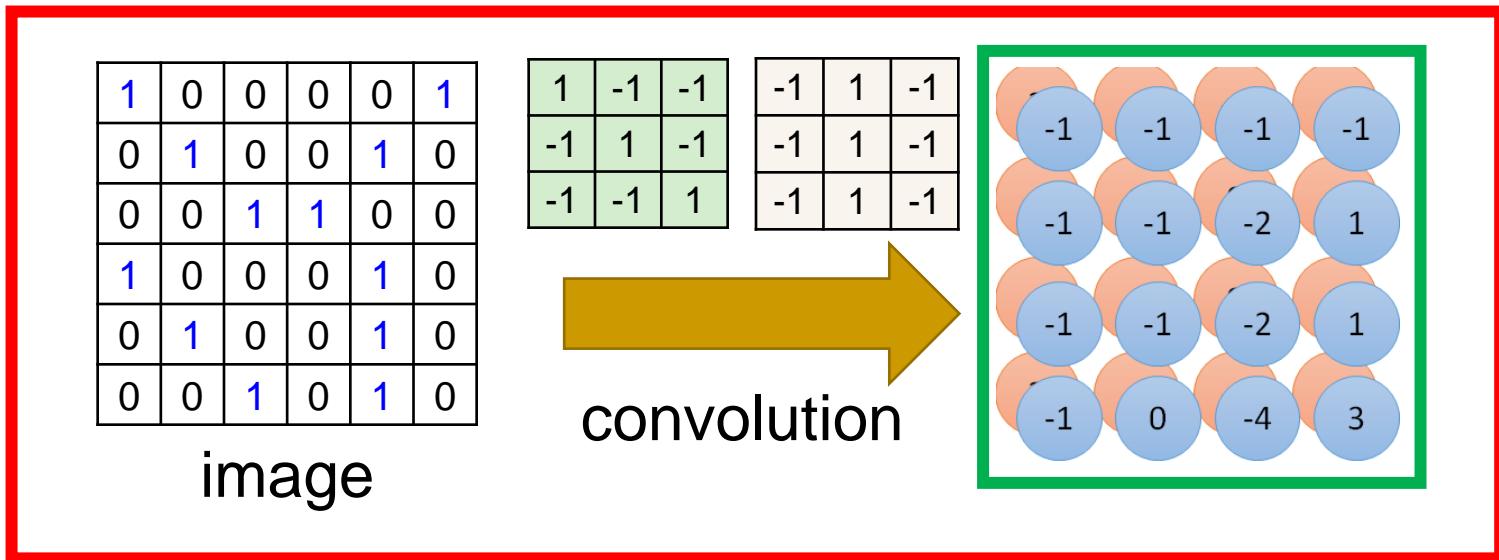


Two 4 x 4 images
Forming 2 x 4 x 4 matrix

Color image: RGB 3 channels

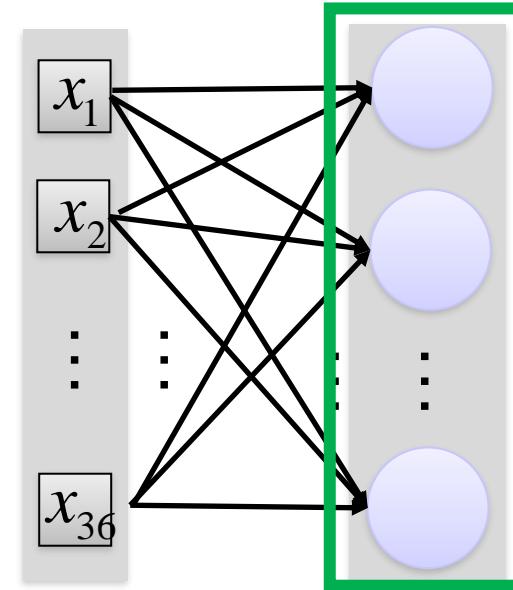


Convolution v.s. Fully Connected

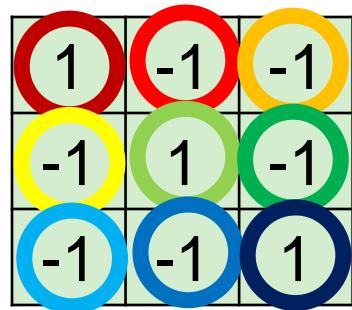


Fully-
connected

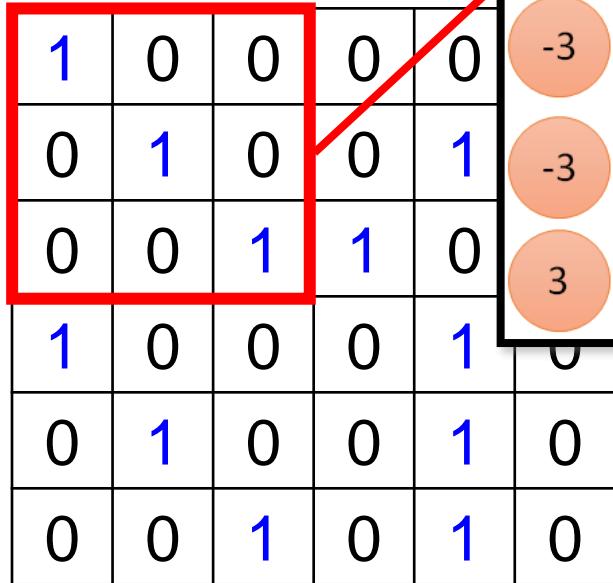
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



How to reduce parameters?

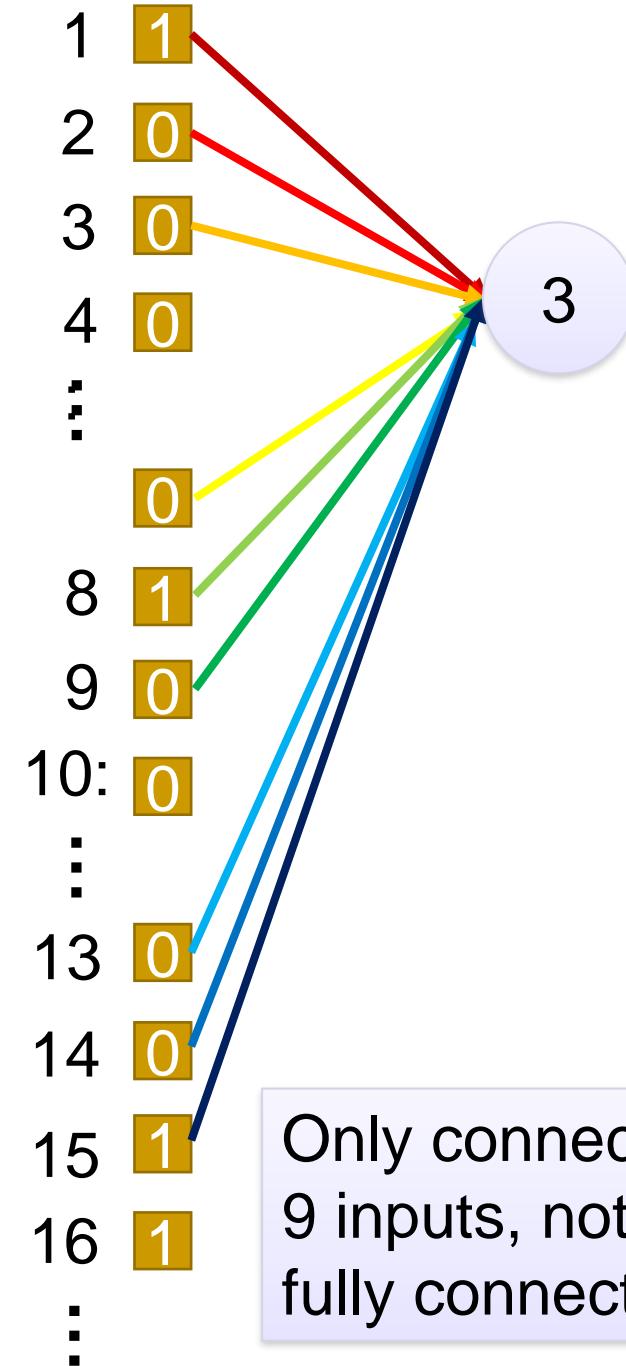
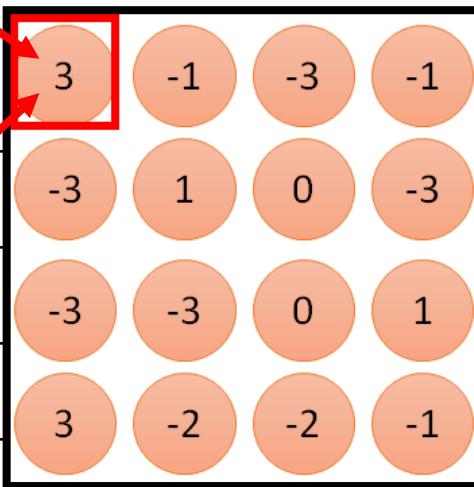


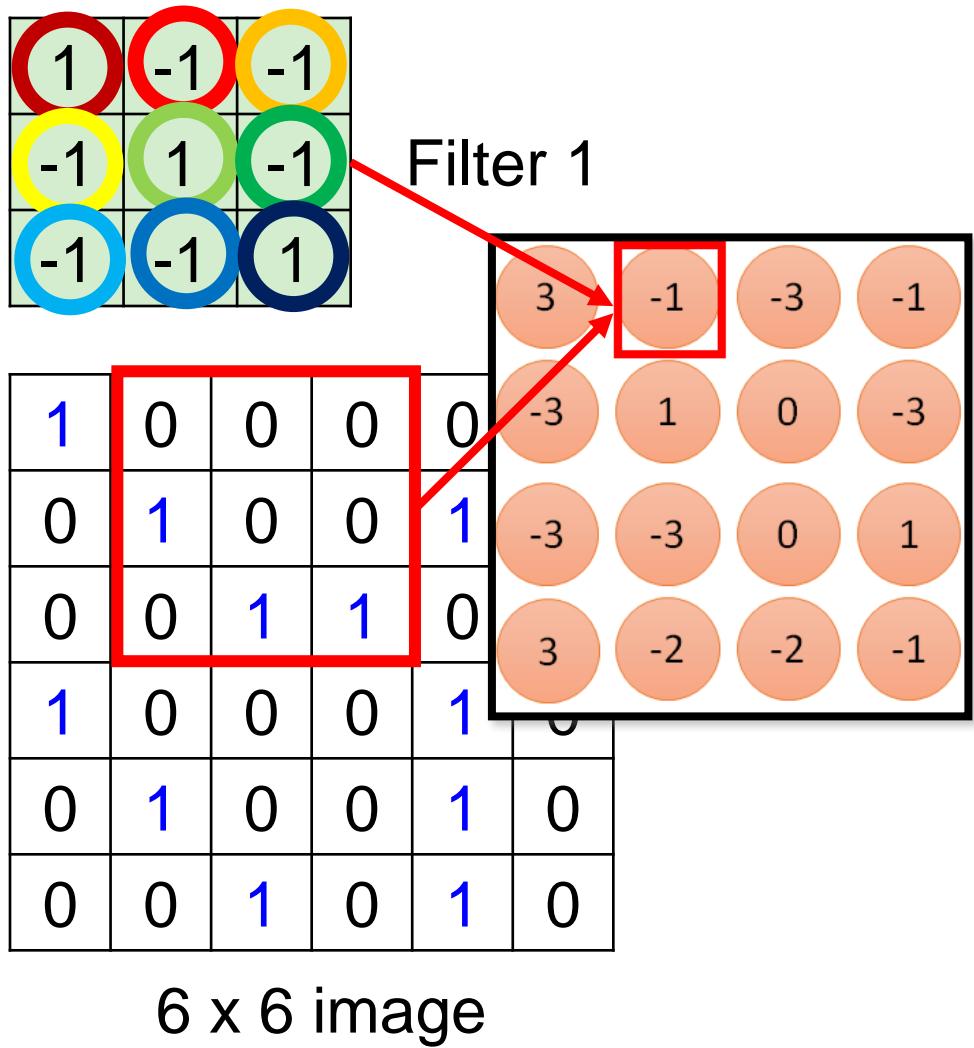
Filter 1



6×6 image

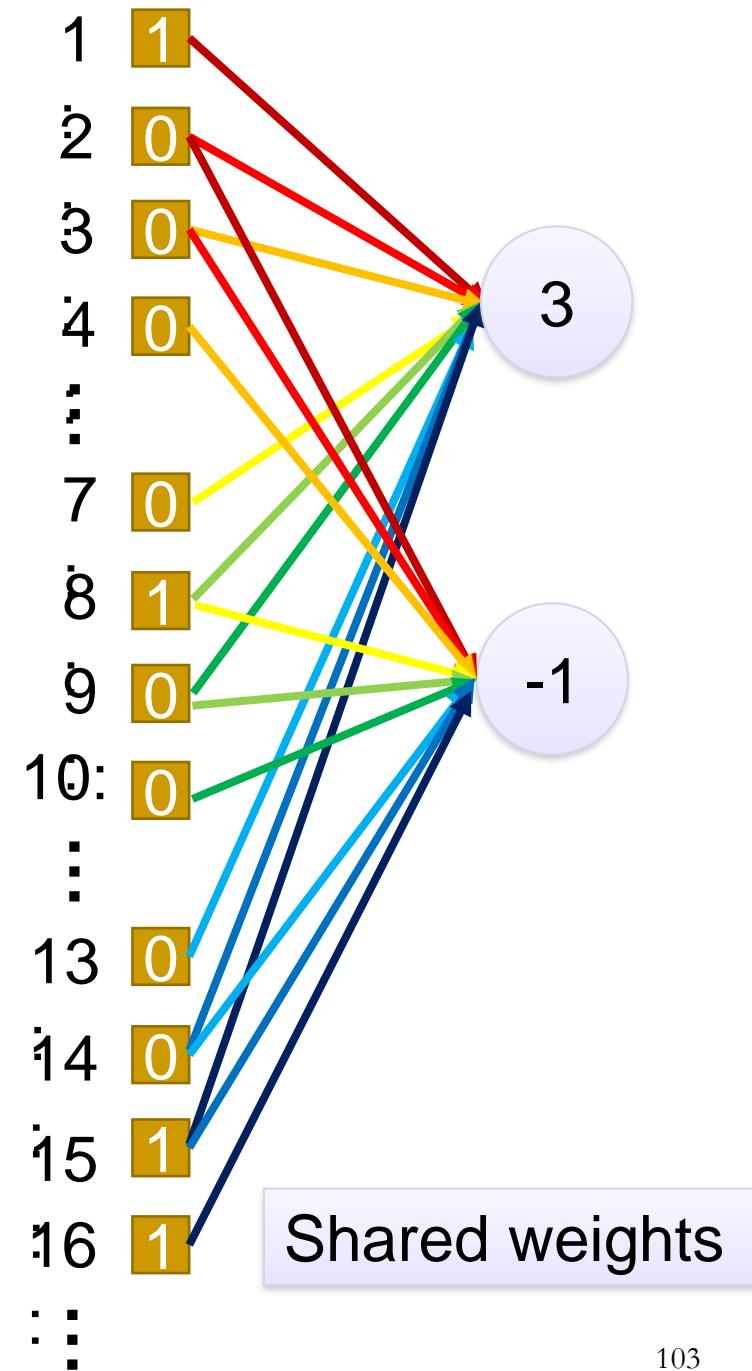
fewer parameters!





Fewer parameters

Even fewer parameters



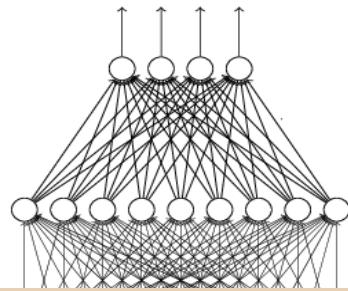
CNN – so far



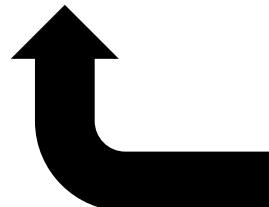
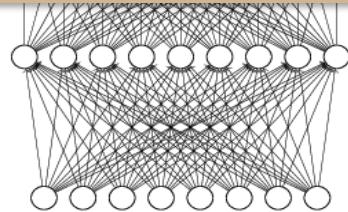
Convolution

The whole CNN

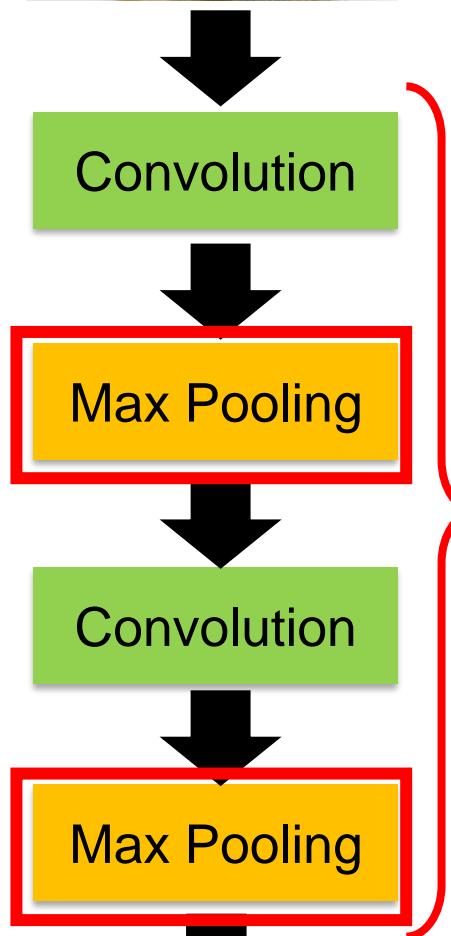
cat dog



Fully Connected
Feedforward network



Flattened



Can
repeat
many
times

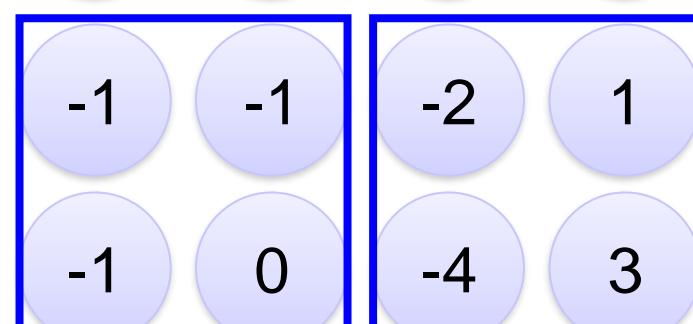
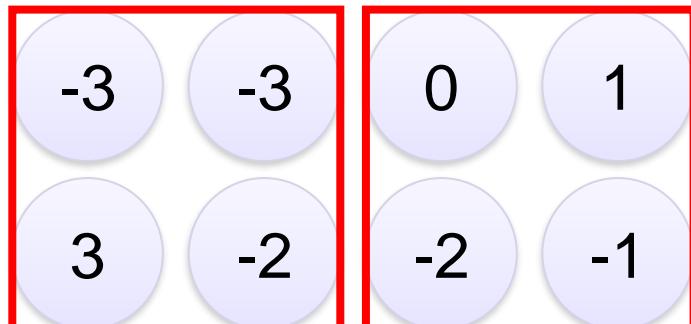
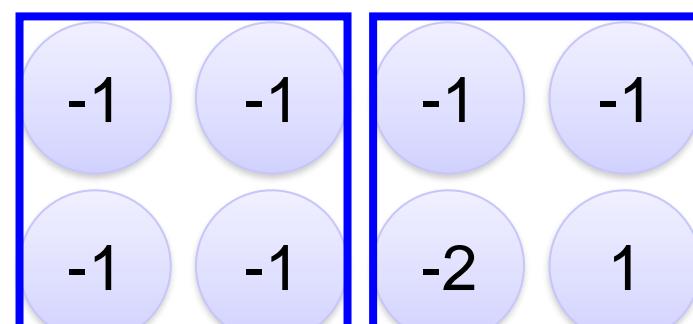
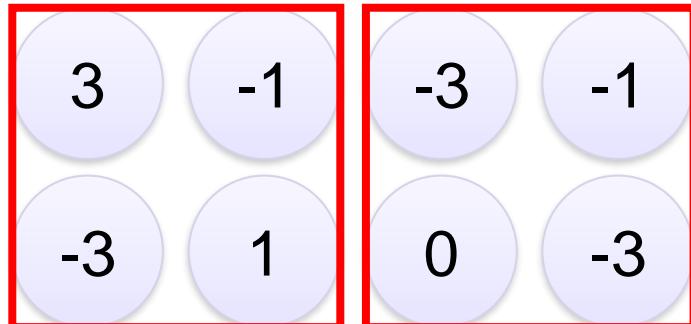
Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2



Why Pooling

- Subsampling pixels will not change the object bird



Subsampling



We can subsample the pixels to make image smaller
→ fewer parameters to characterize the image

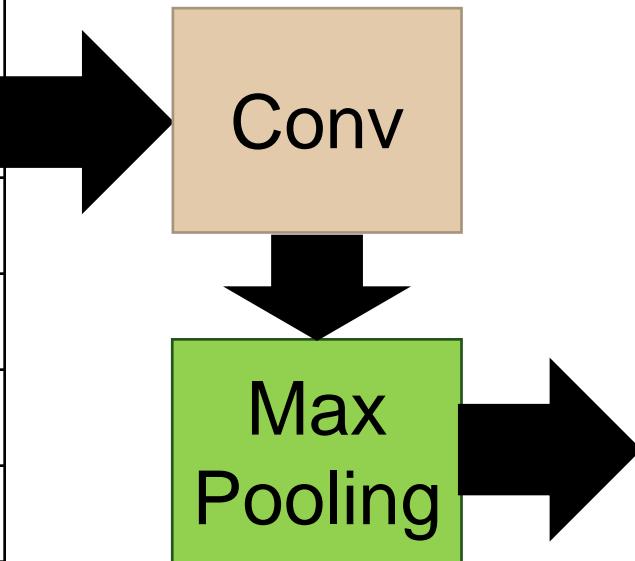
A CNN compresses a fully connected network in two ways:

- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

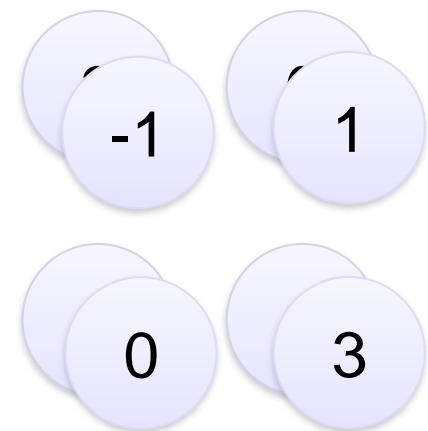
Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



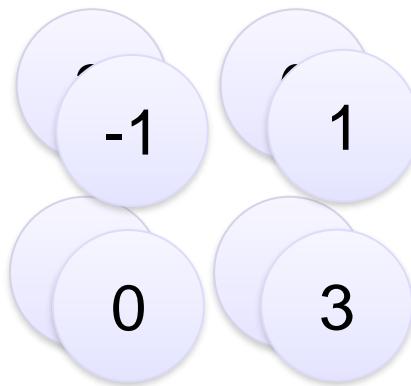
New image
but smaller



2 x 2 image

Each filter
is a channel

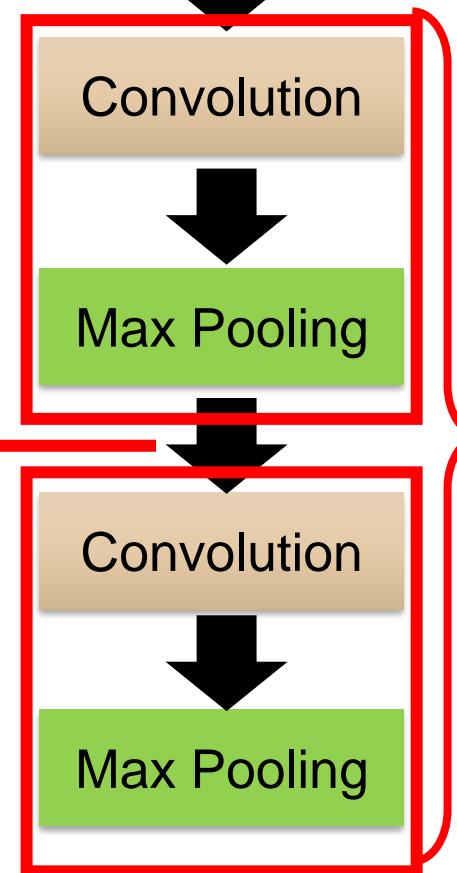
The whole CNN



A new image

Smaller than the original image

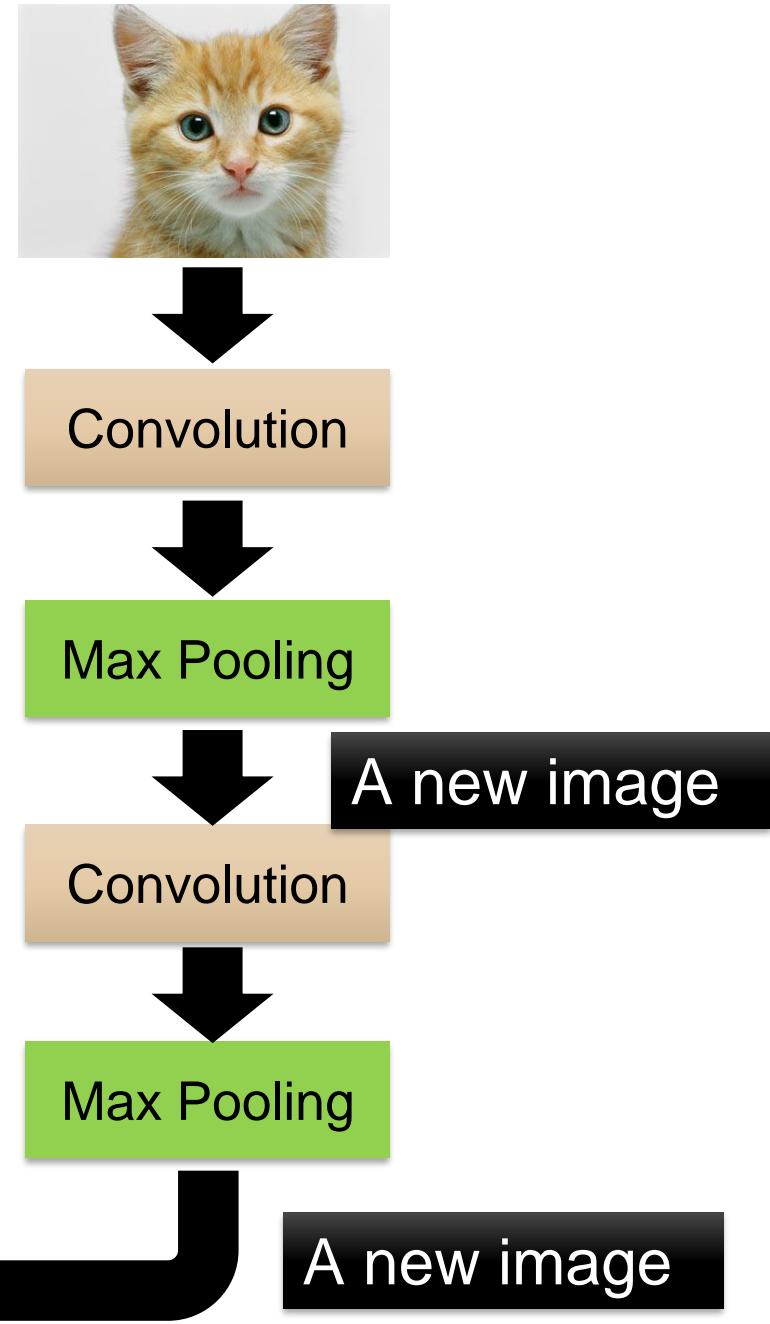
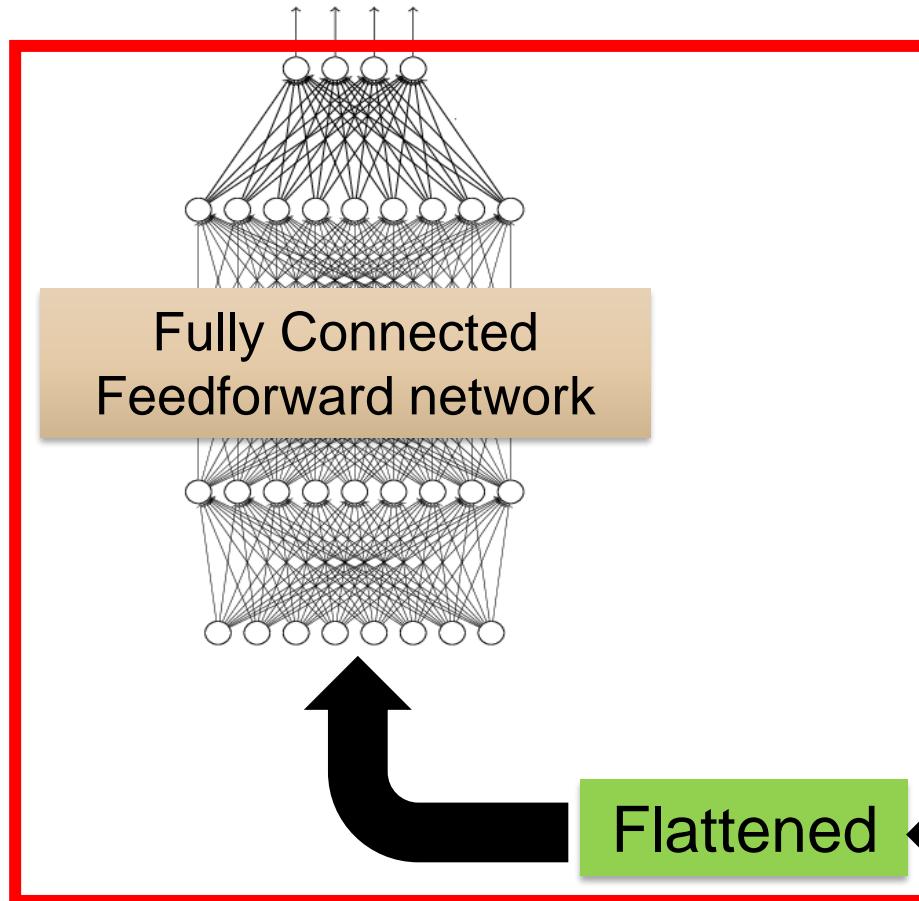
The number of channels is the number of filters



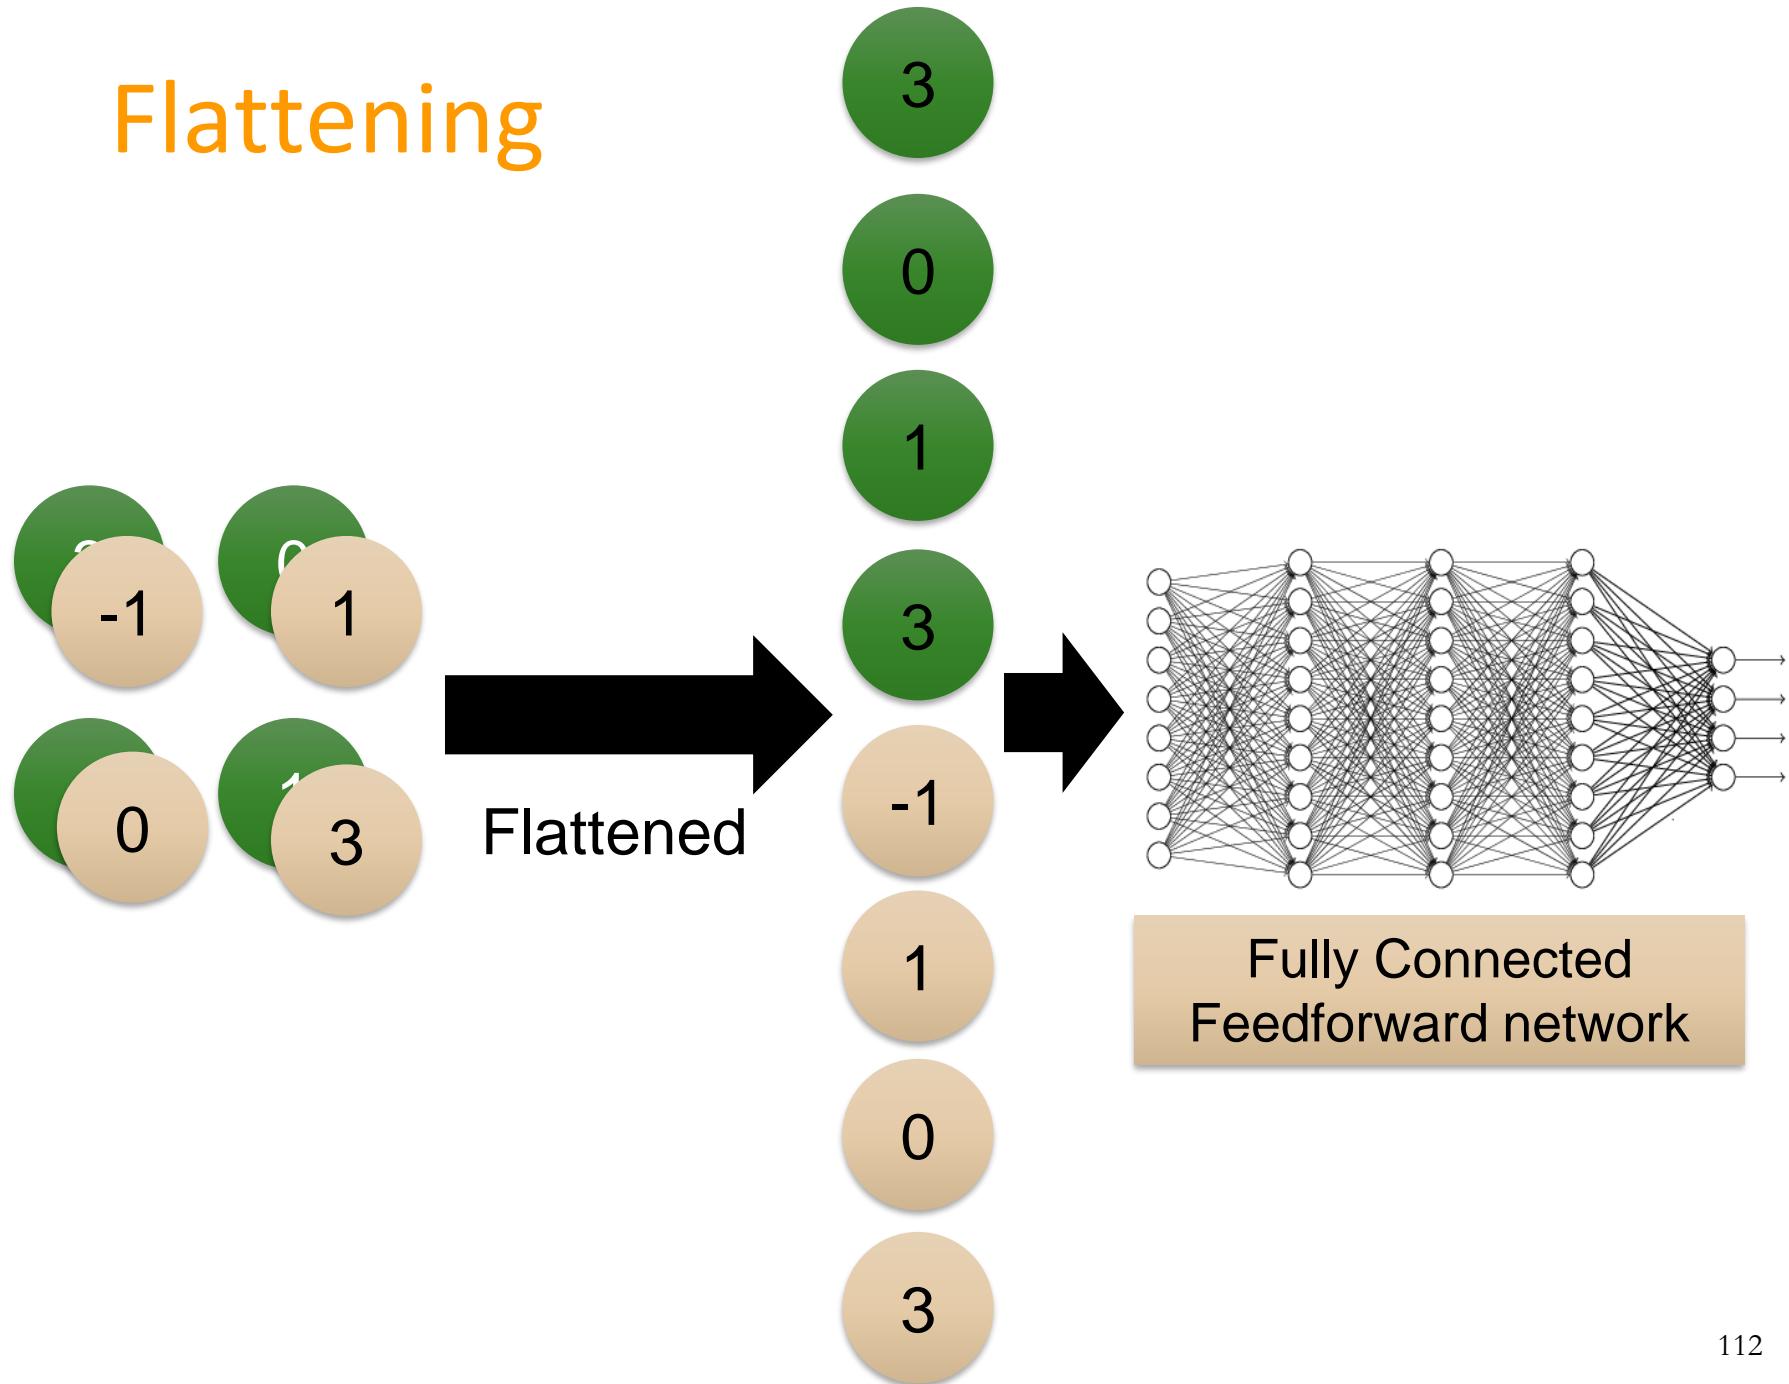
Can
repeat
many
times

The whole CNN

cat dog

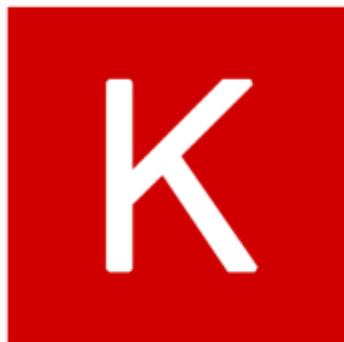


Flattening



Fully Connected
Feedforward network

Keras: The Python Deep Learning library



Keras

You have just found Keras.

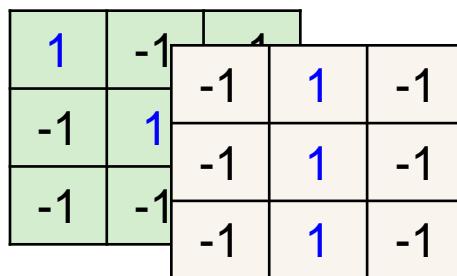
Keras is a high-level neural networks API, written in Python and capable of running on top of [TensorFlow](#), [CNTK](#), or [Theano](#). It was developed with a focus on enabling fast experimentation.
Being able to go from idea to result with the least possible delay is key to doing good research.

<https://keras.io/>

CNN in Keras

Only need to modify the *network structure* and *input format*

```
model2.add( Convolution2D( 25, 3, 3,  
                           input_shape=(28, 28, 1)) )
```



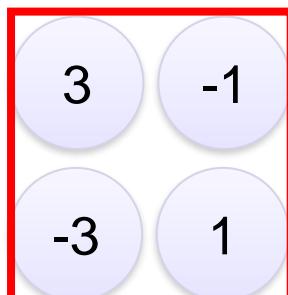
There are
25 3x3
filters.

Input_shape = (28 , 28 , 1)

28 x 28 pixels

1: black/white, 3: RGB

```
model2.add(MaxPooling2D( (2, 2) ))
```



Input
↓

Convolution

↓

Max Pooling

↓

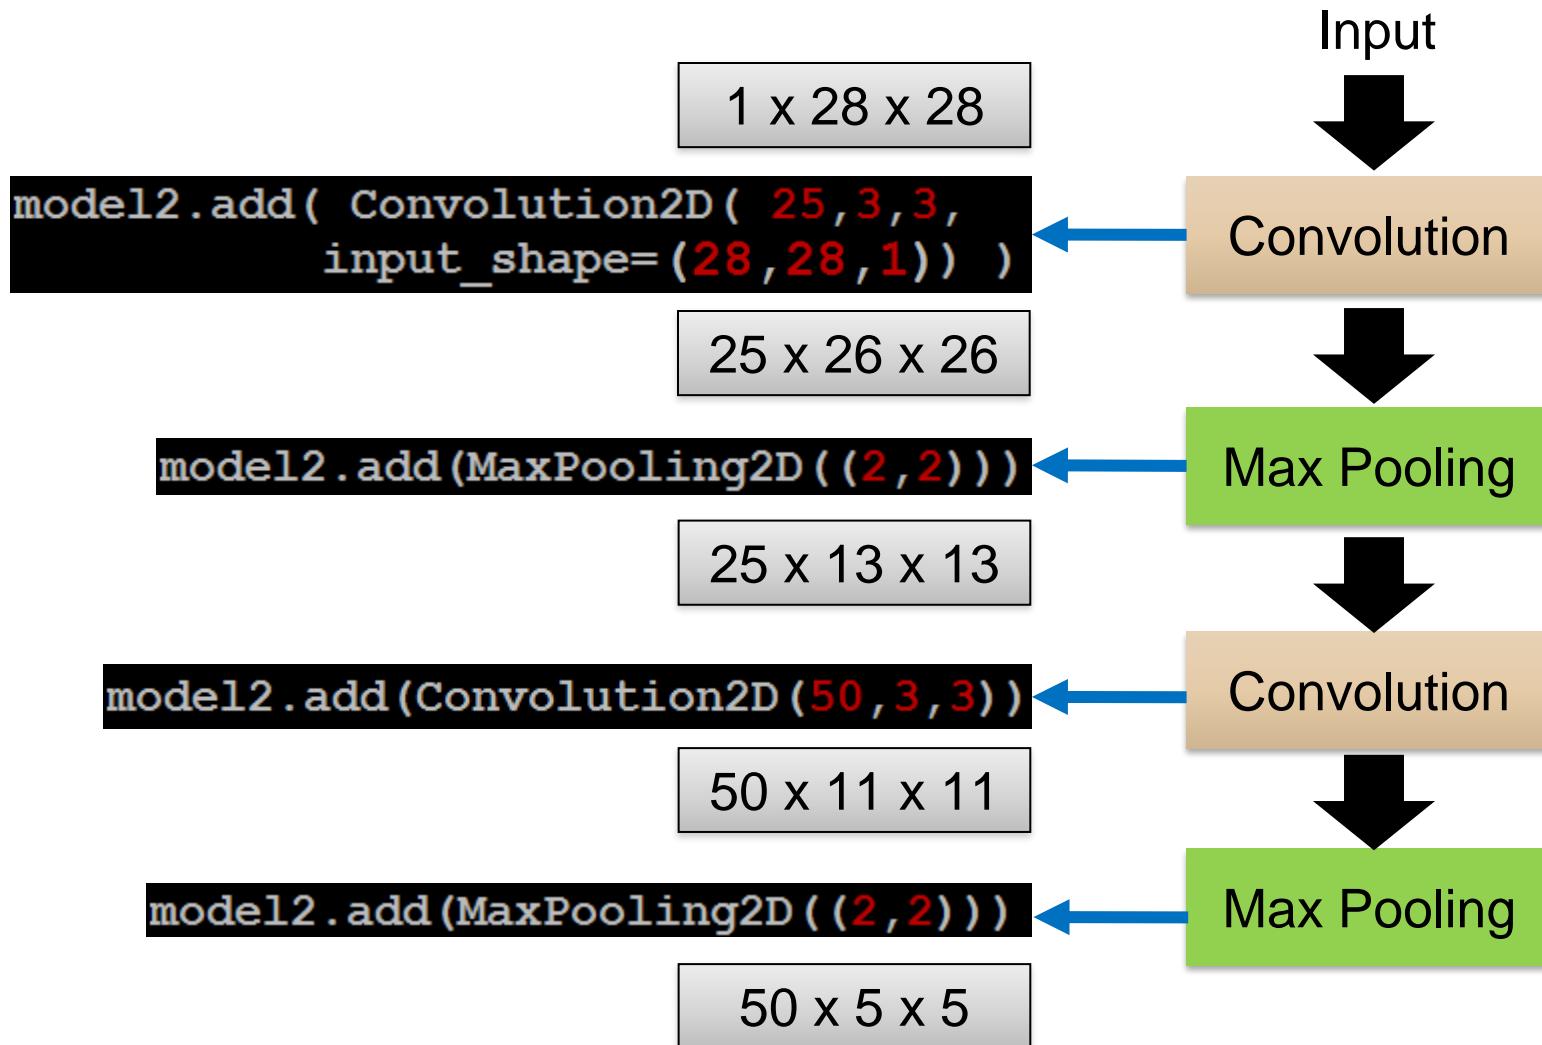
Convolution

↓

Max Pooling

CNN in Keras

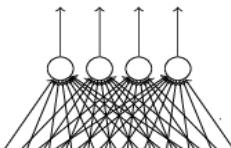
Only need to modify the *network structure* and *input format*



CNN in Keras

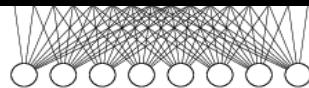
Only need to modify the *network structure* and *input format*

Output



Fully connected
feedforward network

```
model2.add(Dense(output_dim=100))  
model2.add(Activation('relu'))  
model2.add(Dense(output_dim=10))  
model2.add(Activation('softmax'))
```



1250

Flattened

```
model2.add(Flatten())
```

Input

$1 \times 28 \times 28$



Convolution

$25 \times 26 \times 26$



Max Pooling

$25 \times 13 \times 13$



Convolution

$50 \times 11 \times 11$



Max Pooling

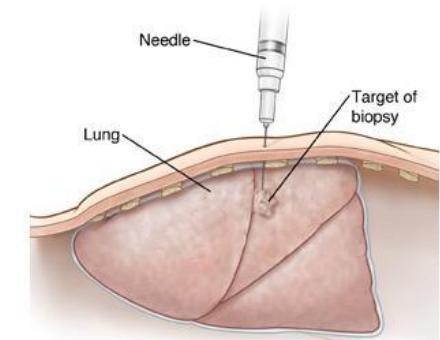
$50 \times 5 \times 5$



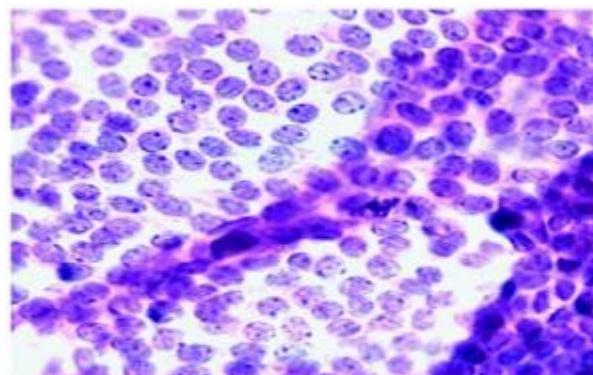
**THIS WEEK:
PROJECT
3**

Project Overview

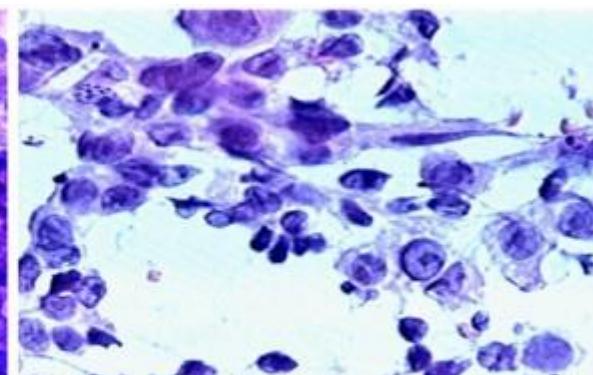
- Imagine, you are given data about patients having some suspicious tissue, e.g. in lung or breast.
- Digitized images are given and have labels “benign” or “malignant”.



For the biopsy, a thin needle is used to remove samples of tissue from an abnormal area in the lung.



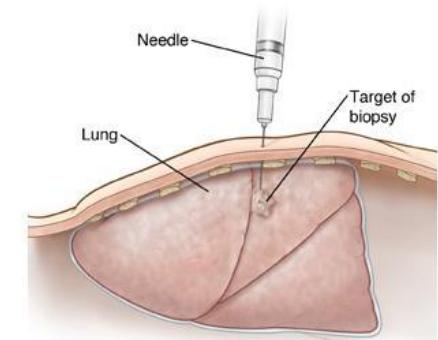
Smear with BENIGN diagnosis – uniform nucleus of cells, symmetrical, homogeneous, with areas within normal size



Smear with MALIGNANT diagnosis – nucleus of cells without uniformity, asymmetrical, not homogeneous (multiple sizes) and with areas above normal size

Project Overview

- Imagine, you are given data about patients having some suspicious tissue, e.g. in lung or breast.
- Digitized images are given and have labels “benign” or “malignant”.
- You want to build a tool that can help medical doctors to decide whether a sample is benign or malign.
- How would you do that?



For the biopsy, a thin needle is used to remove samples of tissue from an abnormal area in the lung.

